



Lecture 1



presentation

DAD – Distributed Applications Development

Cristian Toma

D.I.C.E/D.E.I.C – Department of Economic Informatics & Cybernetics

www.dice.ase.ro



Cristian Toma – Business Card



Cristian Toma

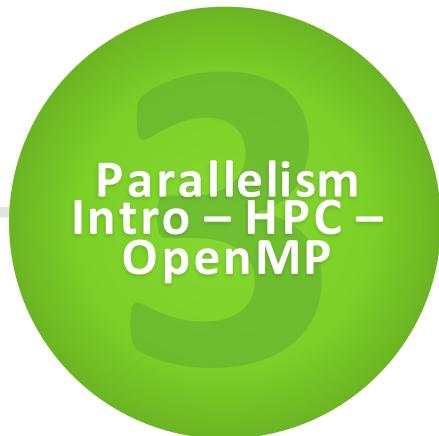
IT&C Security Master

Dorobantilor Ave., No. 15-17
010572 Bucharest - Romania

<http://ism.ase.ro>
cristian.toma@ie.ase.ro
T +40 21 319 19 00 - 310
F +40 21 319 19 00



Agenda for Lecture 1





Distributed Systems and Distributed Applications Development

Distributed Systems

1.1 DAD Lectures Structure

Main administrative details:

Didactic Activities: Lectures 50% + Lab / Seminar 50%
14 meetings **14 meetings**

Evaluation: PC Exam – 70% / Seminars tests & assignments – 30%

E-Framework: VMs – VM-Ware Virtual Machines with:

- Linux Ubuntu LTS + Java – JDK + Eclipse IDE + Apache Tomcat + Apache Kafka + Apache Spark + Vert.x + GCC

E-Learning Platform: **<http://acs.ase.ro/dad>**

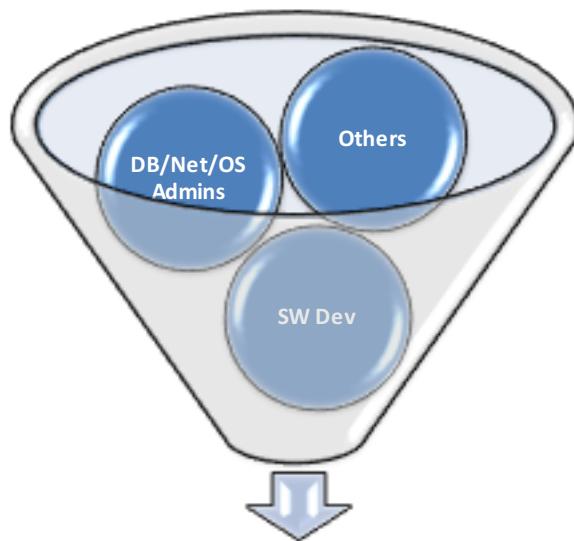
SAKAI – <http://ism.ase.ro> | **<http://acs.ase.ro/java>** | <http://online.ase.ro>

Prerequisites: Fundamentals of Java SE (8 & 9) + C/C++ ('11 & '14) + Networking + Node.js + Linux/Windows OS | Optional – Python

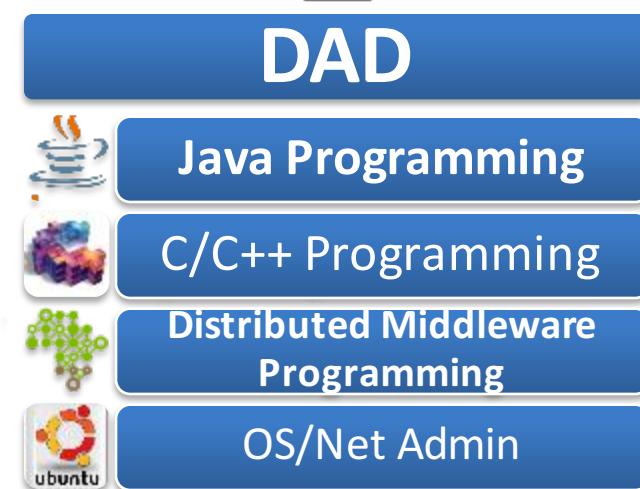
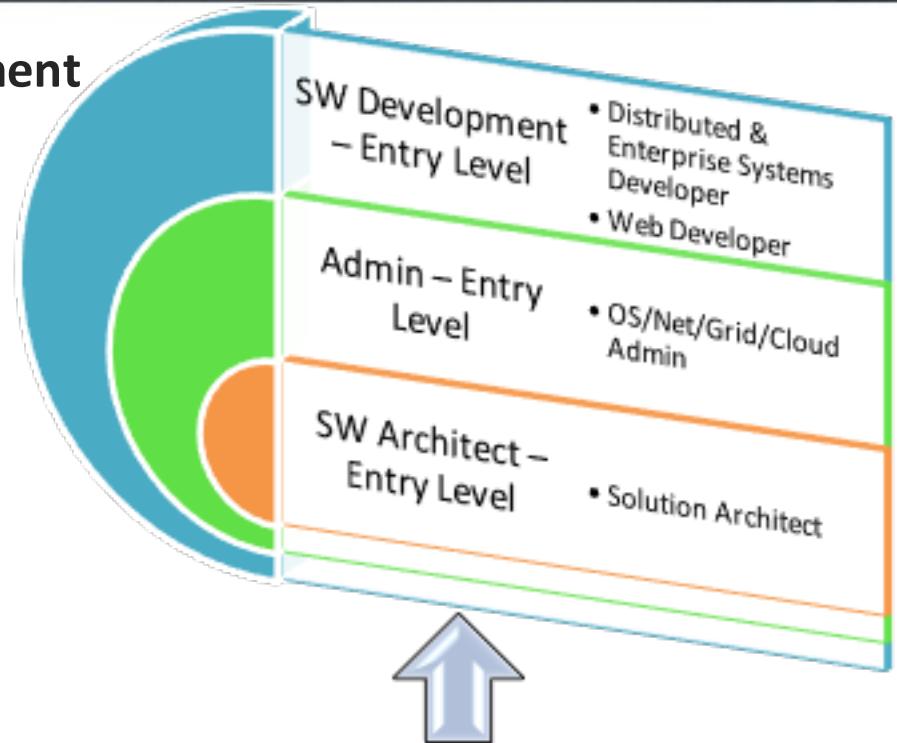
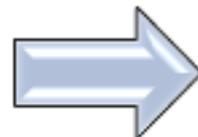
Mission: Technological transfer from university to the students of practical and theoretical issues related with distributed applications development.

1.2 Target Group Profile

DAD – Distributed Applications Development



DAD needs students with
C/C++, C#, Java, Networking, OS
Knowledge of Fundamentals





It's not just about the programming, but providing smart solutions

DAD Sections & References

What about the DAD as it is @ Harvard/MIT?

Could you provide a solution for finding out the biggest mark in the class?

Do we have unicast, multicast, broadcast messages or client-server, P2P / hybrid paradigms?...GREAT...Please upload the solution in Java / C# / C/C++ / Python / Ruby till next week 23:50 in e-learning platform – SAKAI...I'm NOT kidding...

1.3 Distributed Systems

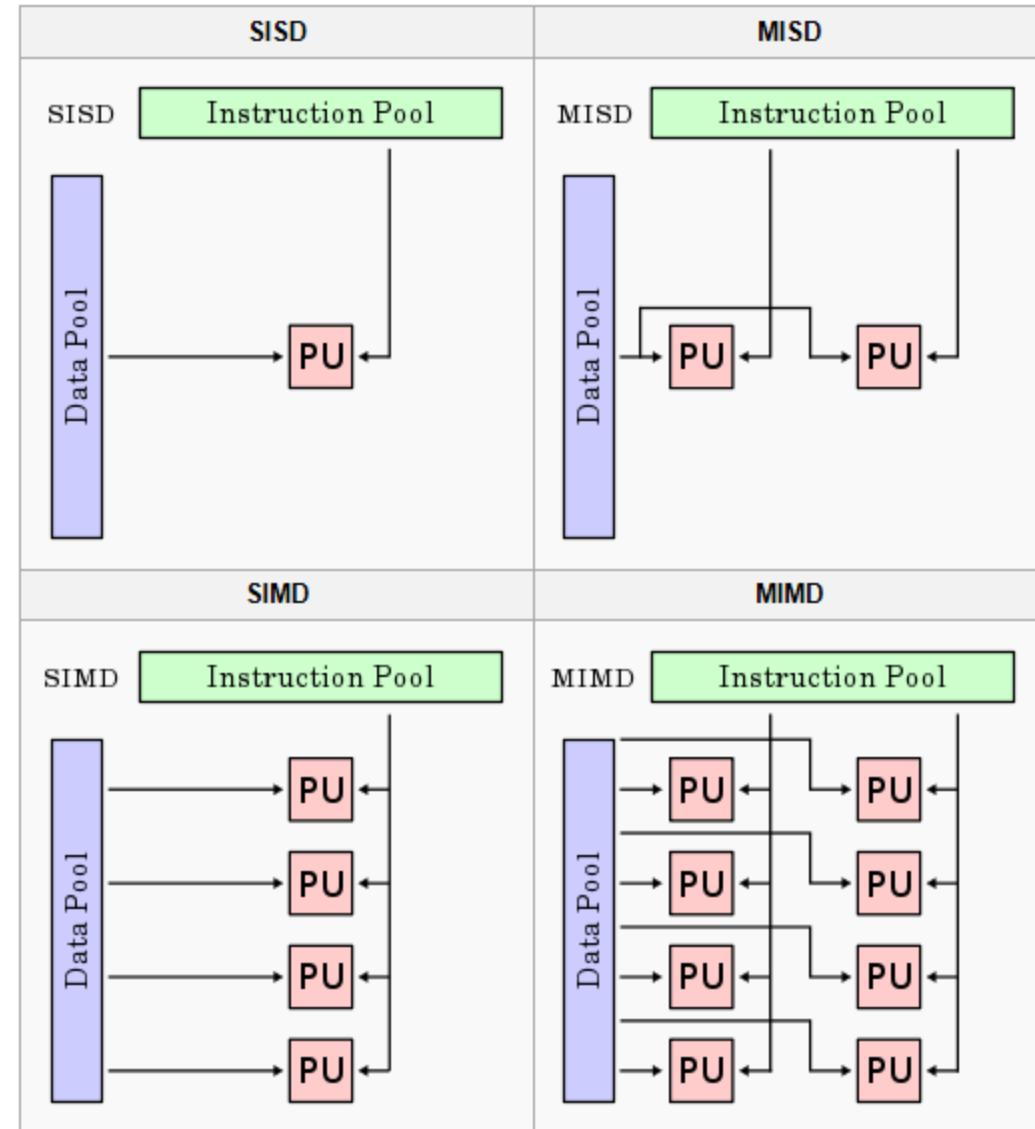
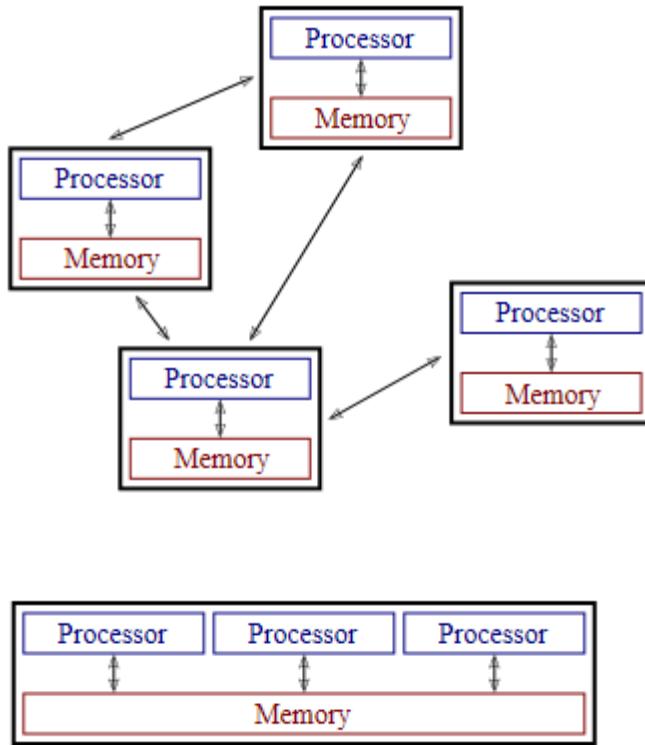
What is a Distributed System?

A distributed system is one in which components located into computers connected in network, communicate and coordinate their actions only by passing messages (sometimes, in addition, by migrating processes), in order to resolve a set of problems.

A distributed system consists of a collection of autonomous computers linked by a computer network and equipped with distributed system software. This software enables computers to coordinate their activities and to share the re- sources of the system hardware, software, and data.

Flynn Taxonomy Parallel vs. Distributed Systems

Parallel vs. Distributed Computing / Algorithms



Where is the picture for:
Distributed System and **Parallel System**?

http://en.wikipedia.org/wiki/Distributed_computing
http://en.wikipedia.org/wiki/Flynn's_taxonomy

1.3 Distributed Systems

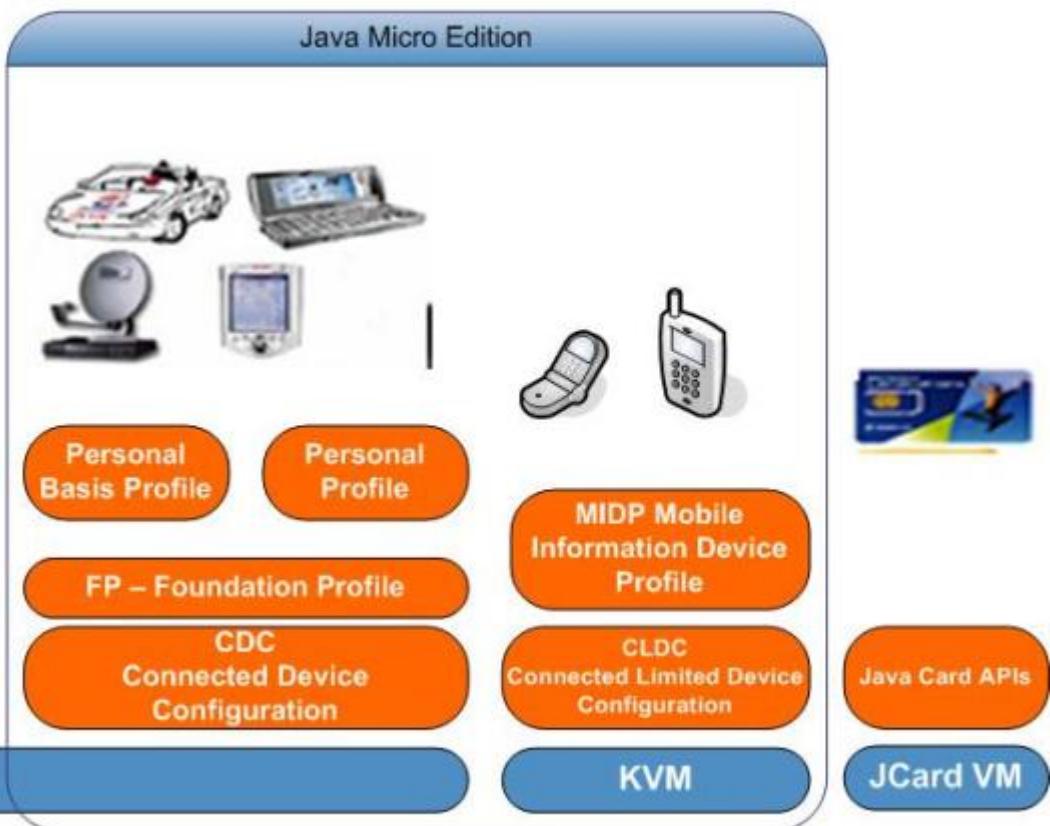
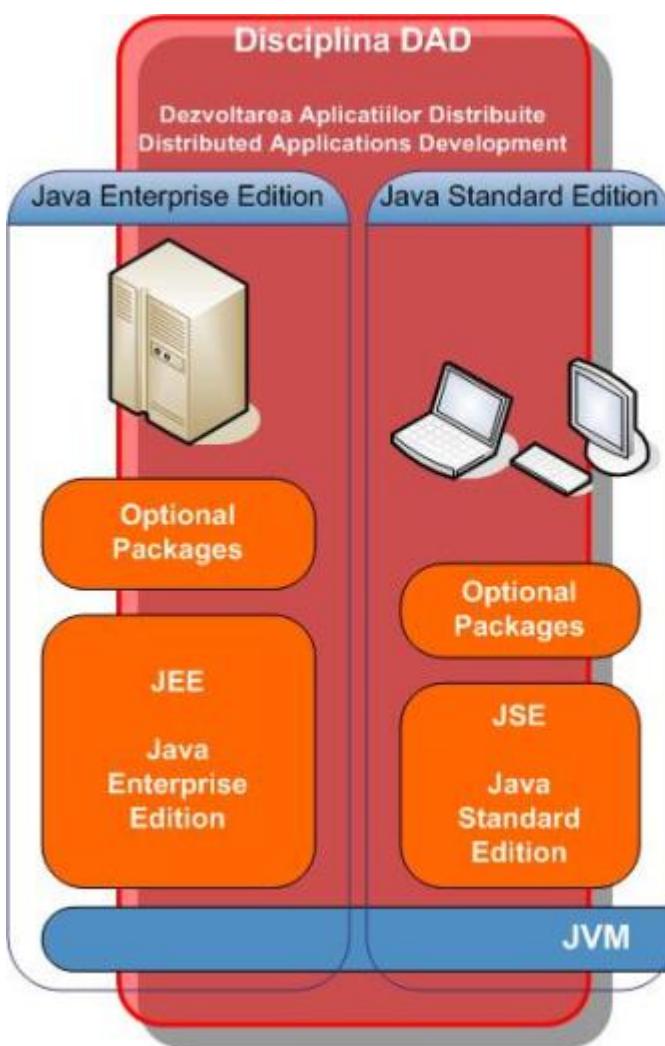
How to characterize a distributed system?

- concurrency of components
- lack of global clock
- independent failures of components

Leslie Lamport :-)

You know you have a distributed system when the crash of a computer you've never heard of stops you from getting any work done!

Prime motivation = to share the resources



1.3 Distributed Systems

What are the challenges?

- heterogeneity of their components
- openness
- security
- scalability – the ability to work well when the load or the number of users increases
- failure handling
- concurrency of components
- transparency
- providing quality of service

1.3 Distributed Systems

What are the resources?

- Hardware
 - Not every single resource is for sharing
- Data
 - Databases
 - Proprietary software
 - Software production
 - Collaboration

Sharing Resources

- Different resources are handled in different ways, there are however some generic requirements:
 - Namespace for identification
 - Name translation to network address
 - Synchronization of multiple access

1.3 Distributed Systems

Concept of Distributed Architecture

A distributed system can be demonstrated by the client-server architecture, which forms the base for multi-tier architectures; alternatives are the broker architecture such as CORBA, and the Service-Oriented Architecture (SOA). In this architecture, information processing is not confined to a single machine rather it is distributed over several independent computers.

There are several technology frameworks to support distributed architectures, including Java EE / Java-Kotlin Reactive Micro-services, .NET, CORBA, Scala Akka, Globus Toolkit Grid services, etc..

Middleware is an infrastructure that appropriately supports the development and execution of distributed applications.

1.3 Distributed Systems

Basis of Distributed Architecture

The basis of a distributed architecture is its transparency, reliability, and availability.

The following table lists the different forms of transparency in a distributed system

Transparency	Description
Access	Hides the way in which resources are accessed and the differences in data platform.
Location	Hides where resources are located.
Technology	Hides different technologies such as programming language and OS from user.
Migration / Relocation	Hide resources that may be moved to another location which are in use.
Replication	Hide resources that may be copied at several location.
Concurrency	Hide resources that may be shared with other users.
Failure	Hides failure and recovery of resources from user.
Persistence	Hides whether a resource (software) is in memory or disk.

1.3 Distributed Systems

Distributed Systems Advantages

It has following advantages:

- Resource sharing – Sharing of hardware and software resources.
- Openness – Flexibility of using hardware and software of different vendors.
- Concurrency – Concurrent processing to enhance performance.
- Scalability – Increased throughput by adding new resources.
- Fault tolerance – The ability to continue in operation after a fault has occurred.

Distributed Systems Disadvantages

Its disadvantages are:

- Complexity – They are more complex than centralized systems.
- Security – More susceptible to external attack.
- Manageability – More effort required for system management.
- Unpredictability – Unpredictable responses depending on the system organization and network load.

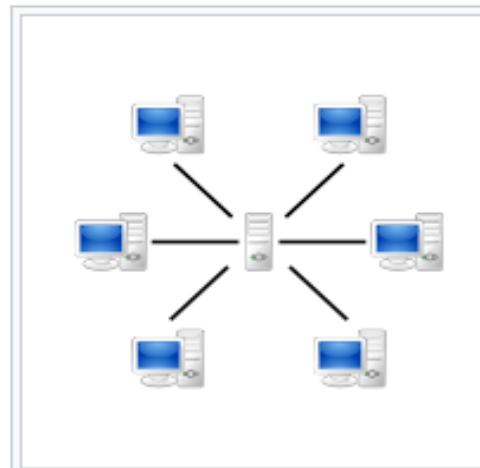
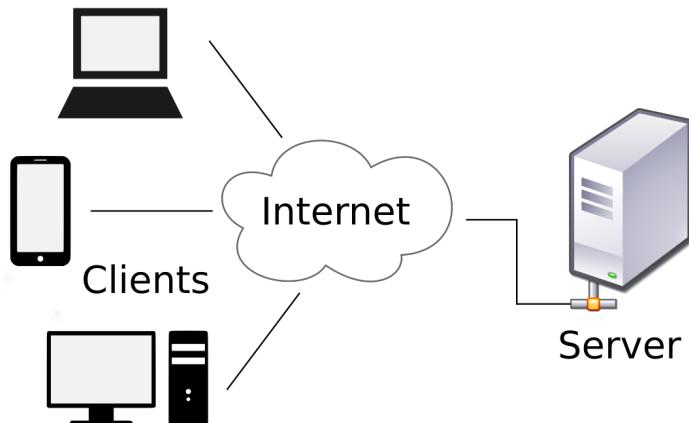
1.3 Distributed Systems

Client-Server Architecture

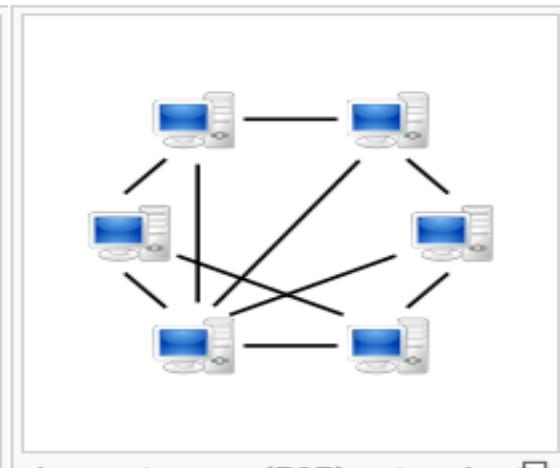
The client-server architecture is the most common distributed system architecture which decomposes the system into two major subsystems or logical processes –

- Client – This is the first process that issues a request to the second process i.e. the server.
- Server – This is the second process that receives the request, carries it out, and sends a reply to the client.

In this architecture, the application is modelled as a set of services those are provided by servers and a set of clients that use these services. The servers need not to know about clients, but the clients must know the identity of servers.



A network based on the **client-server model**, where individual **clients** request services and resources from centralized **servers**.



A **peer-to-peer (P2P) network** in which interconnected nodes ("peers") share resources amongst each other without the use of a centralized administrative system.

1.3 Distributed Systems

Thin-client model

In thin-client model, all the application processing and data management is carried by the server. The client is simply responsible for running the GUI software. It is used when legacy systems are migrated to client server architectures in which legacy system acts as a server in its own right with a graphical interface implemented on a client.

However, A major disadvantage is that it places a heavy processing load on both the server and the network.

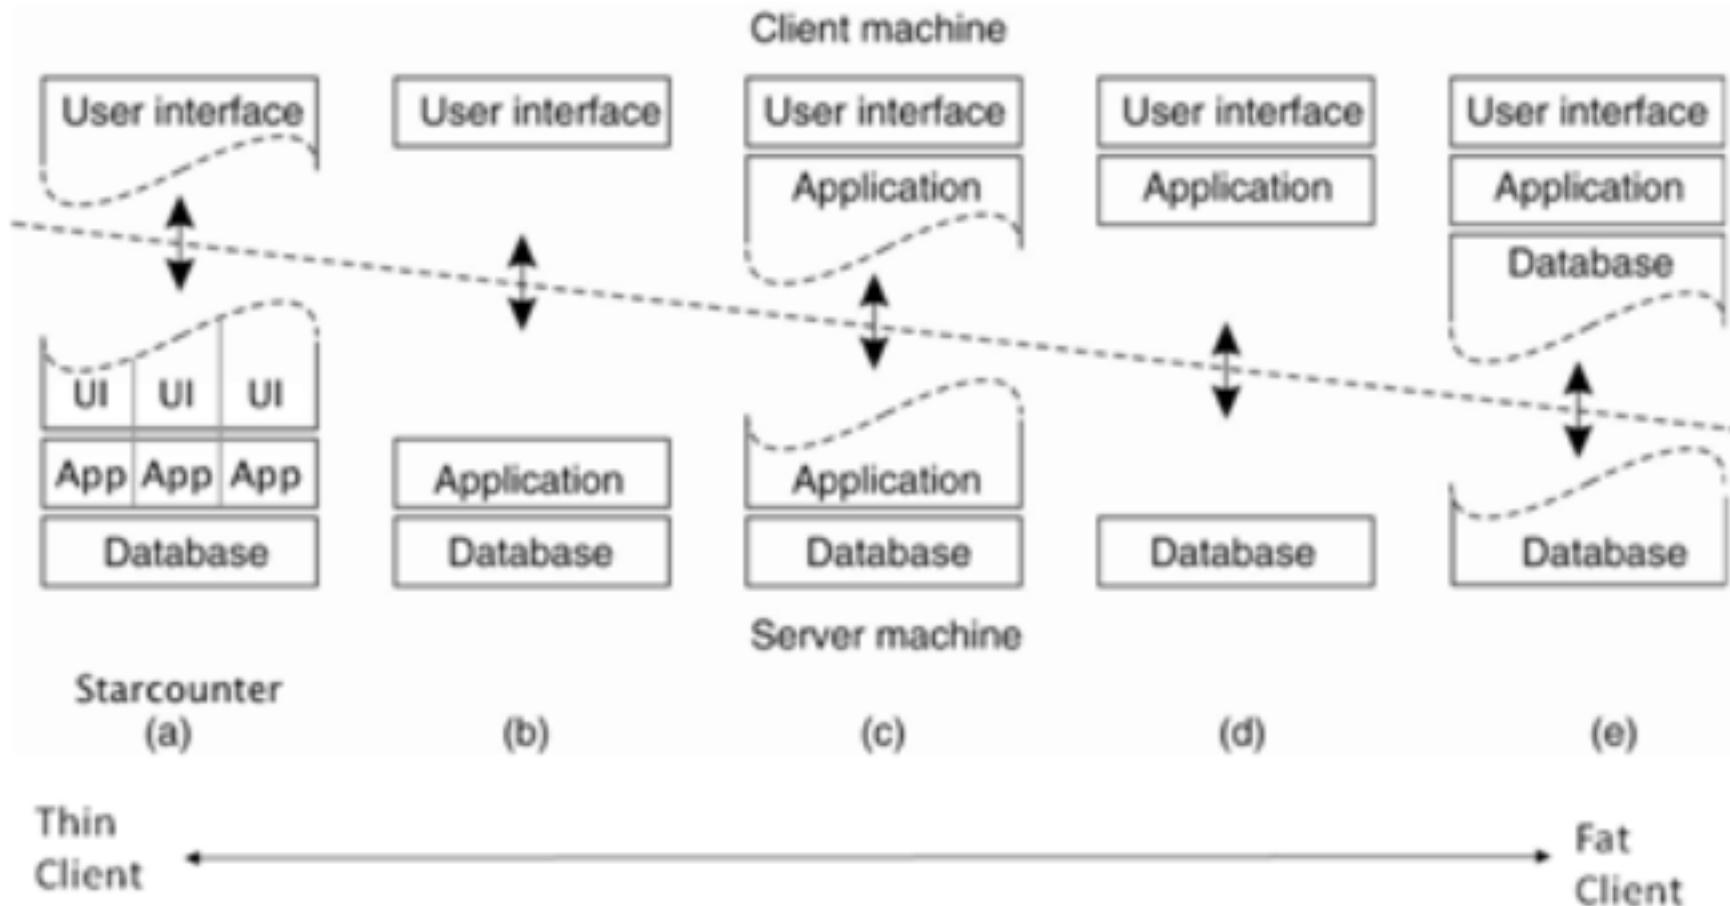
Thick/Fat-client model

In thick-client model, the server is in-charge of the data management. The software on the client implements the application logic and the interactions with the system user. It is most appropriate for new client-server systems where the capabilities of the client system are known in advance.

However, it is more complex than a thin client model especially for management, as all clients should have same copy/version of software application.

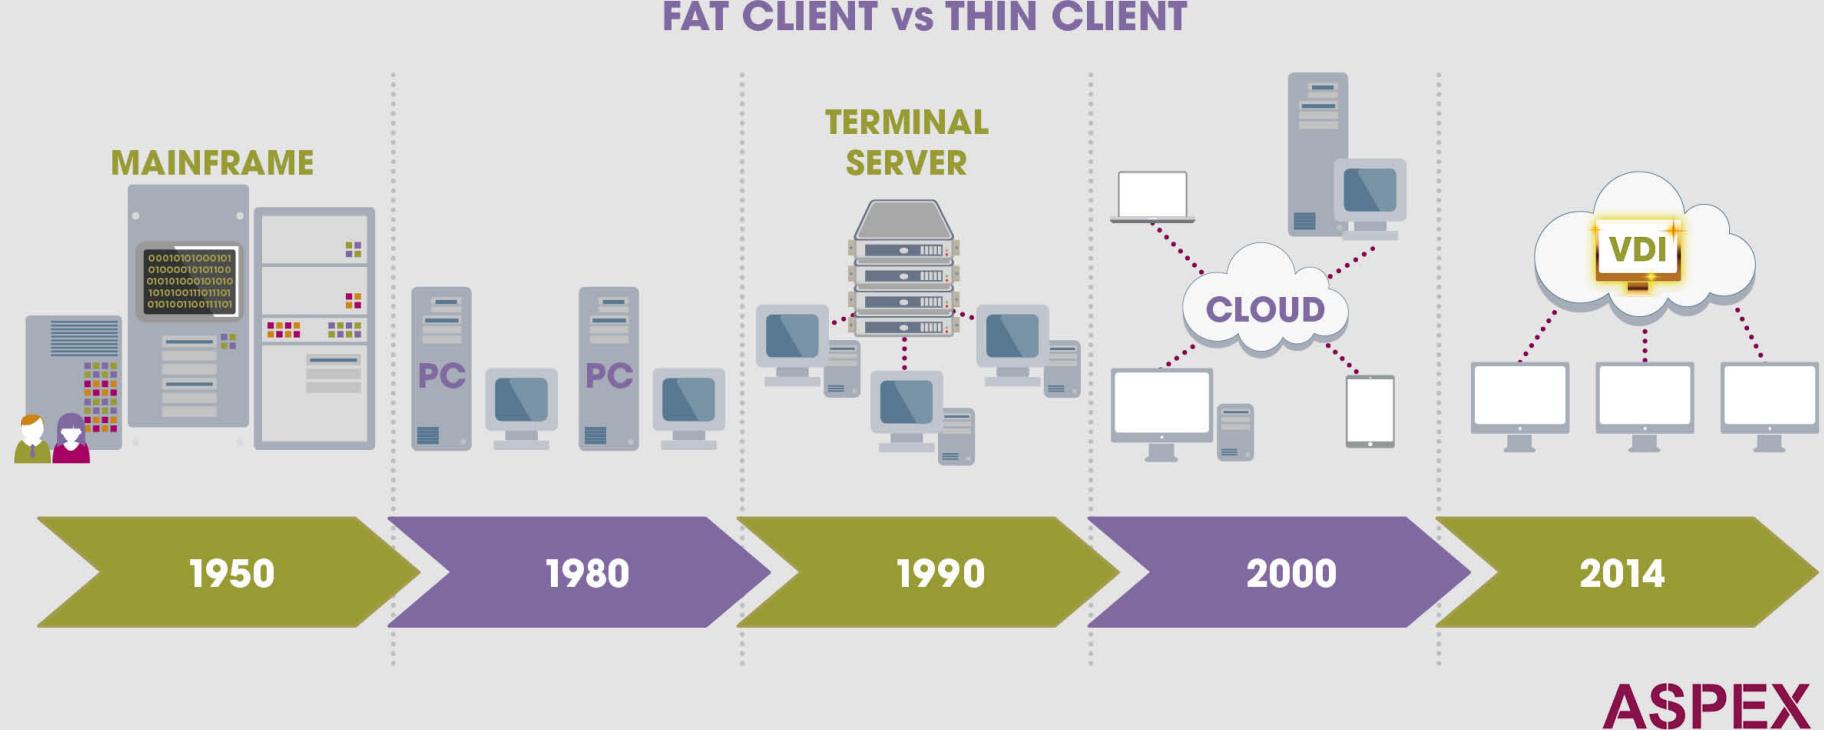
1.3 Distributed Systems

Thin vs. Thick/Fat Client



1.3 Distributed Systems

Thin vs. Thick/Fat Client



1.3 Distributed Systems

Thin/Thick Clients Advantages:

- Separation of responsibilities such as user interface presentation and business logic processing.
- Reusability of server components and potential for concurrency
- Simplifies the design and the development of distributed applications
- It makes it easy to migrate or integrate existing applications into a distributed environment.
- It also makes effective use of resources when a large number of clients are accessing a high-performance server.

Thin/Thick Clients Disadvantages:

- Lack of heterogeneous infrastructure to deal with the requirement changes.
- Security complications.
- Limited server availability and reliability.
- Limited testability and scalability.
- Fat clients with presentation and business logic together.

1.3 Distributed Systems

Multi-Tier Architecture (n-tier Architecture)

Multi-tier architecture is a client–server architecture in which the functions such as presentation, application processing, and data management are physically separated. By separating an application into tiers, developers obtain the option of changing or adding a specific layer, instead of reworking the entire application. It provides a model by which developers can create flexible and reusable applications.

The most general use of multi-tier architecture is the **three-tier architecture**. A three-tier architecture is typically composed of a presentation tier, an application tier, and a data storage tier and may execute on a separate processor.

Presentation Tier

Presentation layer is the topmost level of the application by which users can access directly such as webpage or Operating System GUI (Graphical User interface). The primary function of this layer is to translate the tasks and results to something that user can understand. It communicates with other tiers so that it places the results to the browser/client tier and all other tiers in the network.

Application Tier (Business Logic, Logic Tier, or Middle Tier)

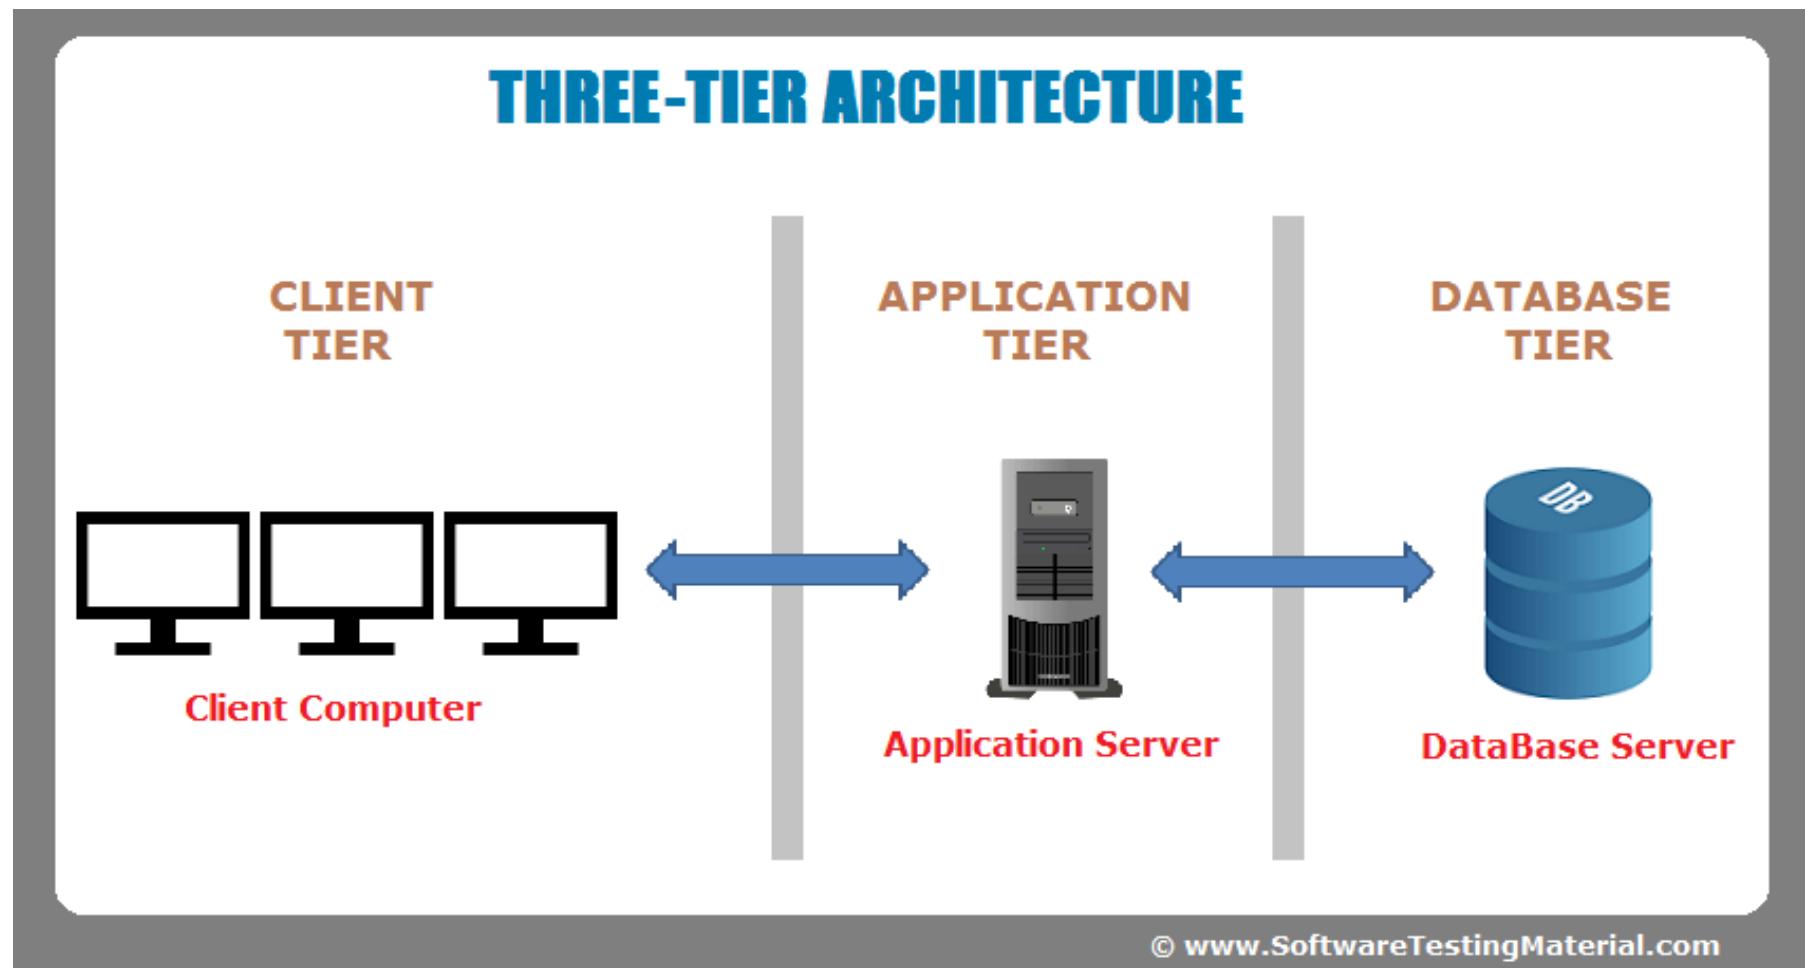
Application tier coordinates the application, processes the commands, makes logical decisions, evaluation, and performs calculations. It controls an application's functionality by performing detailed processing. It also moves and processes data between the two surrounding layers.

Data Tier

In this layer, information is stored and retrieved from the database or file system. The information is then passed back for processing and then back to the user. It includes the data persistence mechanisms (database servers, file shares, etc.) and provides API (Application Programming Interface) to the application tier which provides methods of managing the stored data.

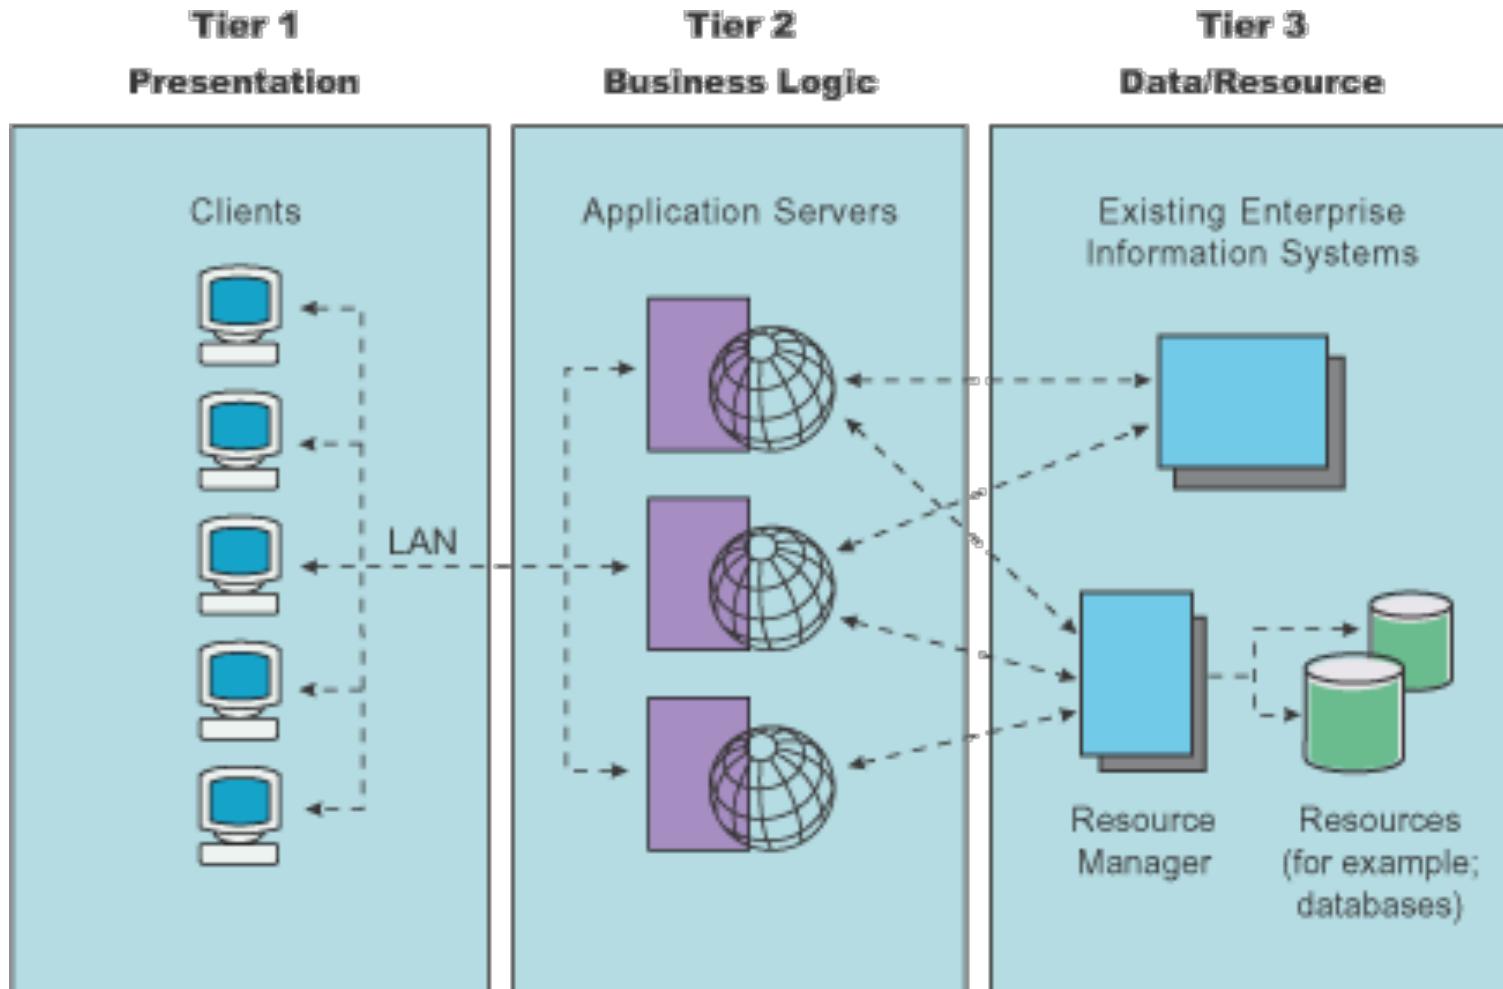
1.3 Distributed Systems

3-Tier Architecture



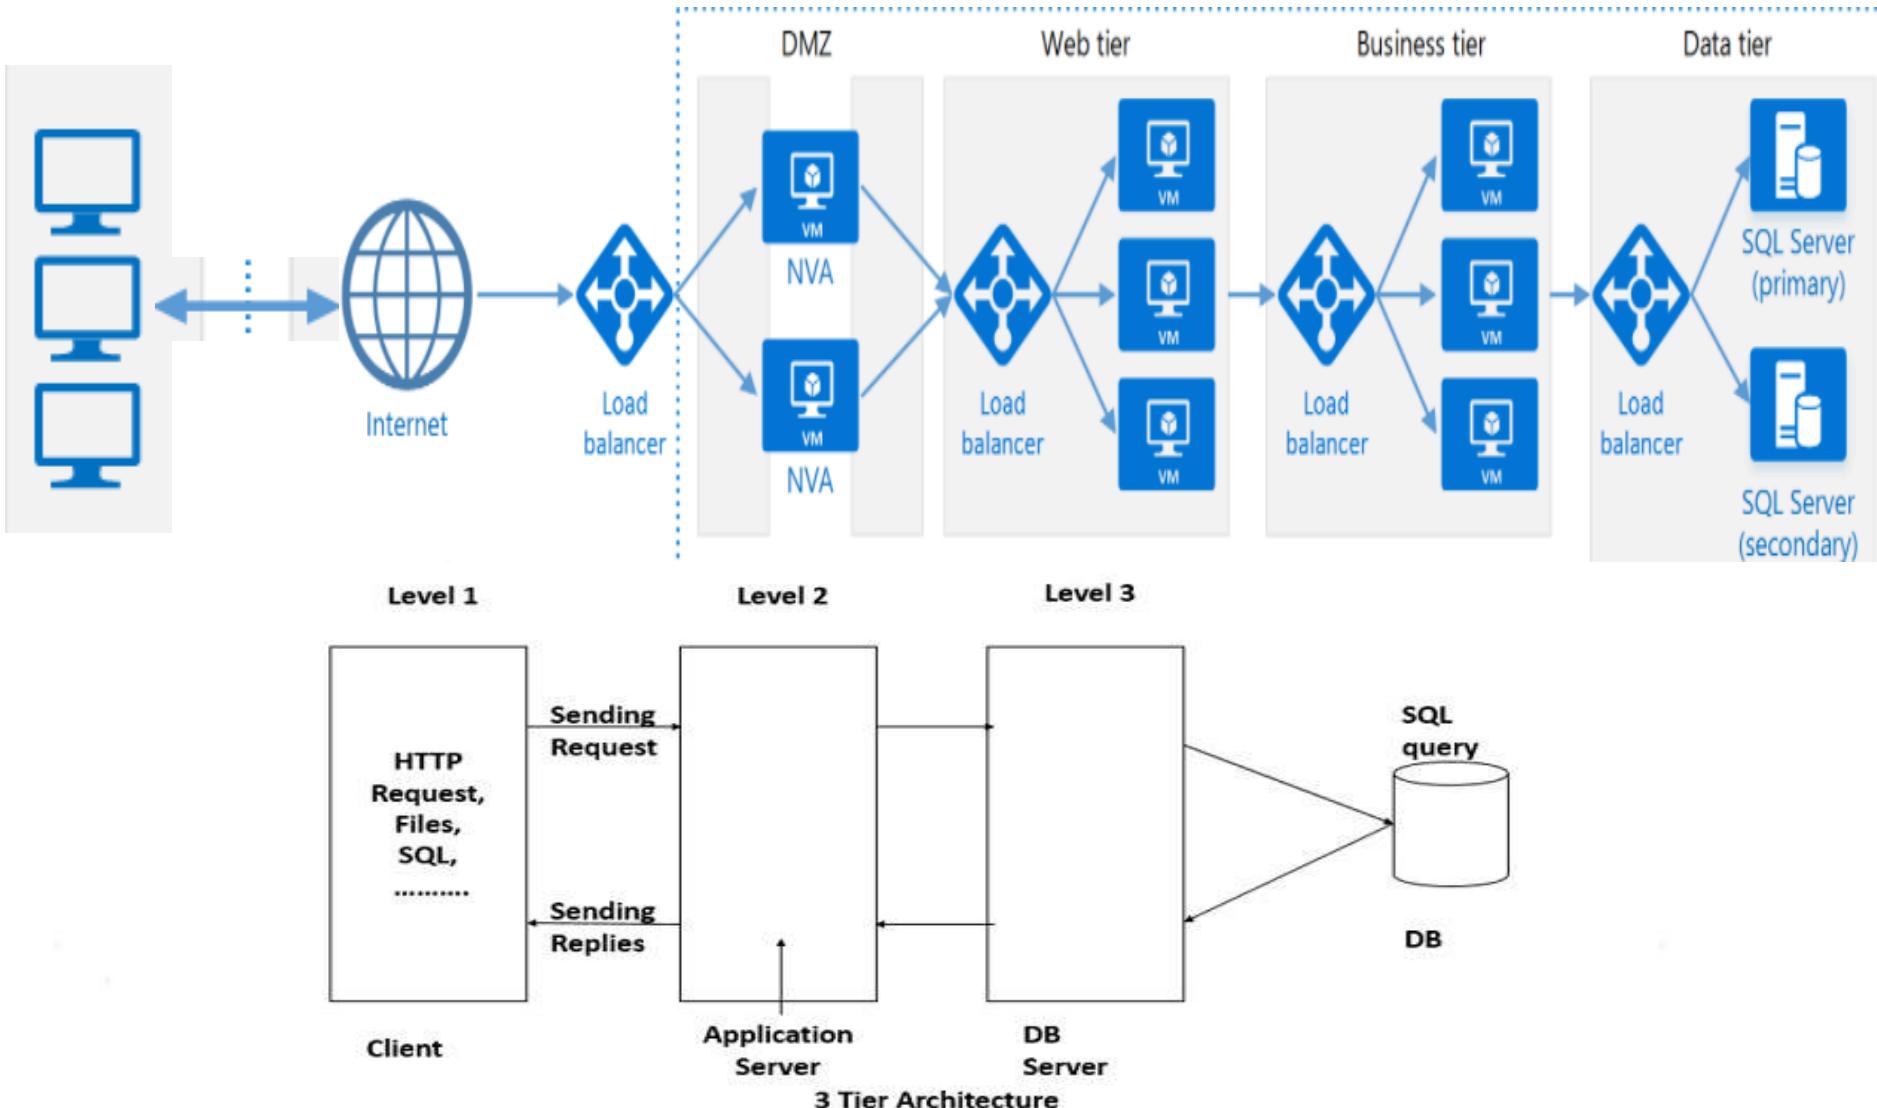
1.3 Distributed Systems

3-Tier Architecture



1.3 Distributed Systems

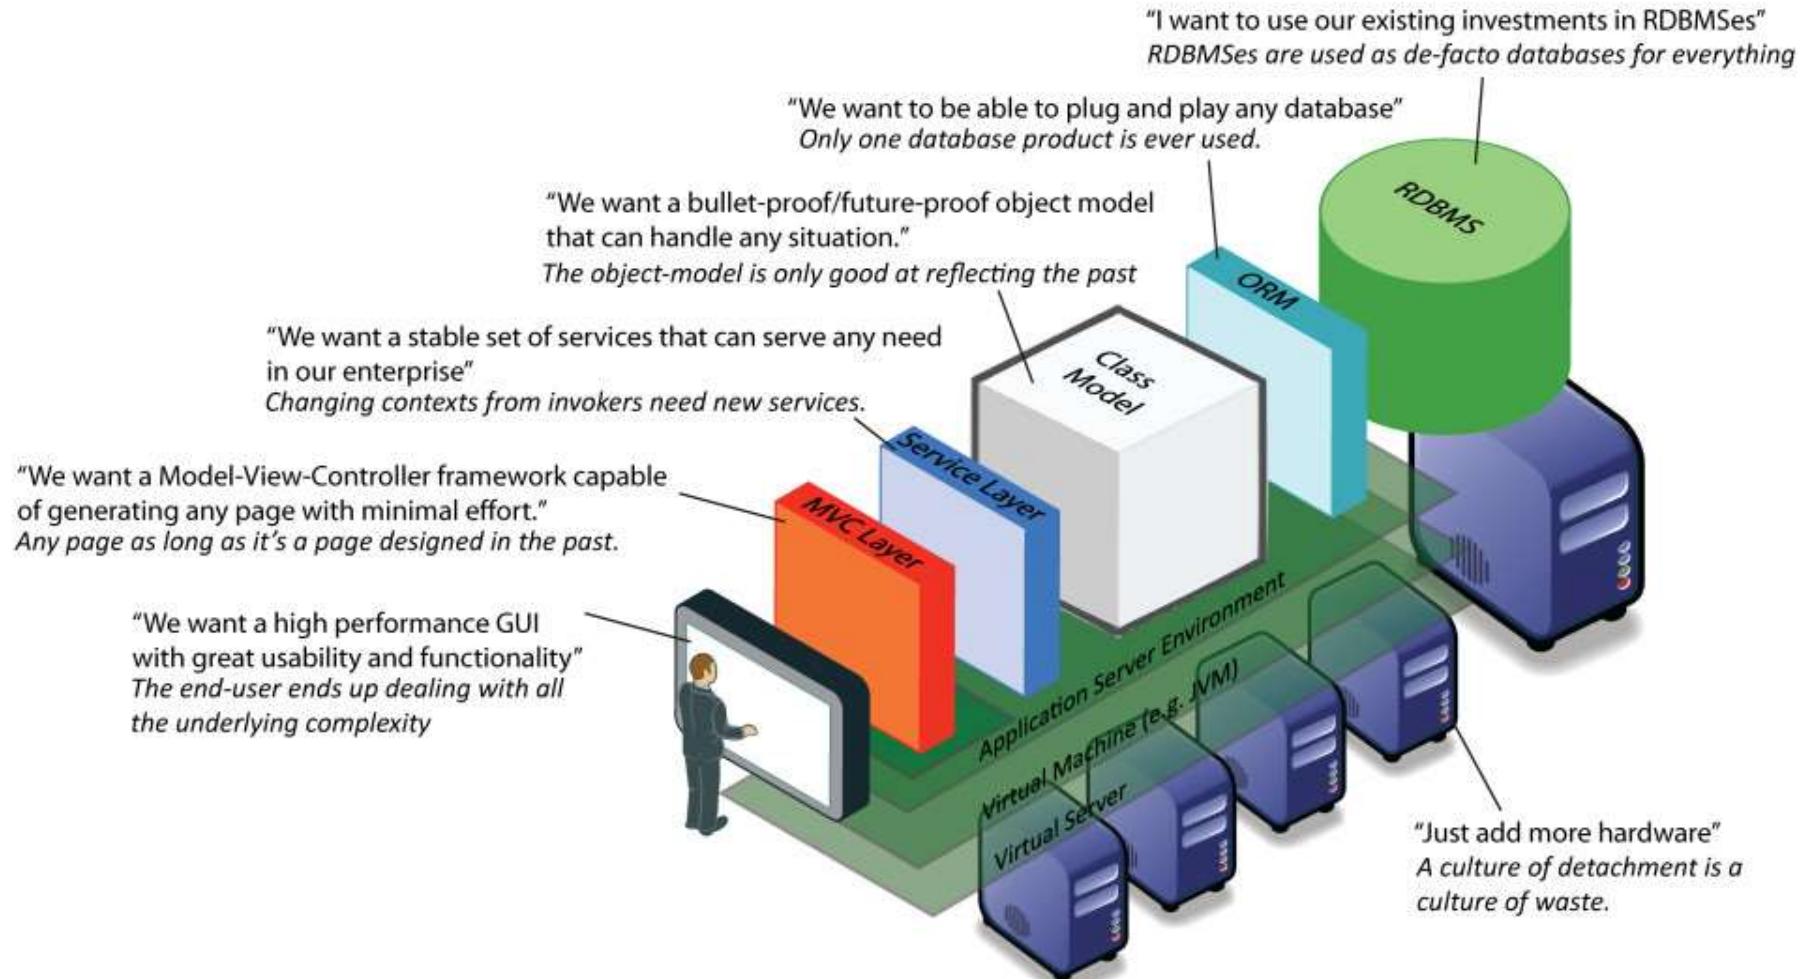
3-Tier/4-Tier Architecture



1.3 Distributed Systems

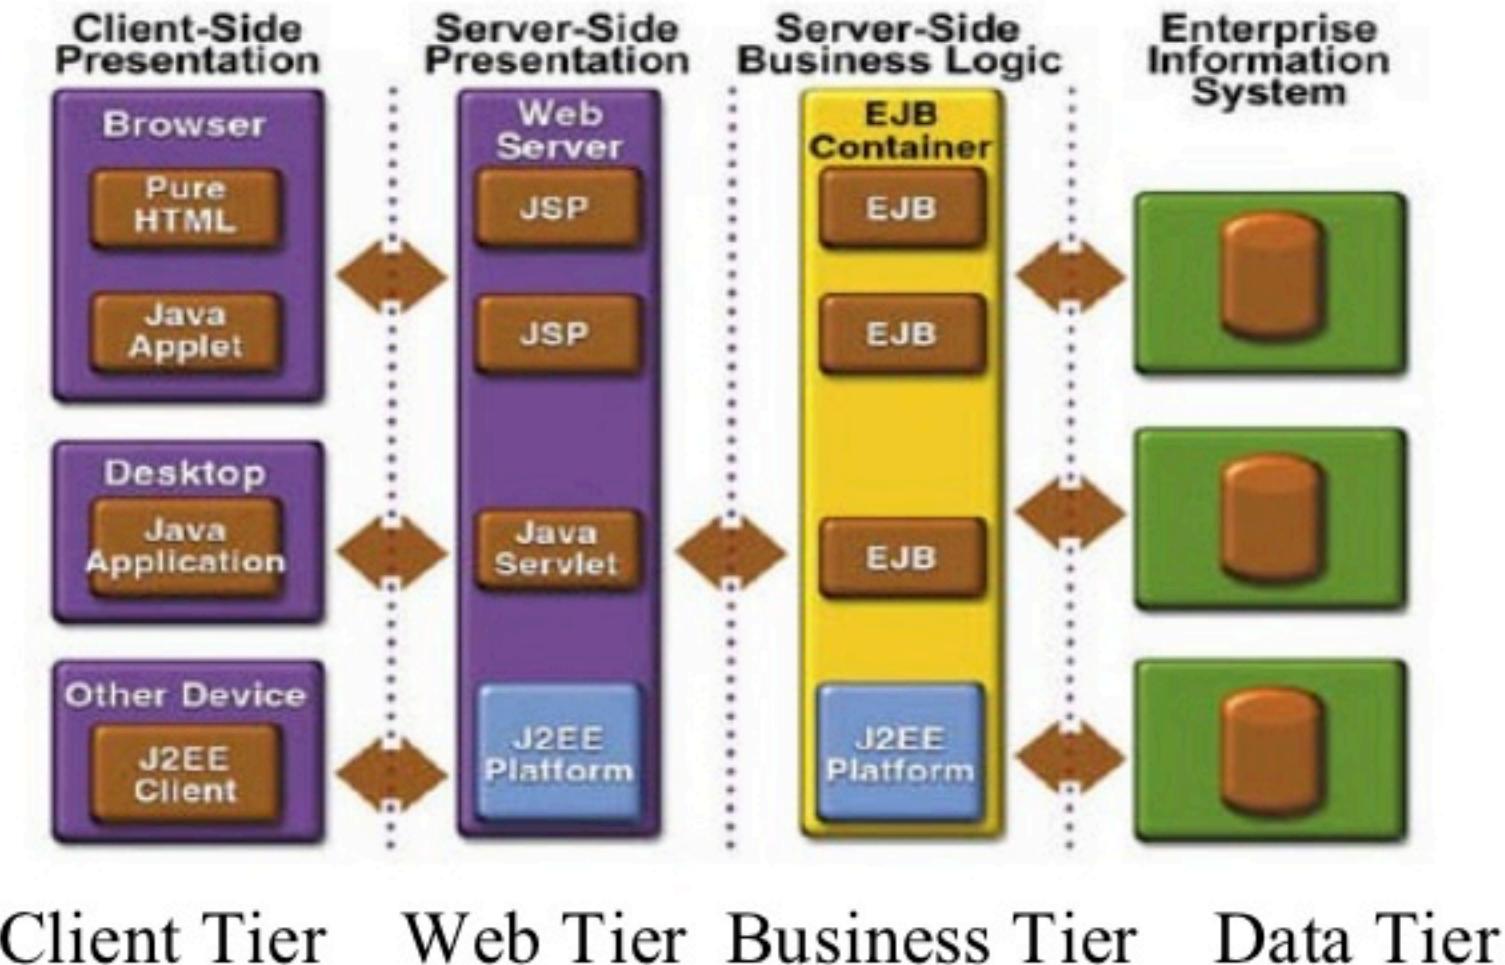
n-Tier Architecture

n-Tier Architecture :: Each layer can add value, just make sure it's really needed.



1.3 Distributed Systems

n-Tier Architecture



1.3 Distributed Systems

n-Tier Architecture

Advantages

- Better performance than a thin-client approach and is simpler to manage than a thick-client approach.
- Enhances the reusability and scalability – as demands increase, extra servers can be added.
- Provides multi-threading support and also reduces network traffic.
- Provides maintainability and flexibility.

Disadvantages

- Unsatisfactory Testability due to lack of testing tools.
- More critical server reliability and availability.

1.3 Distributed Systems

Broker Architectural Style

Broker Architectural Style is a middleware architecture used in distributed computing to coordinate and enable the communication between registered servers and clients. Here, object communication takes place through a middleware system called an object request broker (software bus).

However, client and the server do not interact with each other directly. Client and server have a direct connection to its proxy, which communicates with the mediator-broker. A server provides services by registering and publishing their interfaces with the broker and clients can request the services from the broker statically or dynamically by look-up.

CORBA (Common Object Request Broker Architecture) is a good implementation example of the broker architecture.

1.3 Distributed Systems

Broker Architectural Components

Components of Broker Architectural Style - The components of broker architectural style are discussed through following heads:

Broker

Broker is responsible for brokering the service requests, locating a proper server, transmitting requests, and sending responses back to clients. It also retains the servers' registration information including their functionality and services as well as location information.

Further, it provides APIs for clients to request, servers to respond, registering or unregistering server components, transferring messages, and locating servers.

Stub

Stubs are generated at the static compilation time and then deployed to the client side which is used as a proxy for the client. Client-side proxy acts as a mediator between the client and the broker and provides additional transparency between them and the client; a remote object appears like a local one.

The proxy hides the IPC (inter-process communication) at protocol level and performs marshaling of parameter values and un-marshaling of results from the server.

Skeleton

Skeleton is generated by the service interface compilation and then deployed to the server side, which is used as a proxy for the server. Server-side proxy encapsulates low-level system-specific networking functions and provides high-level APIs to mediate between the server and the broker.

Further, it receives the requests, unpacks the requests, unmarshals the method arguments, calls the suitable service, and also marshals the result before sending it back to the client.

Bridge

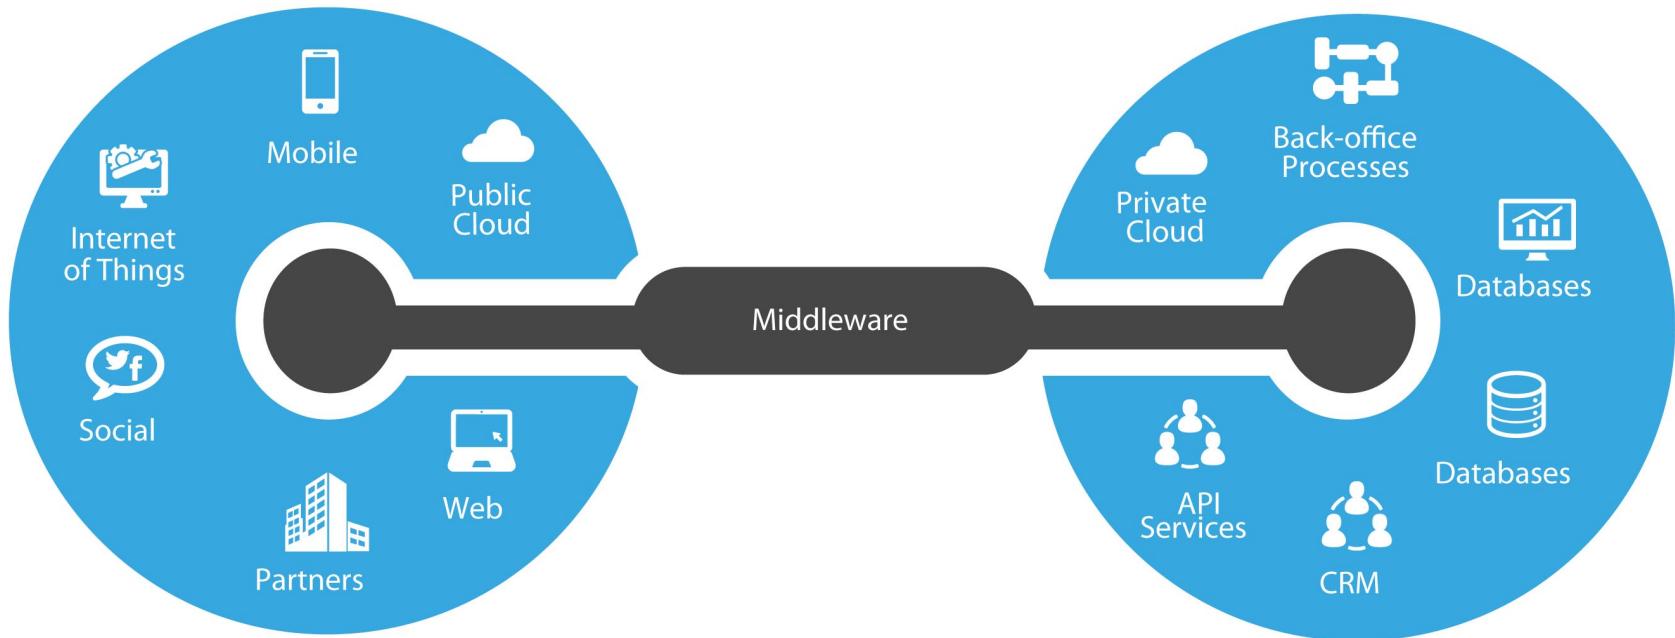
A bridge can connect two different networks based on different communication protocols. It mediates different brokers including DCOM, .NET remote, and Java CORBA brokers.

Bridges are optional component, which hides the implementation details when two brokers interoperate and take requests and parameters in one format and translate them to another format.

1.3 Distributed Systems

Middleware API

Middleware is computer software that provides services to software applications beyond those available from the operating system. It can be described as "software glue".



Systems of Engagement

Systems of Record

1.3 Distributed Systems

Middleware

The term is most commonly used for software that enables communication and management of data in [distributed applications](#). An [IETF](#) workshop in 2000 defined middleware as "those services found above the [transport](#) (i.e. over TCP/IP) layer set of services but below the application environment" (i.e. below application-level [APIs](#)).^[3] In this more specific sense *middleware* can be described as the dash ("") in [client-server](#), or the -to- in [peer-to-peer](#).^[citation needed] Middleware includes [web servers](#), [application servers](#), [content management systems](#), and similar tools that support application development and delivery.

ObjectWeb defines middleware as: "The software layer that lies between the [operating system](#) and applications on each side of a distributed computing system in a network."^[4]

Services that can be regarded as middleware include: [enterprise application integration](#), [data integration](#),

[Remote Procedure Call \(RPC\)](#)

[Message Oriented Middleware \(MOM\)](#),

[Object Request Brokers \(ORBs\)](#),

[Service Oriented Architecture \(SOA\)](#),

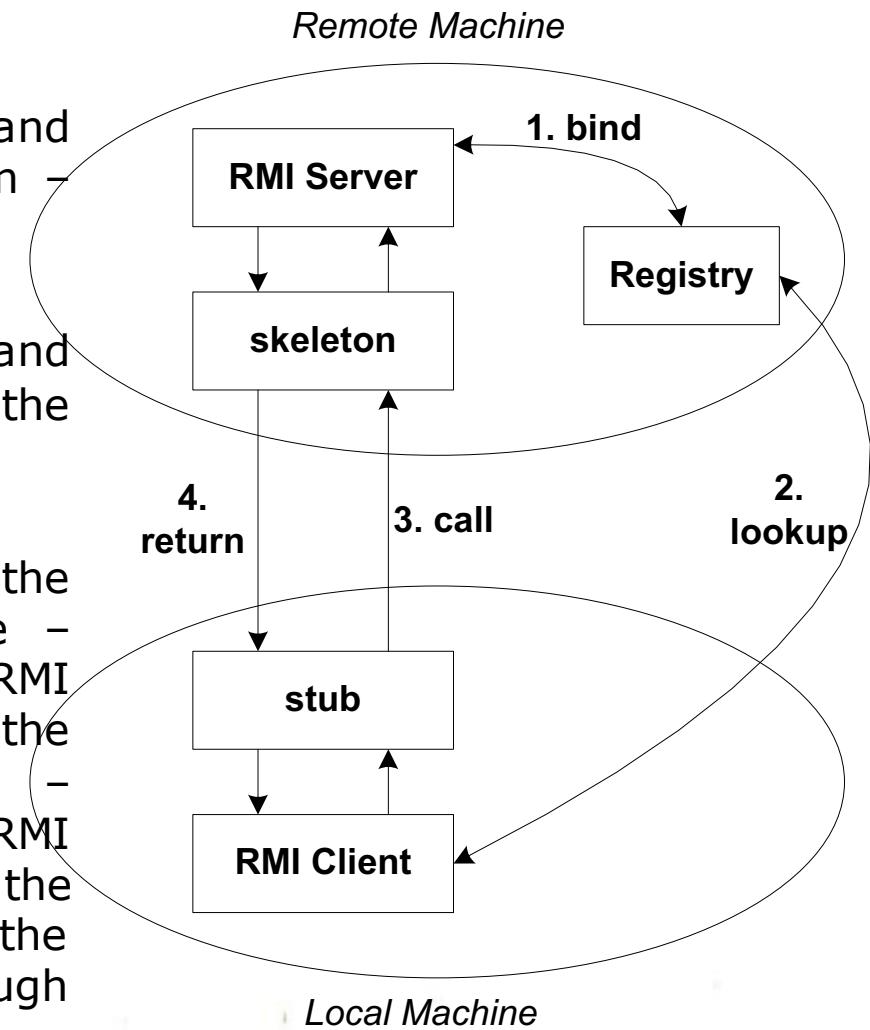
and the [Enterprise Service Bus \(ESB\)](#).

[Database](#) access services are often characterized as middleware. Some of them are language specific implementations and support heterogeneous features and other related communication features. Examples of database-oriented middleware include [ODBC](#), [JDBC](#) and [transaction processing](#) monitors.

1.3 Distributed Systems – Middleware: RPC / Java RMI Architecture

RPC/RMI Architecture

- RMI Server must register its name and address in the RMI Registry program – **bind**
- RMI Client is looking for the address and the name of the RMI server object in the RMI Registry program – **lookup**
- RMI Stub serializes and transmits the parameters of the call in sequence – **“marshalling”** to the RMI Skeleton. RMI Skeleton de-serializes and extracts the parameters from the received call – **“unmarshalling”**. In the end, the RMI Skeleton calls the method inside the server object and send back the response to the RMI Stub through “marshalling” the return’s parameter.

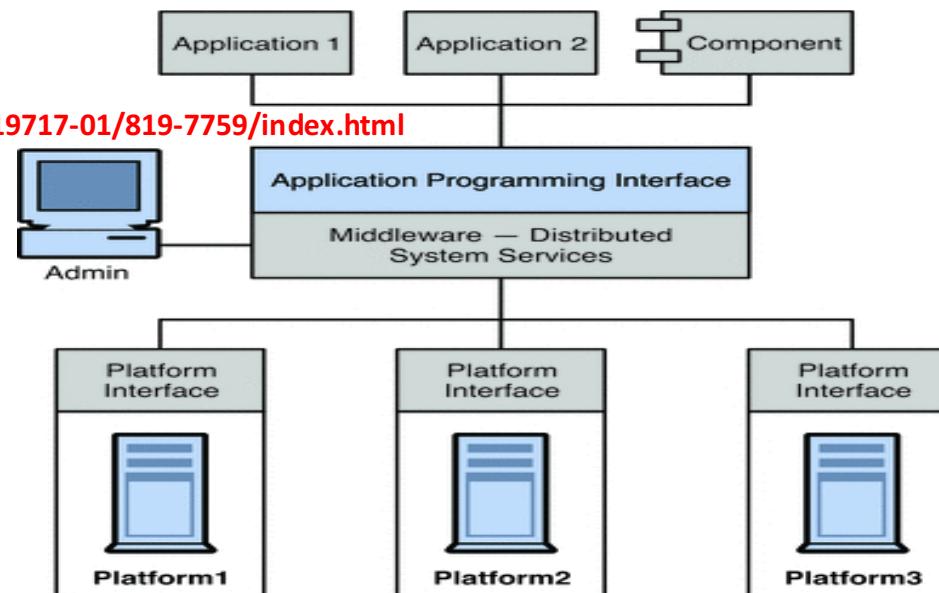


1.3 Distributed Systems: MOM / Middleware & JMS Technology Overview

Middleware Technologies:

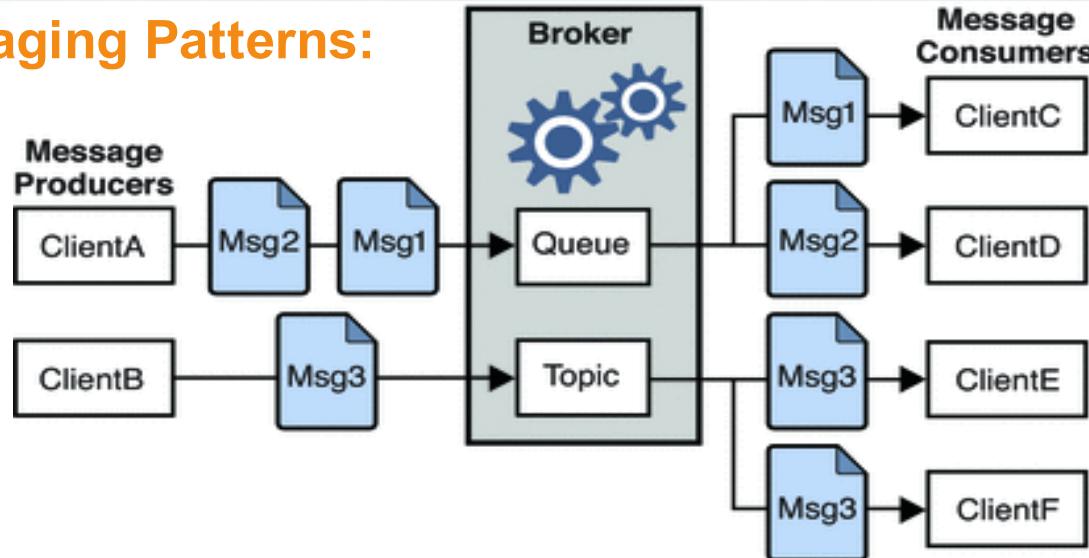
Middleware can be grouped into the following categories:

- **Remote Procedure Call or RPC**-based middleware, which allows procedures in one application to call procedures in remote applications as if they were local calls. The middleware implements a linking mechanism that locates remote procedures and makes these transparently available to a caller. Traditionally, this type of middleware handled procedure-based programs; it now also includes object-based components.
- **Message Oriented Middleware or MOM**-based middleware, which allows distributed applications to communicate and exchange data by sending and receiving messages asynchronously.
- **Object Request Broker or ORB**-based middleware, which enables an application's objects to be distributed and shared across heterogeneous networks.



1.3 Distributed Systems: MOM / Middleware & JMS Technology Overview

JMS Messaging Patterns:



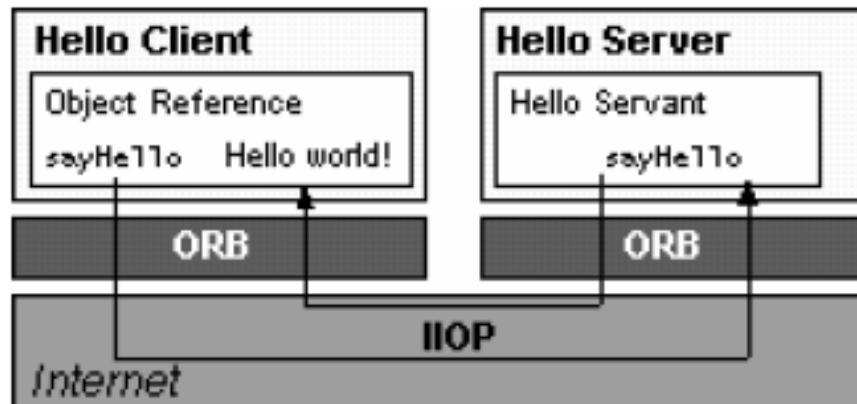
Clients A and B are message producers, sending messages to clients C, D, and E by way of two different kinds of destinations.

- Messaging between clients A, C, and D illustrates the point-to-point pattern. Using this pattern, a client sends a message to a queue destination from which only one receiver may get it. No other receiver accessing that destination can get that message.
- Messaging between clients B, E, and F illustrates the publish/subscribe pattern. Using this broadcast pattern, a client sends a message to a topic destination from which any number of consuming subscribers can retrieve it. Each subscriber gets its own copy of the message.

Message consumers in either domain can choose to get messages synchronously or asynchronously. Synchronous consumers make an explicit call to retrieve a message; asynchronous consumers specify a callback method that is invoked to pass a pending message. Consumers can also filter out messages by specifying selection criteria for incoming messages.

1.3 Distributed Systems: CORBA

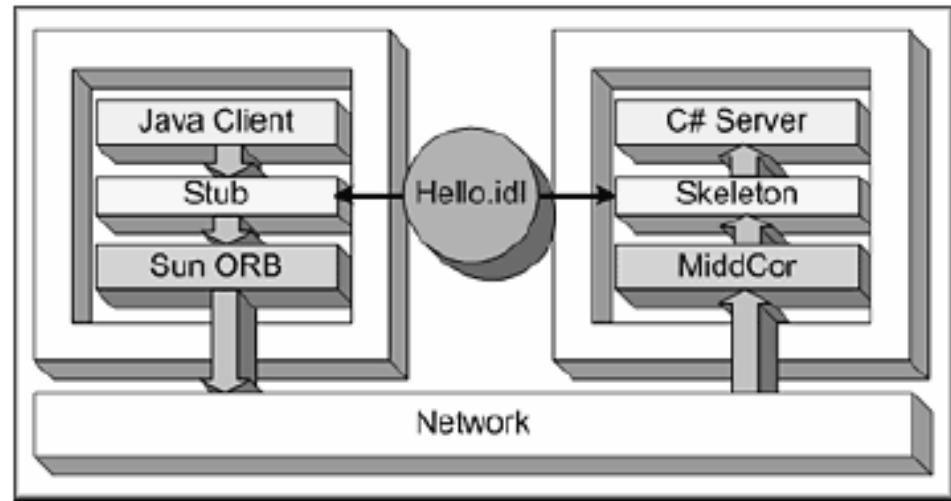
CORBA Architecture



From: <http://java.sun.com/docs/books/tutorial/idl/hello/index.html>

Hello.idl

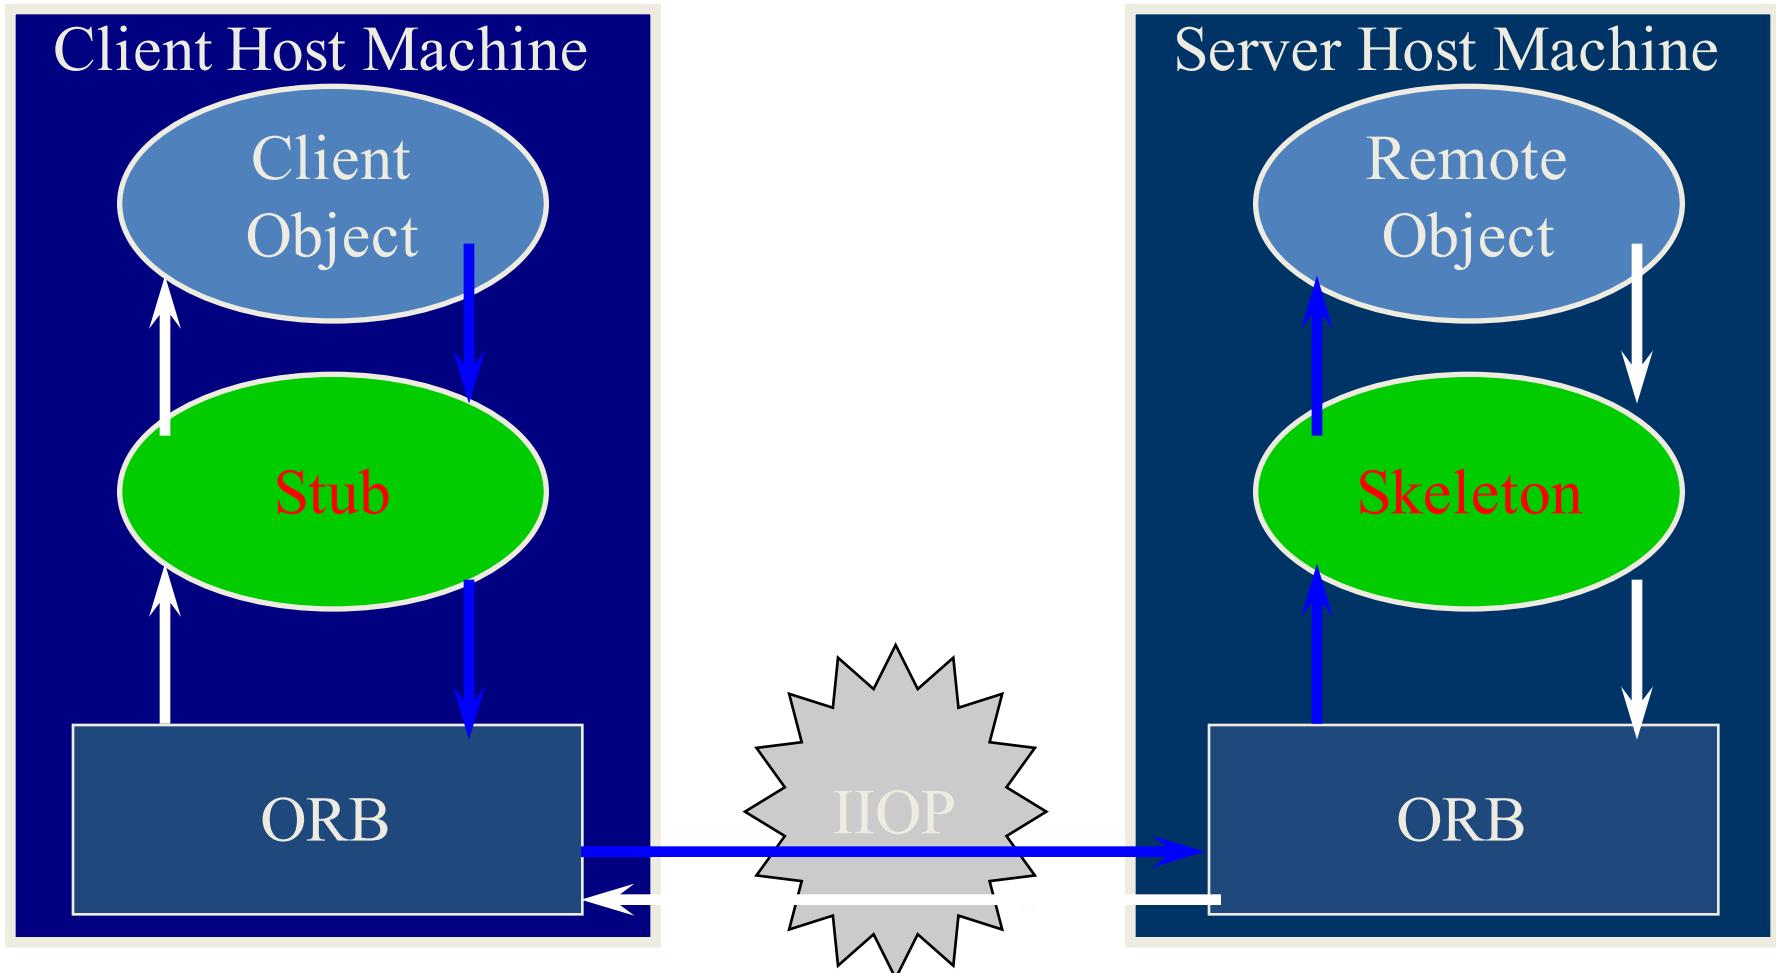
```
interface Hello {  
    string sayHello();  
};
```



From: http://www.molecular.com/news/recent_articles/051904_javaworld.aspx

1.3 Distributed Systems: CORBA

CORBA Architecture

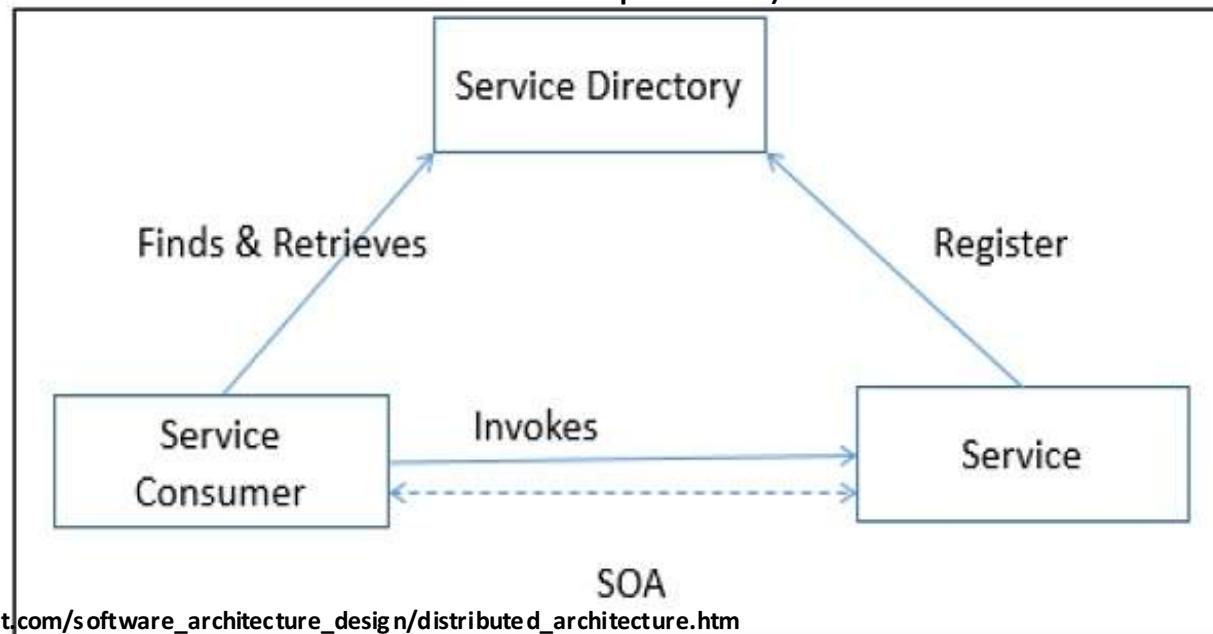


1.3 Distributed Systems: SOA

Service Oriented Architecture – SOA

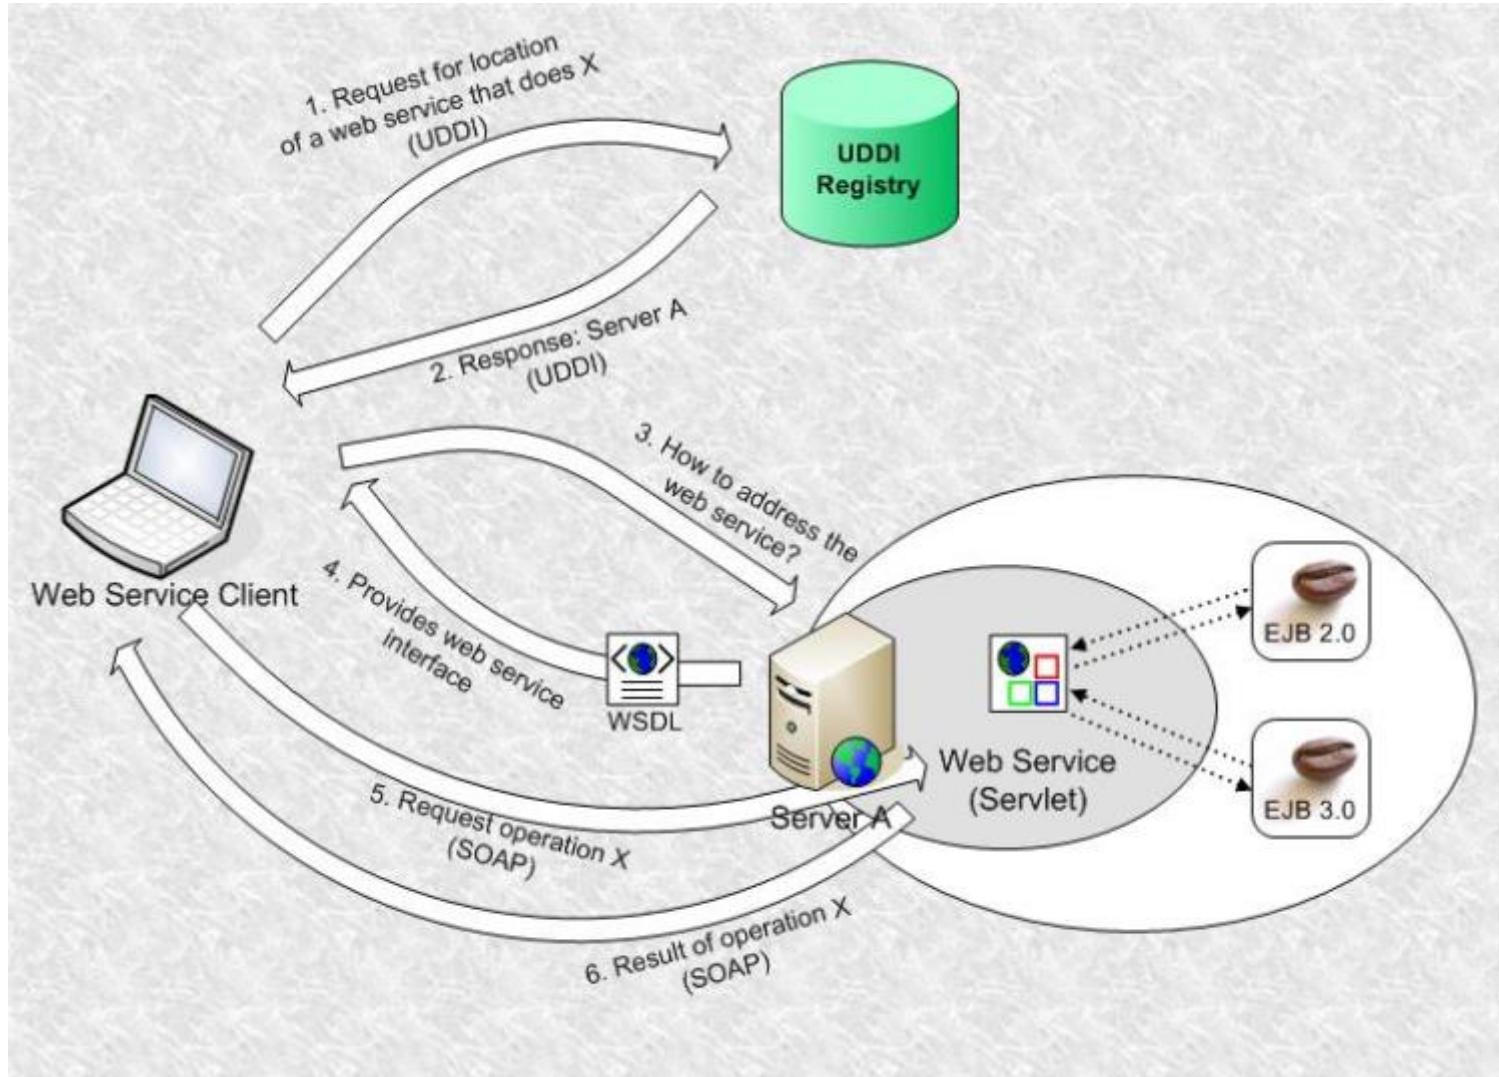
A service is a component of business functionality that is well-defined, self-contained, independent, published, and available to be used via a standard programming interface. The connections between services are conducted by common and universal message-oriented protocols such as the SOAP Web service protocol, which can deliver requests and responses between services loosely.

Service-oriented architecture is a client/server design which support business-driven IT approach in which an application consists of software services and software service consumers (also known as clients or service requesters).



1.3 Distributed Systems: SOA

Service Oriented Architecture – SOA



1.3 Distributed Systems: SOA

Service Oriented Architecture – SOA

Features of SOA

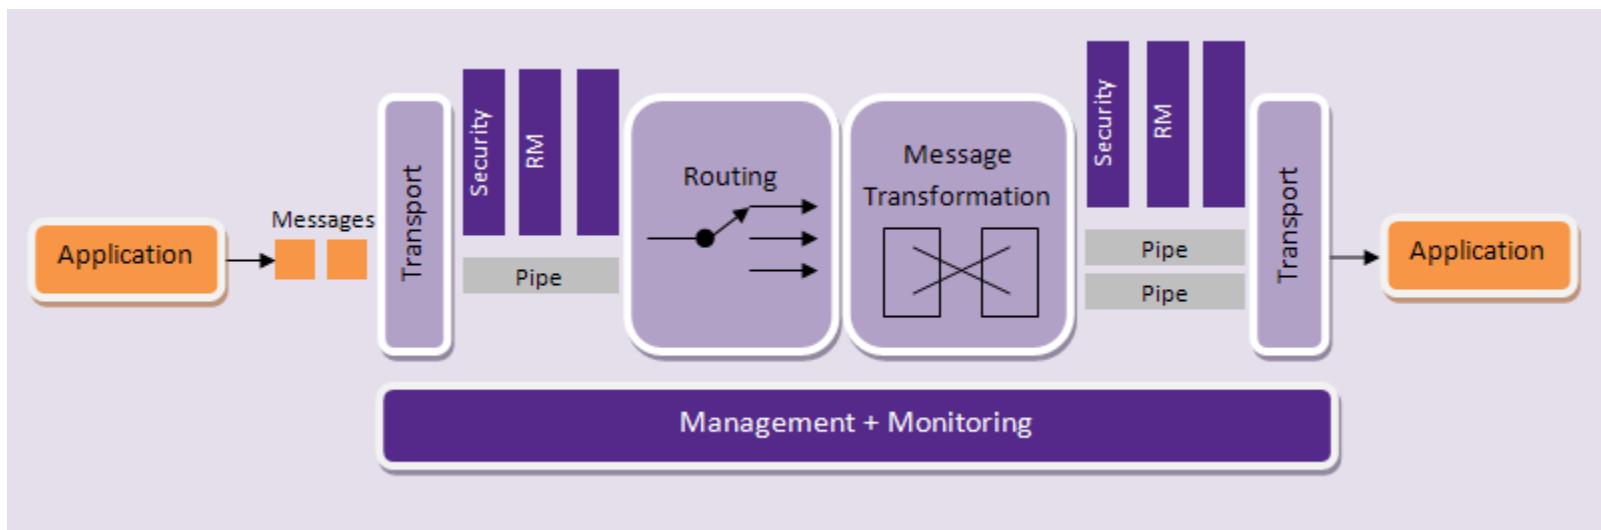
A service-oriented architecture provides the following features –

- Distributed Deployment – Expose enterprise data and business logic as loosely, coupled, discoverable, structured, standard-based, coarse-grained, stateless units of functionality called services.
- Composability – Assemble new processes from existing services that are exposed at a desired granularity through well defined, published, and standard complaint interfaces.
- Interoperability – Share capabilities and reuse shared services across a network irrespective of underlying protocols or implementation technology.
- Reusability – Choose a service provider and access to existing resources exposed as services.

1.3 Distributed Systems: ESB

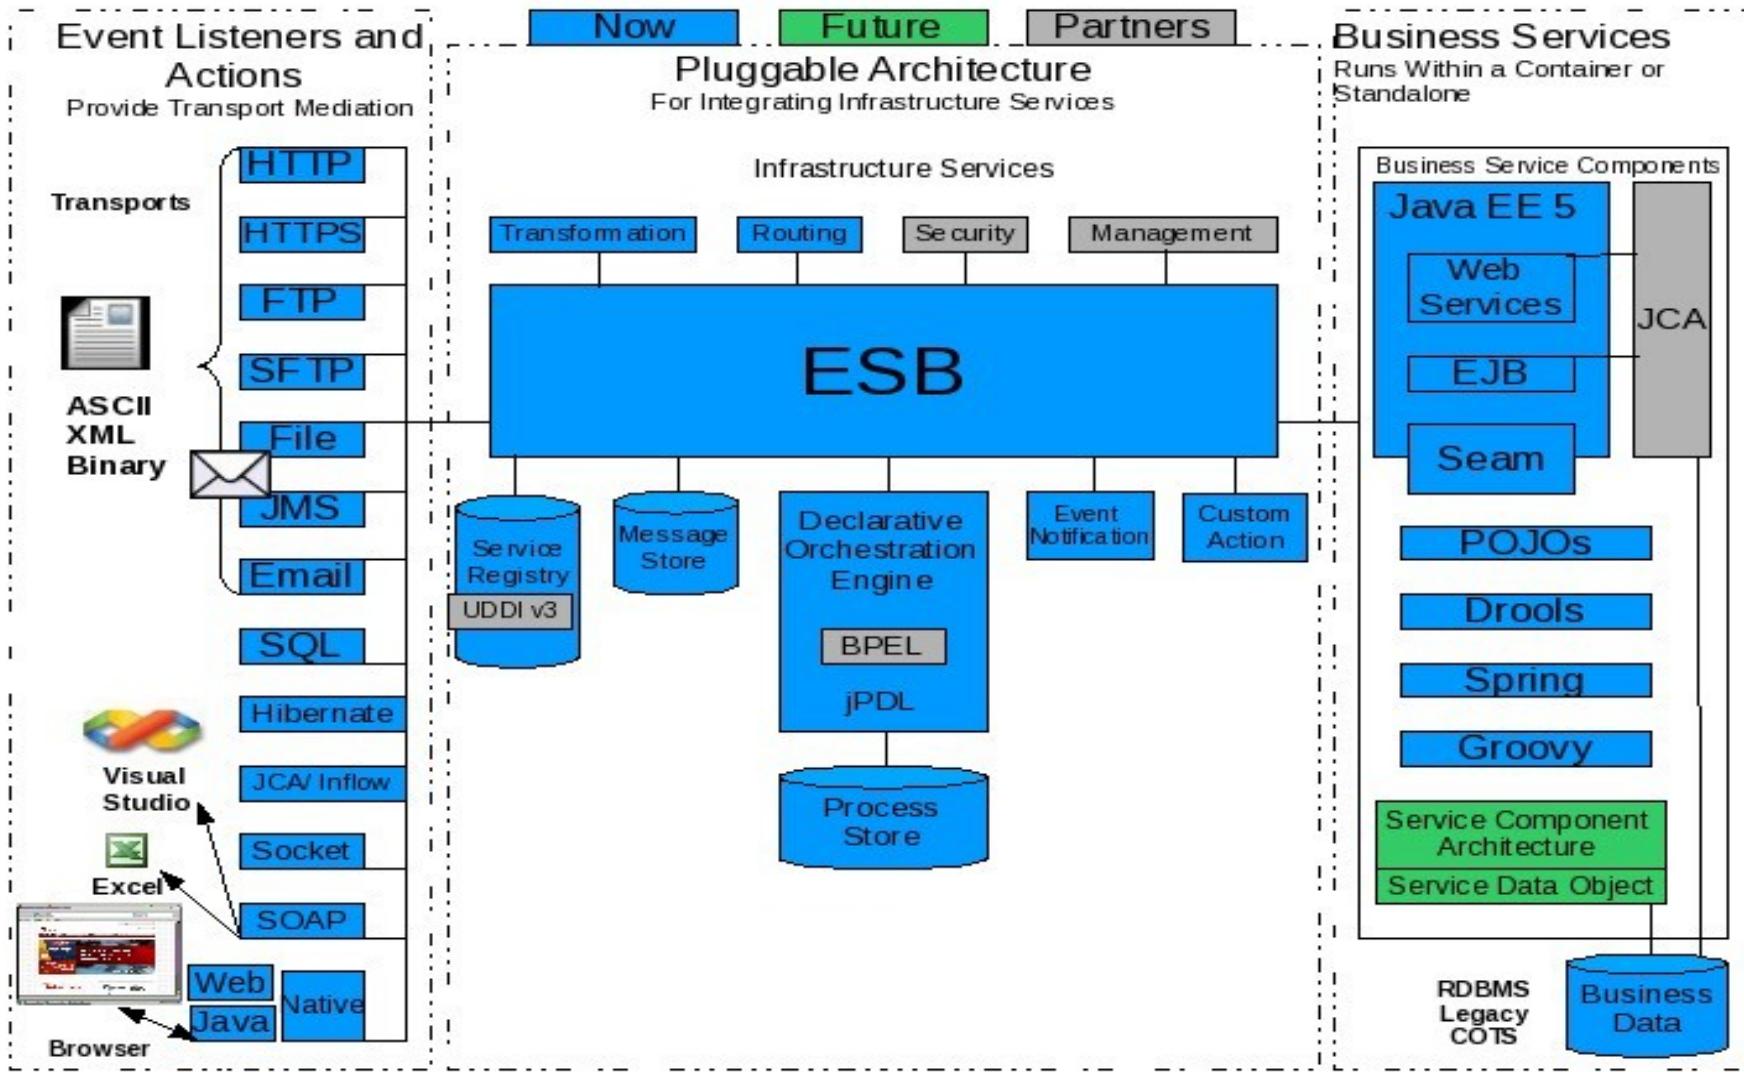
Enterprise Service Bus Architecture – ESB

An **enterprise service bus (ESB)** implements a communication system between mutually interacting software applications in a service-oriented architecture (SOA). As it implements a distributed computing architecture, it implements a special variant of the more general client-server model, wherein, in general, any application using ESB can behave as server or client in turns. ESB promotes agility and flexibility with regard to high-level protocol communication between applications. The primary goal of the high-level protocol communication is enterprise application integration (EAI) of heterogeneous and complex service or application landscapes (a view from the network level).



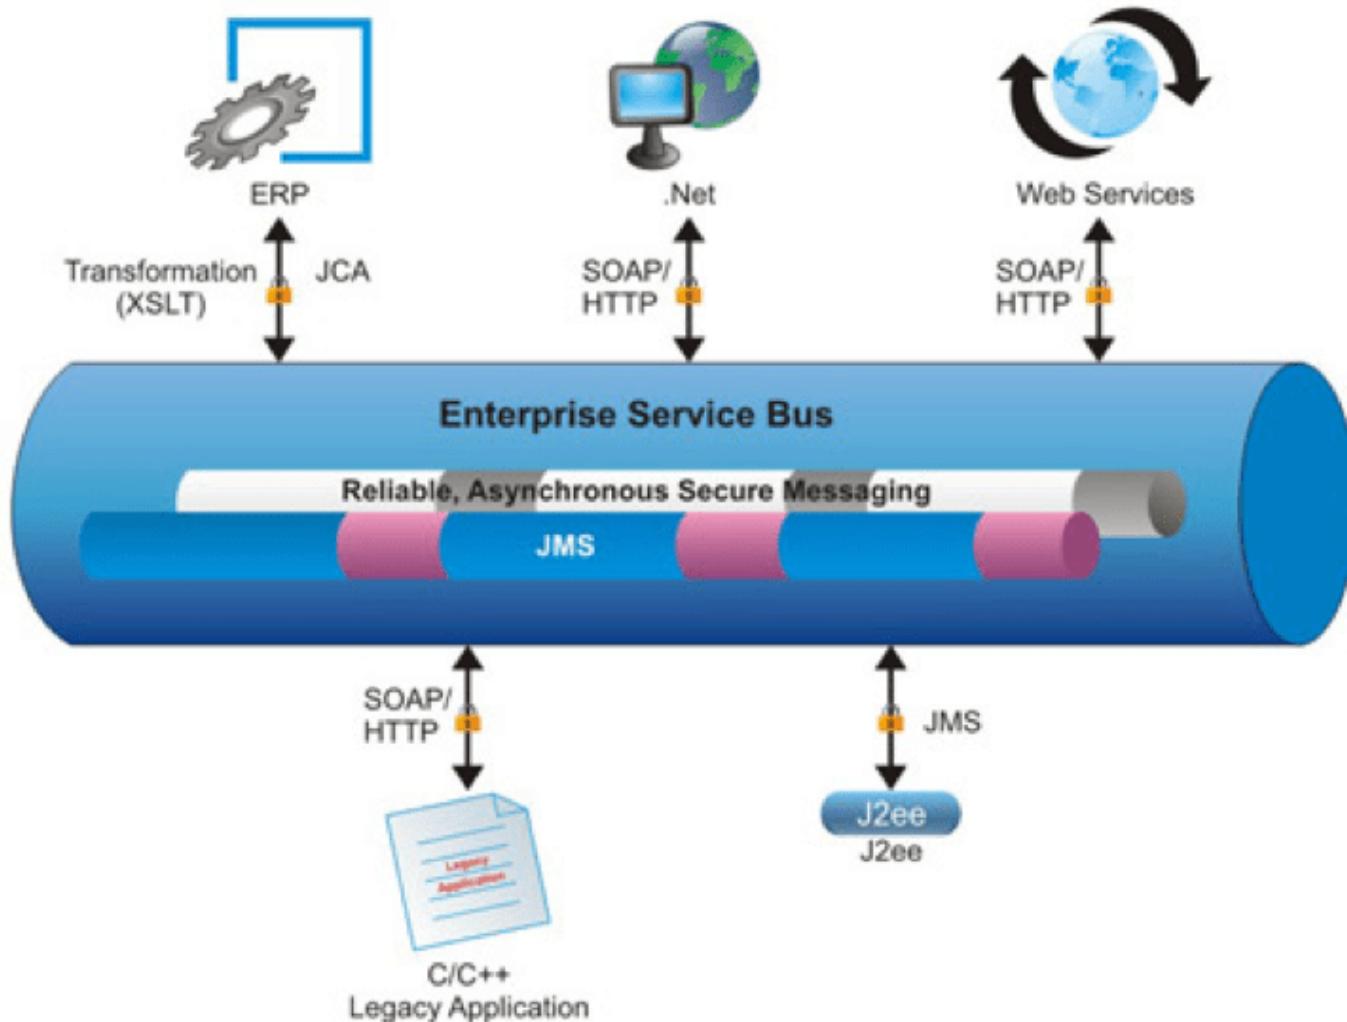
1.3 Distributed Systems: ESB

Enterprise Service Bus Architecture – ESB



1.3 Distributed Systems: ESB

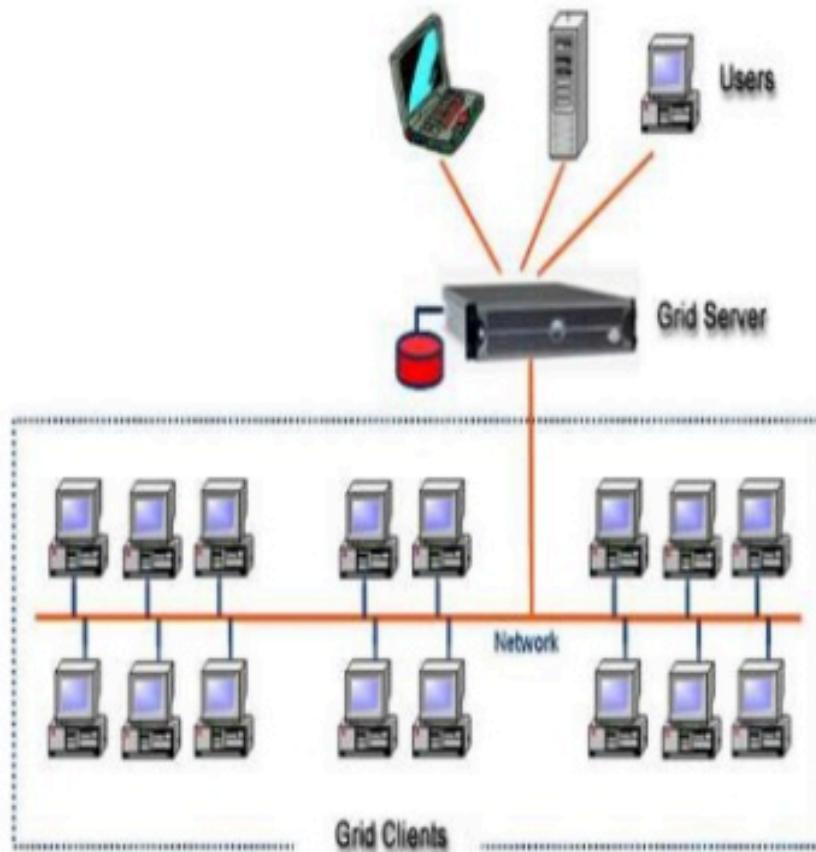
Enterprise Service Bus Architecture – ESB



1.3 Distributed Systems: GRID

Grid Computing

How Grid computing works ?

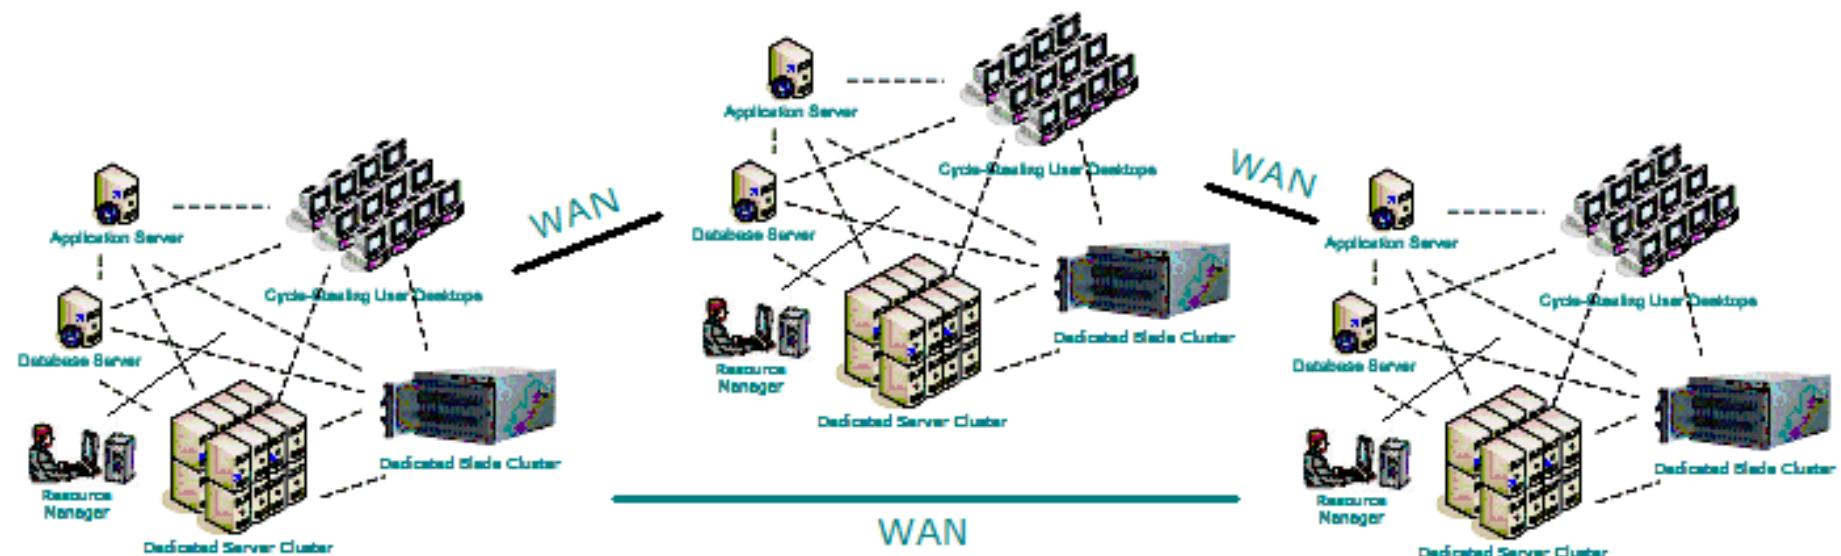


In general, a grid computing system requires:

- **At least one computer, usually a server, which handles all the administrative duties for the System**
- **A network of computers running special grid computing network software.**
- **A collection of computer software called middleware**

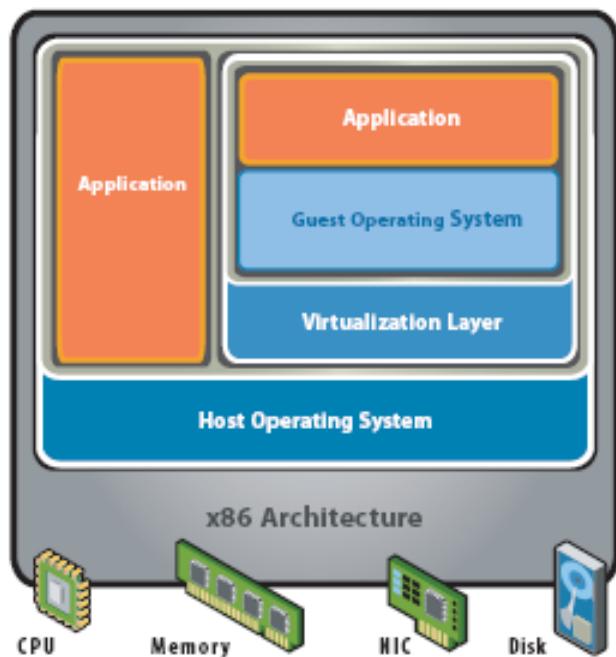
1.3 Distributed Systems: GRID

Grid Computing



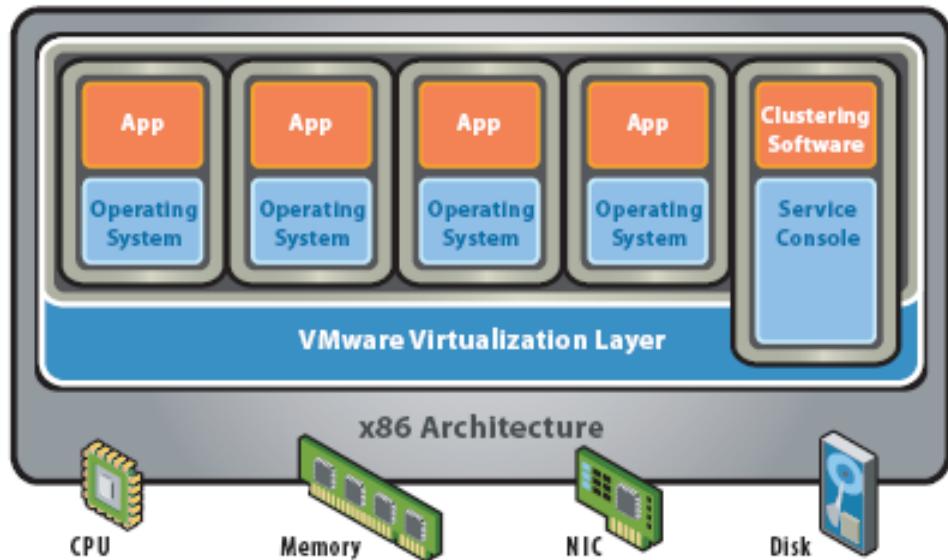
Distributed Systems: Cloud **Architecture**

Cloud Concepts – Intro – Virtualization Overview



Hosted Architecture

- Installs and runs as an application
- Relies on host OS for device support and physical resource management

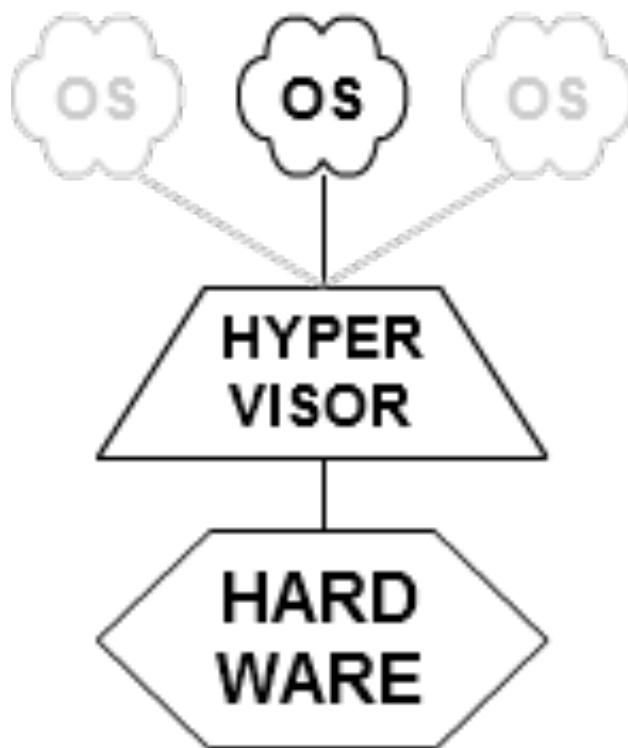


Bare-Metal (Hypervisor) Architecture

- Lean virtualization-centric kernel
- Service Console for agents and helper applications

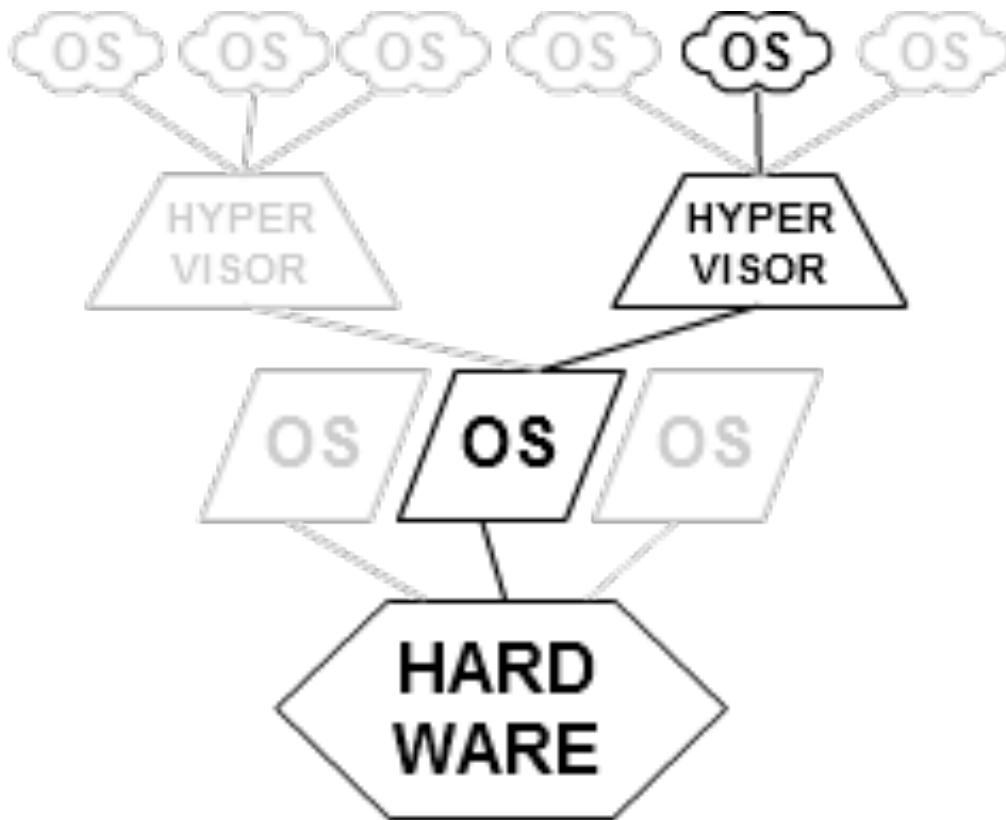
Figure 2: Virtualization Architectures

Cloud Concepts – Intro – Virtualization Overview



TYPE 1

*native
(bare metal)*



TYPE 2

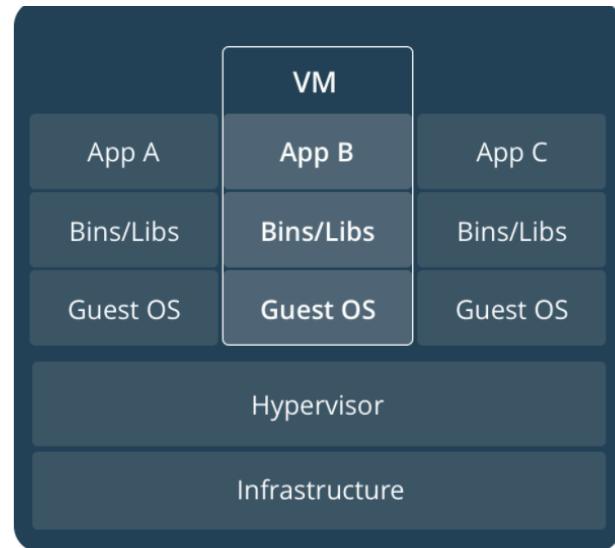
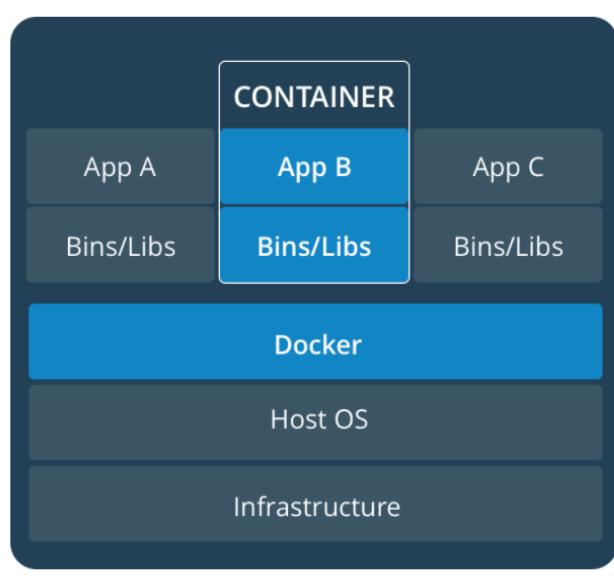
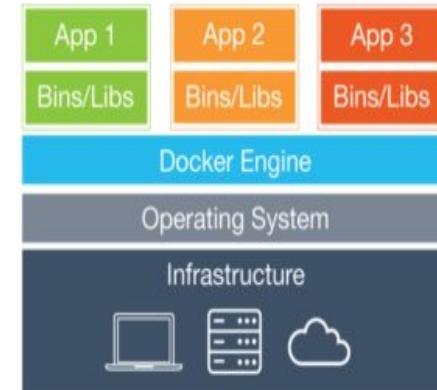
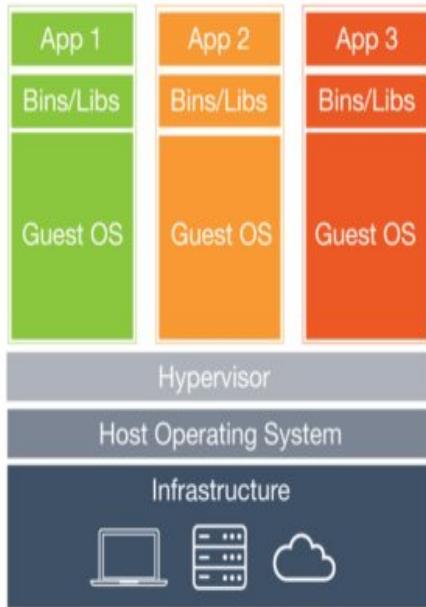
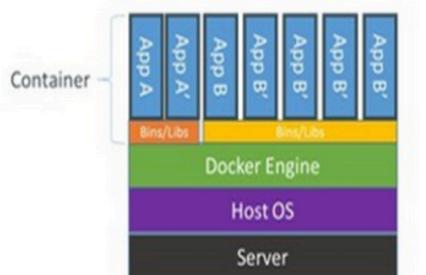
hosted

Deployment: Std. OS vs. VMs vs. Docker Containers

Containers vs. VMs

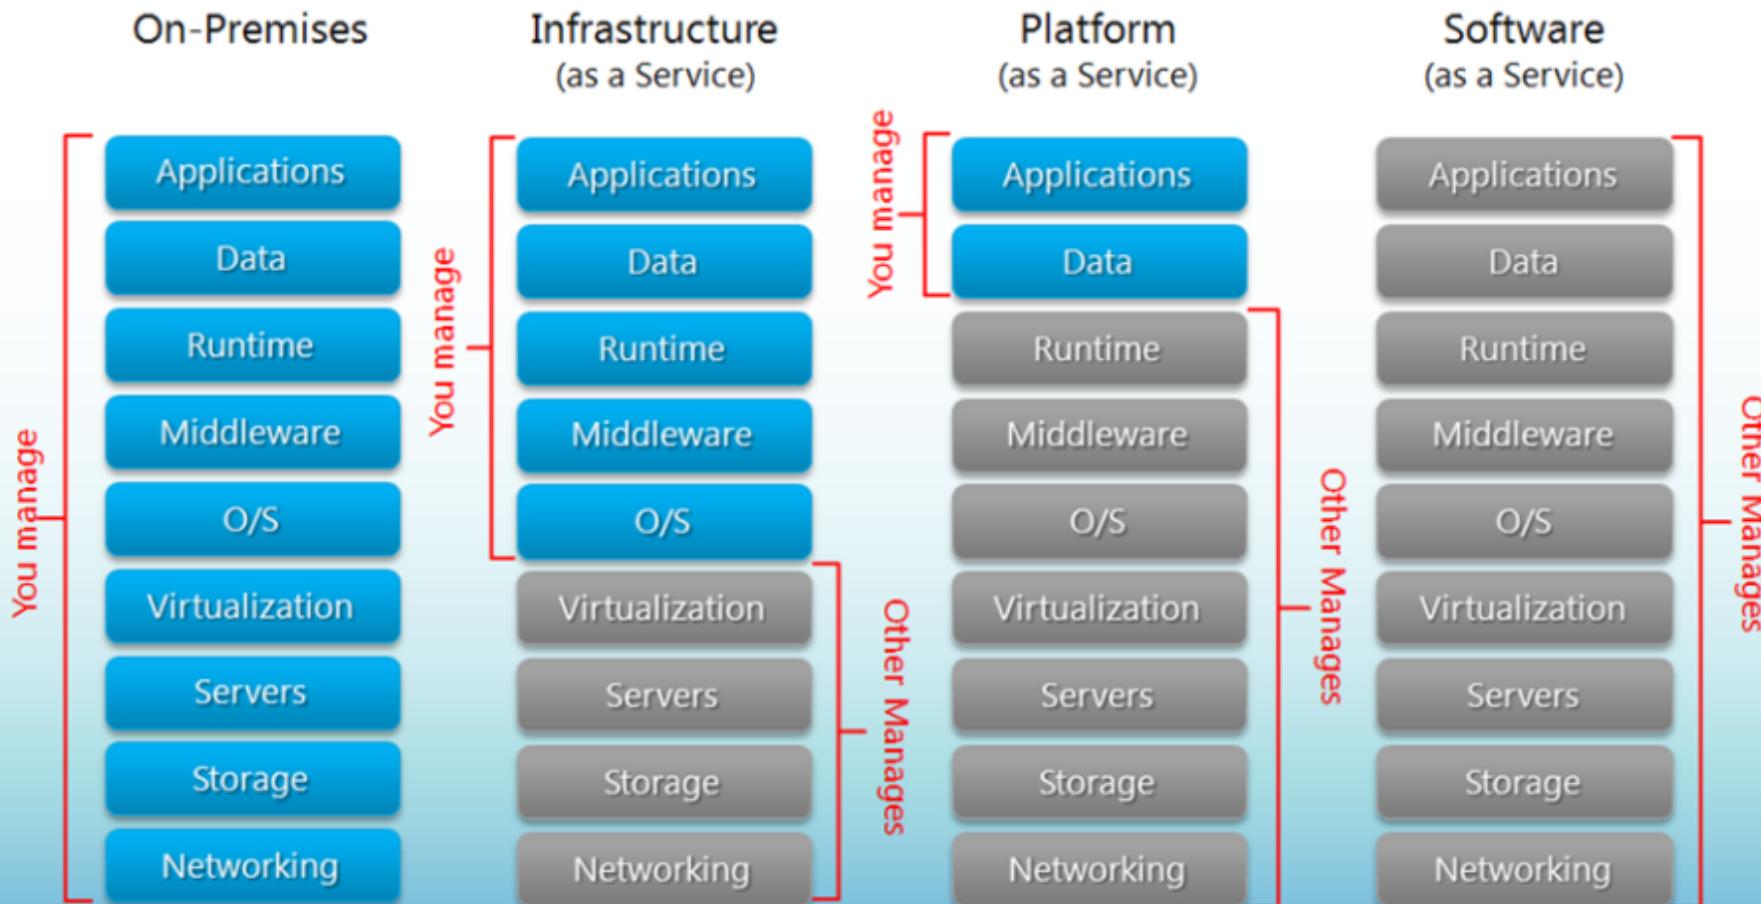


Containers are isolated,
but share OS and, where
appropriate, bins/libraries



1.3 Distributed Systems: Cloud Separation of Responsibilities

Cloud Concepts – Intro – IaaS, PaaS, SaaS



Copyright to

<http://blogs.technet.com/b/kevinremde/archive/2011/04/03/saas-paas-and-iaas-oh-my-quot-cloudy-april-quot-part-3.aspx>

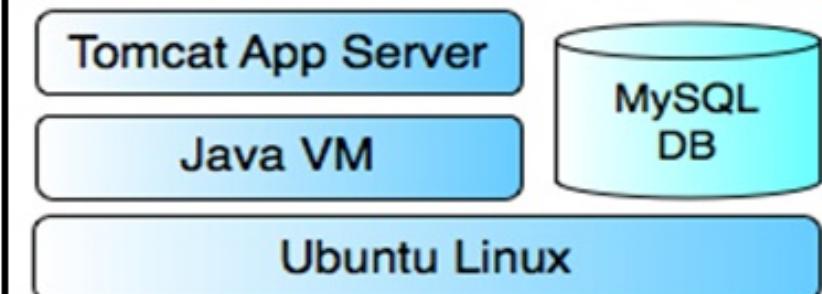
Cloud Concepts – Intro – IaaS, PaaS, SaaS

SaaS

A screenshot of a software application window titled "Accounts". The menu bar includes "File", "Sections", "Lookups", "Tools", "Call Centre", and "Help". The toolbar contains icons for Back, Home, Forward, Accounts, Contacts, Sales, Campaign, Tasks, Documents, Products, Processes, Reports, Library, Email, and Web. The main pane displays a list of accounts with columns for Account, City, Phone 1, Address, and Email. The list includes entries like "Business Solutions" (Kansas City KS), "Tennsoft" (Kansas City MO), "Boyle Inlet Company" (Ottawa), "Maverick Paper" (Kanata City KS), "MacHardware" (Wichita), "May Instruments" (Kansas City KS), "Avalon Technologies" (Overland Park), and "Versent" (Ottawa). Buttons at the bottom include "Add...", "Copy...", "Edit...", "Delete...", and "Script...". A "Details: Maverick Paper" section shows contact information for the company.

SalesForce.com, Google Apps

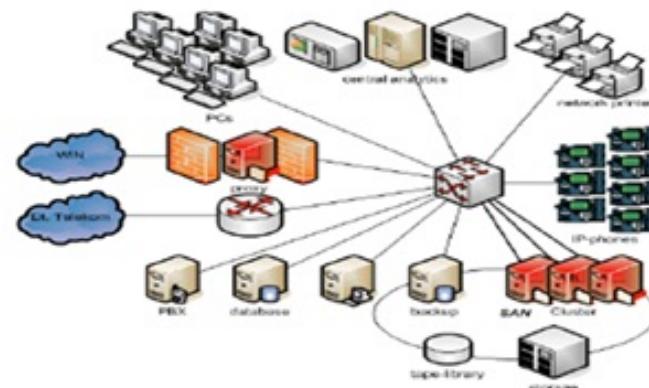
PaaS



Google App Engine for:
Java, Ruby, Python & GO

VMForce.com, MS Azure

IaaS



vCloud Express/Datacenter,
Amazon EC2

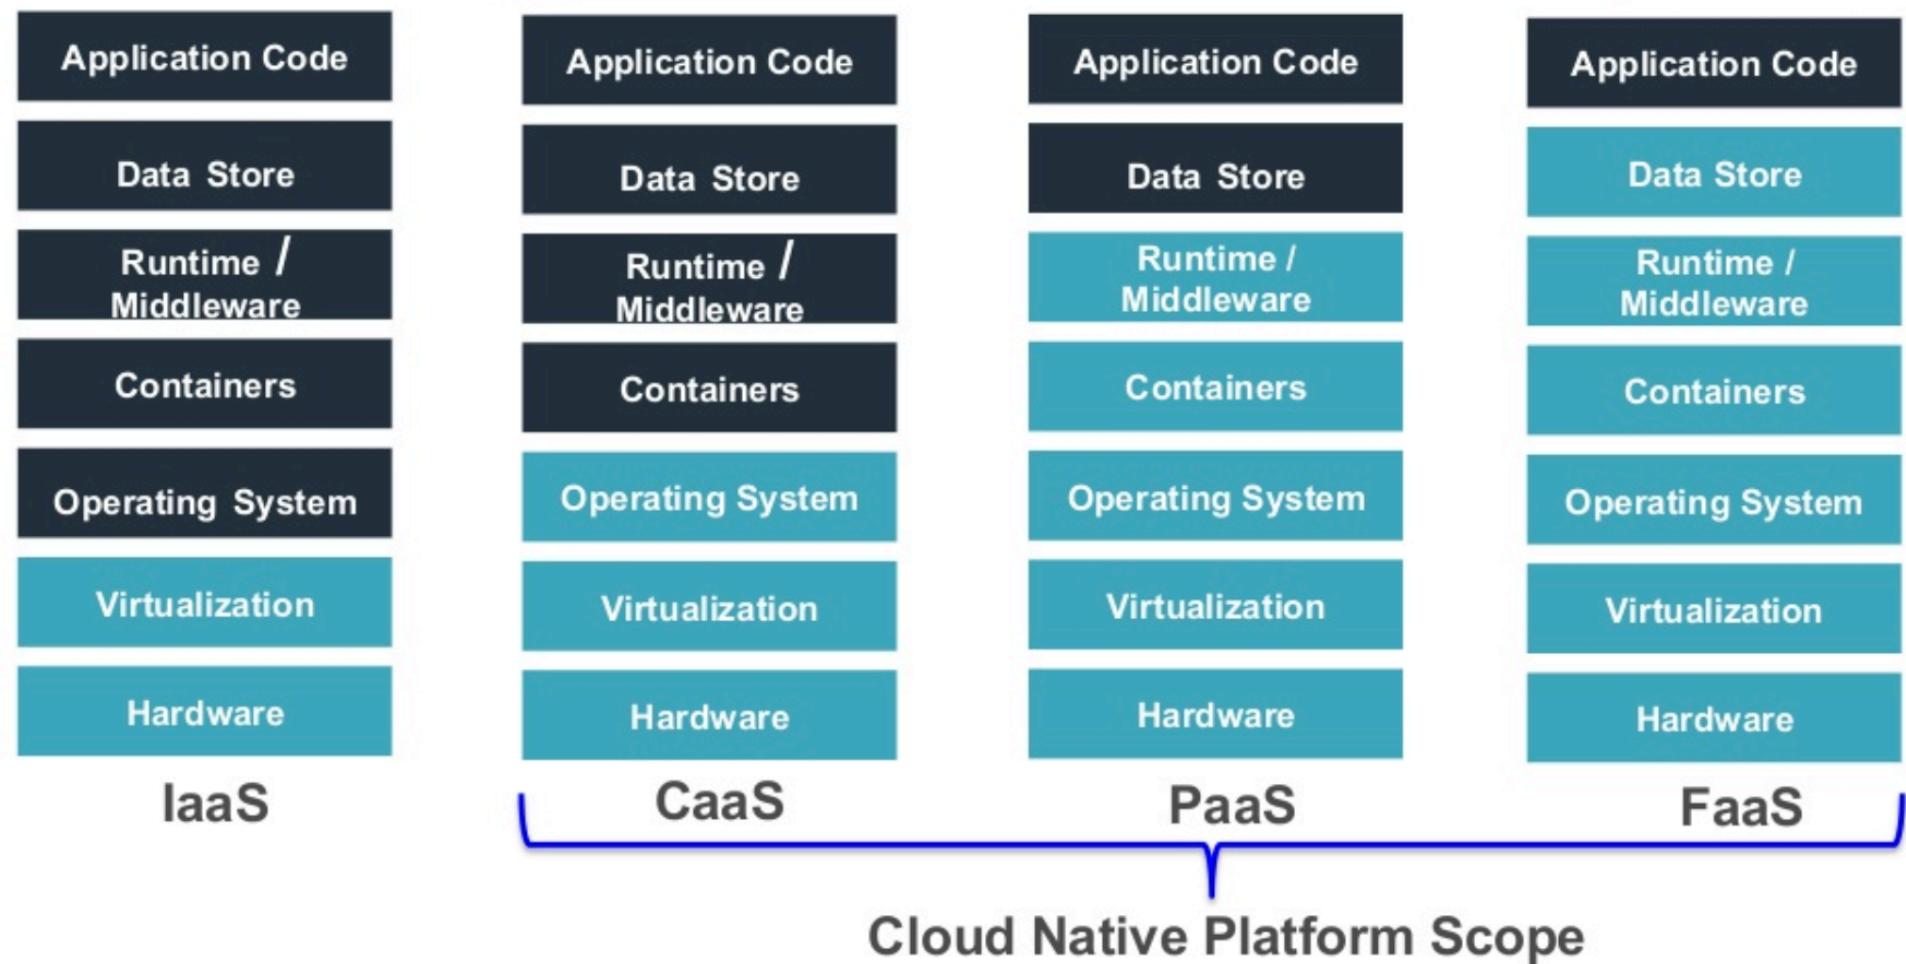
Copyright to

<http://it20.info/2010/11/random-thoughts-and-blasphemies-around-iaas-paas-saas-and-the-cloud-contract/>

Cloud Concepts – Intro – IaaS, PaaS, SaaS

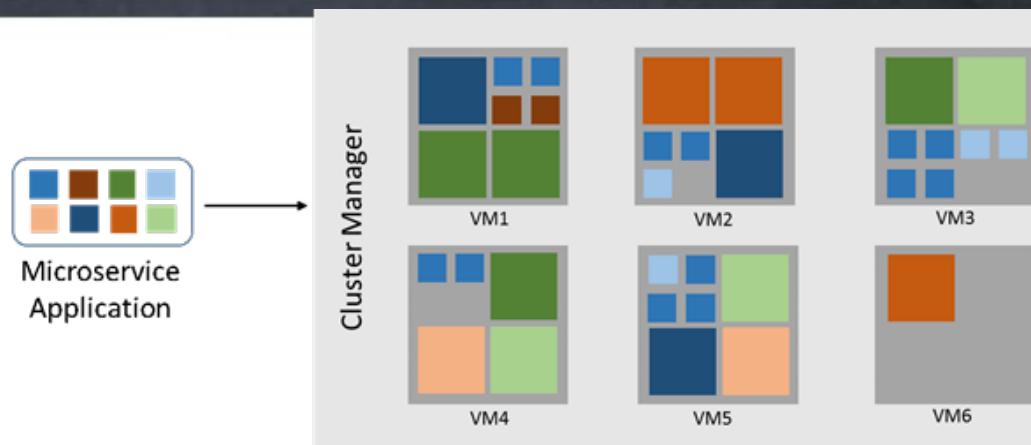


Cloud Concepts – Intro – IaaS, PaaS, CaaS, SaaS



1.3 Distributed Systems: Microservices

Micro-services Architecture



Microservice Architecture

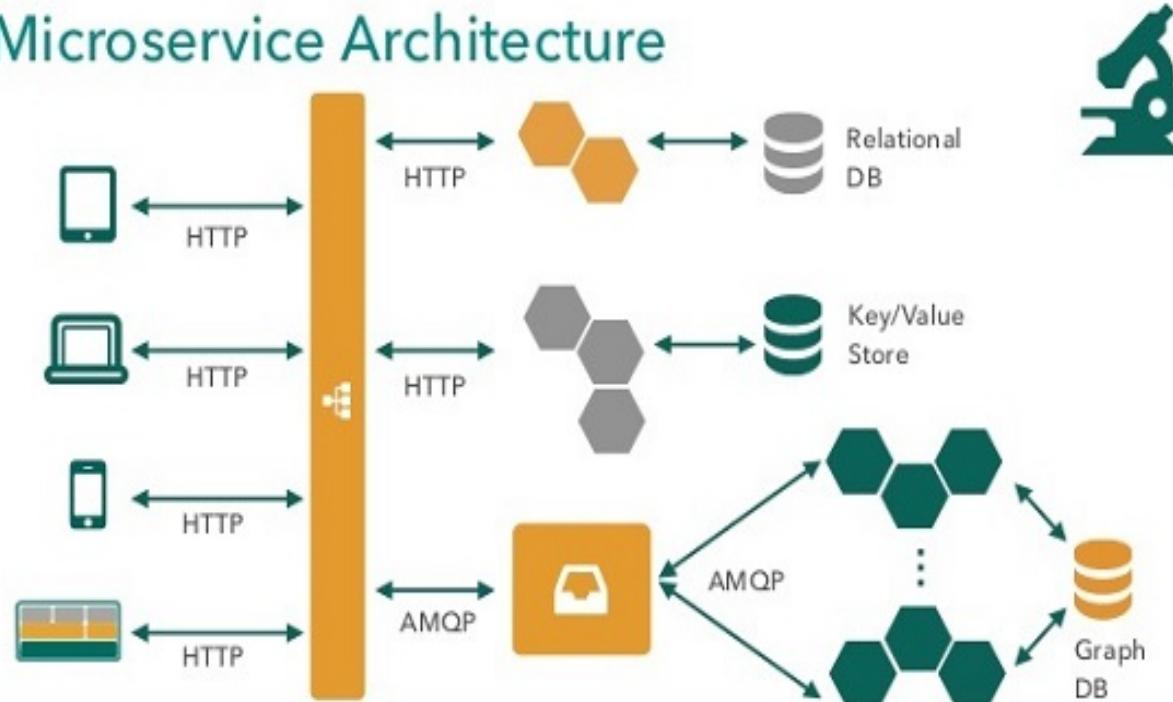
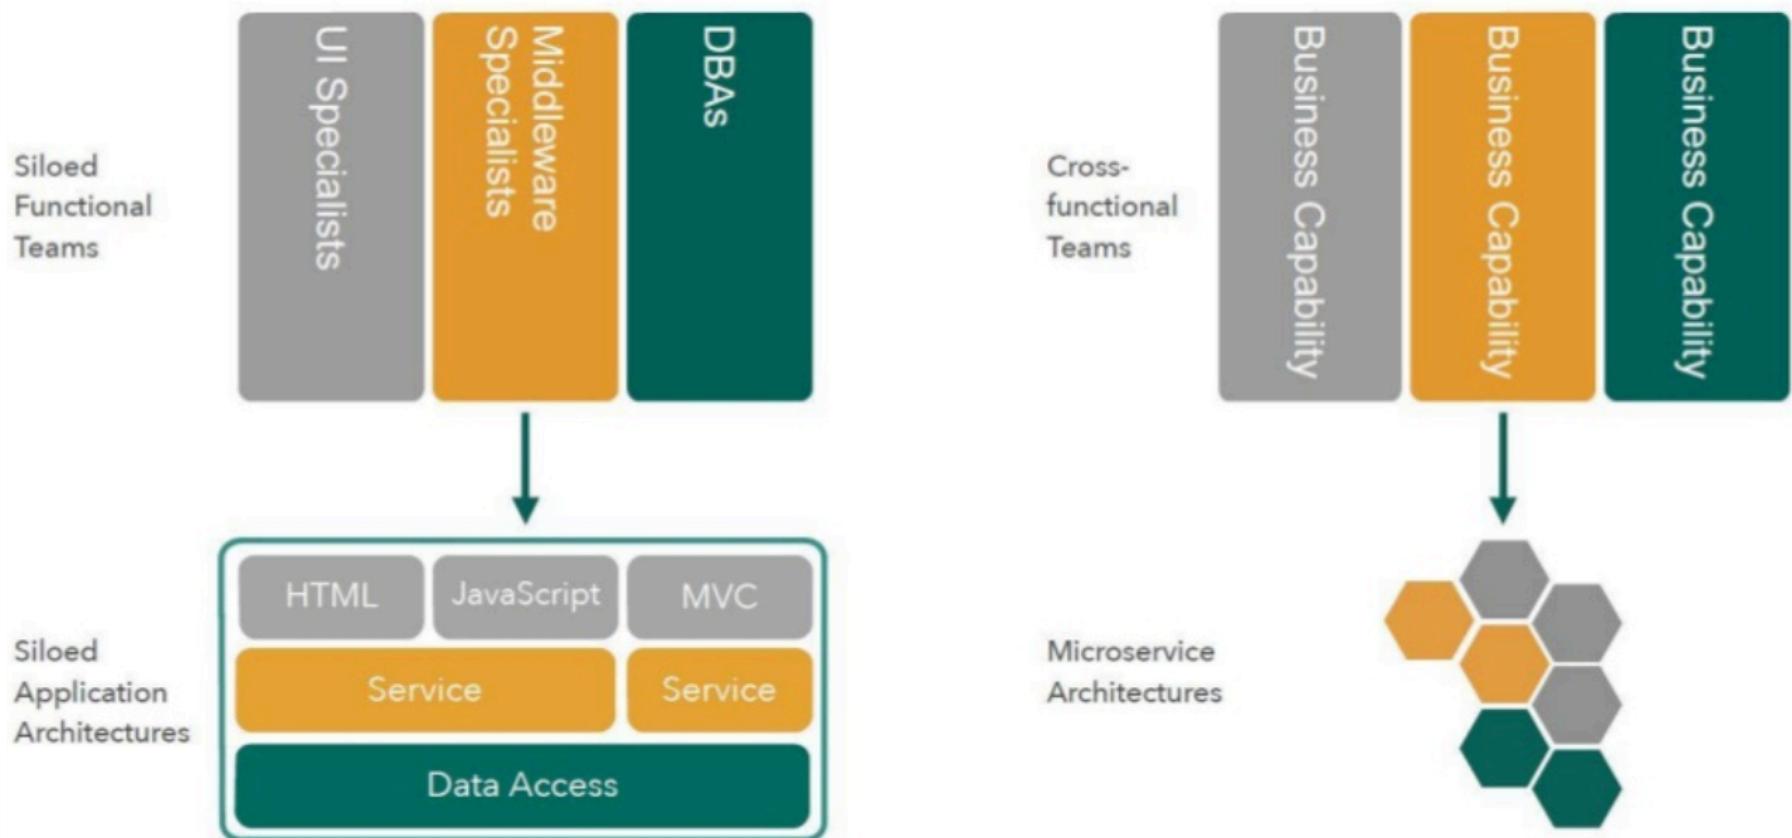


Image courtesy **Pivotal**.

1.3 Distributed Systems

Micro-services Architecture



1.4 Distributed Systems Challenges

1. Heterogeneity – variety and difference in:

- networks
- computer hardware
- OS
- programming languages
- implementations by different developers

2. OPENNESS

- computer system - can the system be extended and re-implemented in various ways?
- distributed system - can new resource-sharing services be added and made available for use by variety of client programs?

3. Security

- Confidentiality
- Integrity
- Availability

4. Scalability

–the ability of the system system to work well when the system load or the number of users increases

Challenges with building scalable distributed systems:

- Controlling the cost of physical resources
- Controlling the performance loss
- Preventing software resources running out (like 32-bit internet addresses, which are being replaced by 128 bits)
- Avoiding performance bottlenecks

5. Failure handling

Techniques for dealing with failures

- Detecting failures
- Masking failures
 - messages can be retransmitted
 - disks can be replicated in a synchronous action
- Tolerating failures
- Recovery from failures
- Redundancy – redundant components
 - at least two different routes
 - database can be replicated in several servers

Main goal: High availability

1.4 Distributed Systems Challenges

6. Concurrency – Several clients trying to access shared resource at the same time.

Any object with shared resources in a DS must be responsible that it operates correctly in a concurrent environment.

7. Transparency

Transparency – concealment from the user and the applications programmer of the separation of components in a Distributed System for the system to be perceived as a whole rather than a collection of independent components:

- Access transparency – access to local and remote resources identical
- Location transparency – resources accessed without knowing their physical or network location
- Concurrency transparency – concurrent operation of processes using shared resources without interference between them
- Replication transparency – multiple instances seem like one
- Failure transparency – fault concealment
- Mobility transparency – movement of resources/clients within a system without affecting the operation of users or programs

8. Quality of service

Main nonfunctional properties of systems that affect *Quality of Service (QoS)*:

- reliability
- security
- performance

Section Conclusion

Fact: **DAD needs Java**

In few **samples** it is simple to remember:
Types of the distributed architectural
patterns of the distributed systems.





Java OOP + Build Automation: ANT/Maven/Gradle, Annotation, Reflection, I/O, JNI, Linux IPC – Inter-Process Communication, light-weight processes / process thread in C/C++ Linux – pthread library and C++ '11, JVM and OS threads, Java Multi-threading issues

Java RECAP, Linux IPC, Multi-threading & Java

2.1.1 Java Reflection

- **Java Reflection** is an “introspective technique” that allows a computer program to examine and modify the structure and behavior (specifically the values, meta-data, properties and functions) of an object at runtime. [WIKI]
- **Java Reflection** is an advanced technique and should be used by experienced programmers that have good knowledge of Java and JVM.
- **Java Reflection** is a technique which allows different applications to do various operations that are quit impossible otherwise. It is a common approach for high-level programming languages as Java or C#.

2.1.1 Java Reflection

Samples for objects and objects arrays in **Java Reflection**:

- Operator ***instanceof***
- Displaying class methods
- Obtaining info about constructors methods
- Obtaining info about class fields
- Invoking methods by name
- Creation of new objects
- Changing the value from various field
- Using the arrays/vectors in Java Reflection context

2.1.2 Java Annotations

- Java Annotation “is the meta-tags that you will use in your code to give it some life.”
- There are two types: “***annotation type***” and “***annotation***”
- Define annotation – “***annotation type***”:

```
public @interface MyAnnotation {  
    String doSomething();  
}
```

- Use annotation – “***annotation***”:

```
@MyAnnotation (doSomething="What to do")  
public void mymethod() { .... }
```

2.1.2 Java Annotations

Three kind of “**annotation type**”:

- **1. Marker** – does NOT have internal elements

Sample:

```
public @interface MyAnnotation { }
```

Usage:

```
@MyAnnotation
```

```
public void mymethod() { .... }
```

- **2. Single Element** – has a single element represented by key=value

Sample:

```
public @interface MyAnnotation {  
    String doSomething();  
}
```

Usage:

```
@MyAnnotation ("What to do")  
public void mymethod() { .... }
```

2.1.2 Java Annotations

Kind of “***annotation type***”:

- **3. Full-Value / Multi-Value** – has multiple internal elements

Sample:

```
public @interface MyAnnotation {  
    String doSomething();  
    int count;  
    String date();  
}
```

Usage:

```
@MyAnnotation (doSomething="What to do", count=1, date="09-09-  
2005")  
public void mymethod() { .... }
```

2.1.2 Java Annotations

Rules for defining – “**annotation type**” :

1. The defining of the annotation should start with ‘@interface’ keyword.
2. The declared methods has no parameters.
3. The declared methods has no “throw exception” statements.
4. The data types of the method are:
 - * primitive – byte, char, int, float, double, etc.
 - * String
 - * Class
 - * enum
 - * arrays of one of the types from above – int[], float[], etc.

In JDK 5.0 there are predefined / simple – “**annotation**” :

1. @Override
2. @Deprecated
3. @SupressWarnings

2.1.2 Java Annotations

Starting with JDK 5.0 there are “*meta-annotation*” that can be applied only to the “*annotation type*”:

1. Target

```
@Target(ElementType.TYPE)
@Target(ElementType.FIELD)
@Target(ElementType.METHOD)
@Target(ElementType.PARAMETER)
@Target(ElementType.CONSTRUCTOR)
@Target(ElementType.LOCAL_VARIABLE)
@Target(ElementType.ANNOTATION_TYPE)
```

2. Retention

- `@Retention(RetentionPolicy.SOURCE)` – retinute la nivel cod sursa si sunt ignoreate de compilator
- `@Retention(RetentionPolicy.CLASS)` – retinute la nivel de compilare dar ignoreate de VM la run-time
- `@Retention(RetentionPolicy.RUNTIME)` – sunt retinute si utilizeaza doar la run-time

3. Documented – `@Documented`

4. Inherited – `@Inherited`

2.1.3 Java Library

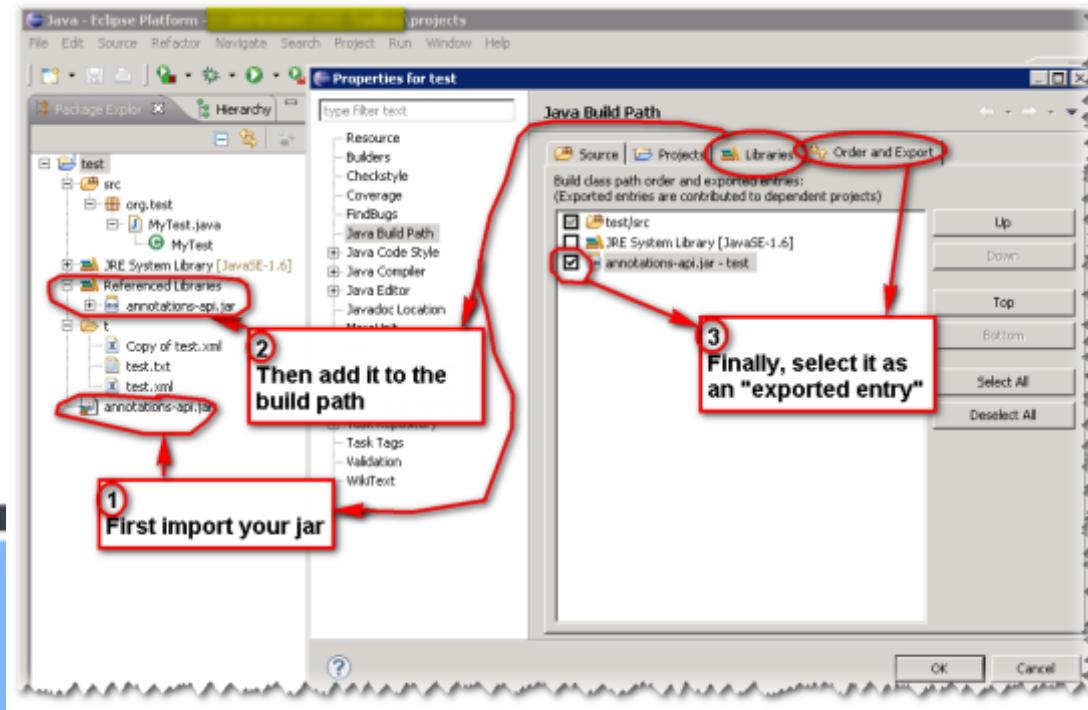
What is a **Java library**?

What are the advantages and disadvantages of Java libraries?

How can be solved in Java multiple dependencies or inclusions of the same class in the compilation phase? How were solved these problems in C/C++?

How should be created a Java library and how should be used – command line vs. IDE?

```
>jar -cvf archive_name.jar files_names_to_compress  
>javac -classpath .:archive_name.jar *.java  
>java -classpath .:archive_name.jar file_with_main_class
```



2.2.1 Build Automation: Apache ANT

Optional steps to build the lectures with Gradle – all lectures are independent by any Build

Automation system (e.g. Gradle, MVN, ANT, etc.) and the labs are projects in Eclipse IDE:

Apache ANT is a Java based build tool from Apache Software Foundation. Apache Ant's build files are written in XML and they take advantage of being open standard, portable and easy to understand.

Features:

- Ant is the most complete Java build and deployment tool available.
- Ant is platform neutral and can handle platform specific properties such as file separators.
- Ant can be used to perform platform specific tasks such as modifying the modified time of a file using 'touch' command.
- Ant scripts are written using plain XML. If you are already familiar with XML, you can learn Ant pretty quickly.
- Ant is good at automating complicated repetitive tasks.
- Ant comes with a big list of predefined tasks.
- Ant provides an interface to develop custom tasks.
- Ant can be easily invoked from the command line and it can integrate with free and commercial IDEs.

2.2.1 Build Automation: Apache ANT

```
<project>

    <target name="clean">
        <delete dir="build"/>
    </target>

    <target name="compile">
        <mkdir dir="build/classes"/>
        <javac srcdir="src" destdir="build/classes"/>
    </target>

    <target name="jar">
        <mkdir dir="build/jar"/>
        <jar destfile="build/jar/HelloWorld.jar" basedir="build/classes">
            <manifest>
                <attribute name="Main-Class" value="oata.HelloWorld"/>
            </manifest>
        </jar>
    </target>

    <target name="run">
        <java jar="build/jar/HelloWorld.jar" fork="true"/>
    </target>

</project>
```

ant compile jar run

2.2.1 Build Automation: Apache ANT

While having a look at the buildfile, we will see some similar steps between Ant and the java-only commands:

java-only	Ant
md build\classes javac -sourcepath src -d build\classes src\oata\HelloWorld.java echo Main-Class: oata.HelloWorld>mf md build\jar jar cfm build\jar\HelloWorld.jar mf -C build\classes . java -jar build\jar\HelloWorld.jar	<mkdir dir="build/classes"/> <javac srcdir="src" destdir="build/classes"/> <!-- automatically detected --> <!-- obsolete; done via manifest tag --> <mkdir dir="build/jar"/> <jar destfile="build/jar/HelloWorld.jar" basedir="build/classes"> <manifest> <attribute name="Main-Class" value="oata.HelloWorld"/> </manifest> </jar> <java jar="build/jar/HelloWorld.jar" fork="true"/>

ant compile

ant jar

ant run

2.2.2 Build Automation: Apache MVN - Maven

Optional steps to build the lectures with Gradle – all lectures are independent by any Build Automation system (e.g. Gradle, MVN, ANT, etc.) and the labs are projects in Eclipse IDE:

Apache Maven (MVN) is a software project management and comprehension tool. Based on the concept of a project object model (POM), Maven can manage a project's build, reporting and documentation from a central piece of information. Using maven we can build and manage any Java based project.

Features:

- Maven uses **Convention over Configuration**, which means developers are not required to create build process themselves. Developers do not have to mention each and every configuration detail. Maven provides sensible default behavior for projects. When a Maven project is created, Maven creates default project structure. Developer is only required to place files accordingly and he/she need not to define any configuration in **pom.xml**.
- Simple project setup that follows best practices.
- Consistent usage across all projects.
- Dependency management including automatic updating.

2.2.2 Build Automation: Apache MVN - Maven

Apache Maven (MVN) Features:

- A large and growing repository of libraries.
- Extensible, with the ability to easily write plugins in Java or scripting languages.
- Instant access to new features with little or no extra configuration.
- **Model-based builds** – Maven is able to build any number of projects into predefined output types such as jar, war, metadata.
- **Coherent site of project information** – Using the same metadata as per the build process, maven is able to generate a website and a PDF including complete documentation.
- **Release management and distribution publication** – Without additional configuration, maven will integrate with your source control system such as CVS and manages the release of a project.
- **Backward Compatibility** – You can easily port the multiple modules of a project into Maven 3 from older versions of Maven. It can support the older versions also.
- **Automatic parent versioning** – No need to specify the parent in the sub module for maintenance.
- **Parallel builds** – It analyzes the project dependency graph and enables you to build schedule modules in parallel. Using this, you can achieve the performance improvements of 20-50%.
- **Better Error and Integrity Reporting** – Maven improved error reporting, and it provides you with a link to the Maven wiki page where you will get full description of the error.

2.2.2 Build Automation: Apache MVN - Maven

Creating a Project

You will need somewhere for your project to reside, create a directory somewhere and start a shell in that directory. On your command line, execute the following Maven goal:

```
1. mvn archetype:generate -DgroupId=com.mycompany.app -DartifactId=my-app -DarchetypeArtifactId=maven-archetype-quickstart -DinteractiveMode=false
```

If you have just installed Maven, it may take a while on the first run. This is because Maven is downloading the most recent artifacts (plugin jars and other files) into your local repository. You may need to execute the command a couple of times before it succeeds. This is because the remote server may time out before your downloads are complete. Don't worry, there are ways to fix that.

You will notice that the generate goal created a directory with the same name given as the artifactId. Change into that directory.

```
1. cd my-app
```

Under this directory you will notice the following [standard project structure](#).

```
1. my-app
2. |-- pom.xml
3. '-- src
4.   |-- main
5.   |   '-- java
6.   |       '-- com
7.   |           '-- mycompany
8.   |               '-- app
9.   |                   '-- App.java
10.  '-- test
11.    '-- java
12.     '-- com
13.         '-- mycompany
14.             '-- app
15.                 '-- AppTest.java
```

The `src/main/java` directory contains the project source code, the `src/test/java` directory contains the test source, and the `pom.xml` file is the project's Project Object Model, or POM.

2.2.2 Build Automation: Apache MVN - Maven

The POM

The pom.xml file is the core of a project's configuration in Maven. It is a single configuration file that contains the majority of information required to build a project in just the way you want. The POM is huge and can be daunting in its complexity, but it is not necessary to understand all of the intricacies just yet to use it effectively. This project's POM is:

```
1. <project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
2.   xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/xsd/maven-4.0.0.xsd">
3.   <modelVersion>4.0.0</modelVersion>
4.
5.   <groupId>com.mycompany.app</groupId>
6.   <artifactId>my-app</artifactId>
7.   <version>1.0-SNAPSHOT</version>
8.   <packaging>jar</packaging>
9.
10.  <name>Maven Quick Start Archetype</name>
11.  <url>http://maven.apache.org</url>
12.
13.  <dependencies>
14.    <dependency>
15.      <groupId>junit</groupId>
16.      <artifactId>junit</artifactId>
17.      <version>4.8.2</version>
18.      <scope>test</scope>
19.    </dependency>
20.  </dependencies>
21. </project>
```

mvn clean build
mvn package

java -cp target/my-app-1.0-SNAPSHOT.jar com.mycompany.app.App

2.3 Build Automation: Gradle

Optional steps to build the lectures with Gradle – all lectures are independent by any Build Automation system (e.g. Gradle, MVN, ANT, etc.) and the labs are projects in Eclipse IDE:

Gradle is an advanced general purpose **build management system** based on Groovy.

Features:

- It supports automatic download and configuration of project dependencies (libraries).
- It supports Maven and Ivy (ANT) repositories for retrieving project dependencies, allowing reuse of the artifacts of existing build systems.
- It supports multi-project and multi-artifact builds.

(Optional full Gradle presentation @ISM – Cyber-Security Master – www.ism.ase.ro, by Marius Popa @ Oracle)

<https://gradle.org/install/>

<https://plugins.gradle.org/>

<https://github.com/jabedhasan21/java-hello-world-with-gradle/blob/master/README.md>

<https://www.tutorialspoint.com/gradle/index.htm>

2.3.1. Gradle Intro and Basics

Projects and tasks in **Gradle**:

- A Gradle build consists of one or more projects.
- A **project** using Gradle describes its build via a **build.gradle** file. **build.gradle** file is located in the root folder of the project. **build.gradle** file is based on a *Domain Specific Language* (DSL) written in Groovy.
- **build.gradle** build file defines a project and its tasks.
- A **task** represents a piece of work which a build performs.

```
task hello {  
    doLast {  
        println 'Hello Gradle'  
        println 'Project name: ' + rootProject.name  
    }  
}
```

```
mepopa@MEPOPA-RO /d/work/P006_IoT_GW/GradleTraining  
$ gradle project  
:projects  
-----  
Root project  
-----  
Root project 'GradleTraining'  
No sub-projects  
To see a list of the tasks of a project, run gradle <project-path>:tasks  
For example, try running gradle :tasks  
BUILD SUCCESSFUL  
Total time: 3.676 secs
```

```
mepopa@MEPOPA-RO /d/Work/P006_IoT_GW/GradleTraining  
$ gradle hello  
:hello  
Hello Gradle  
Project name: GradleTraining  
BUILD SUCCESSFUL  
Total time: 3.162 secs
```

2.3.2. Installing and configuring Gradle

- **Requirement:** JDK (Java Development Kit) installation.
- **Installing:** Download the latest release of Gradle from <http://gradle.org/gradle-download/> | <https://gradle.org/install/> for the usage on the command line. Add the /bin folder to PATH environment variable.
- **Gradle daemon:** avoid starting the Java virtual machine for every build. Usage: `org.gradle.daemon=true` in `gradle.properties` in the `${HOME}/.gradle`; Executing gradle with the `--daemon` parameter on the command line or `gradle --stop` to stop it.
- Set specific JVM options: `GRADLE_OPTS` environment variable. Eg. `export GRADLE_OPTS=-Xmx1024m` defines 1 GB as maximum heap size.

2.3.3. Gradle plug-ins

- Plug-ins extend Gradle core functionality. Extension to Gradle adding some preconfigured tasks.
- Gradle ships with a number of plug-ins (**java**), and custom plug-ins can be developed.
- Adding a plug-in in ***build.gradle*** file: **apply plugin: 'pluginname'** (Eg. **apply plugin: 'java'**).
- Registry for Gradle plug-ins: <https://plugins.gradle.org/>.

2.3.3. Gradle plug-ins

build.gradle

```
task hello {  
    doLast {  
        println 'Hello Gradle'  
        println 'Project name: ' + rootProject.name  
    }  
}
```

```
mepopa@MEPOPA-RO /d/Work/P006_IoT_GW/GradleTraining  
$ gradle tasks  
:tasks  
  
All tasks runnable from root project  
  
Build Setup tasks  
init - Initializes a new Gradle build. [incubating]  
wrapper - Generates Gradle wrapper files. [incubating]  
  
Help tasks  
components - Displays the components produced by root project 'GradleTraining'. [incubating]  
dependencies - Displays all dependencies declared in root project 'GradleTraining'.  
dependencyInsight - Displays the insight into a specific dependency in root project 'GradleTraining'.  
help - Displays a help message.  
model - Displays the configuration model of root project 'GradleTraining'. [incubating]  
projects - Displays the sub-projects of root project 'GradleTraining'.  
properties - Displays the properties of root project 'GradleTraining'.  
tasks - Displays the tasks runnable from root project 'GradleTraining'.  
  
Other tasks  
hello  
  
To see all tasks and more detail, run gradle tasks --all  
To see more detail about a task, run gradle help --task <task>  
BUILD SUCCESSFUL  
Total time: 3.418 secs
```

build.gradle

```
apply plugin: 'java'  
task hello {  
    doLast {  
        println 'Hello Gradle'  
        println 'Project name: ' + rootProject.name  
    }  
}
```

```
mepopa@MEPOPA-RO /d/Work/P006_IoT_GW/GradleTraining  
$ gradle tasks  
:tasks  
  
All tasks runnable from root project  
  
Build tasks  
assemble - Assembles the outputs of this project.  
build - Assembles and tests this project.  
buildDependents - Assembles and tests this project and all projects that depend on it.  
buildNeeded - Assembles and tests this project and all projects it depends on.  
classes - Assembles classes 'main'.  
jar - Assembles a jar archive containing the main classes.  
testClasses - Assembles classes 'test'.  
  
Build Setup tasks  
init - Initializes a new Gradle build. [incubating]  
wrapper - Generates Gradle wrapper files. [incubating]  
  
Documentation tasks  
javadoc - Generates Javadoc API documentation for the main source code.
```

```
mepopa@MEPOPA-RO /d/Work/P006_IoT_GW/GradleTraining  
$ gradle tasks  
:tasks  
  
All tasks runnable from root project  
  
Build tasks  
assemble - Assembles the outputs of this project.  
build - Assembles and tests this project.  
buildDependents - Assembles and tests this project and all projects that depend on it.  
buildNeeded - Assembles and tests this project and all projects it depends on.  
classes - Assembles classes 'main'.  
jar - Assembles a jar archive containing the main classes.  
testClasses - Assembles classes 'test'.  
  
Build Setup tasks  
init - Initializes a new Gradle build. [incubating]  
wrapper - Generates Gradle wrapper files. [incubating]  
  
Help tasks  
components - Displays the components produced by root project 'GradleTraining'. [incubating]  
dependencies - Displays all dependencies declared in root project 'GradleTraining'.  
dependencyInsight - Displays the insight into a specific dependency in root project 'GradleTraining'.  
help - Displays a help message.  
model - Displays the configuration model of root project 'GradleTraining'. [incubating]  
projects - Displays the sub-projects of root project 'GradleTraining'.  
properties - Displays the properties of root project 'GradleTraining'.  
tasks - Displays the tasks runnable from root project 'GradleTraining'.  
  
Verification tasks  
check - Runs all checks.  
clean - Deletes the build directory.  
test - Runs the unit test.  
  
Other tasks  
hello  
  
Rules  
Pattern: clean<taskName> - Cleans the output files of a task.  
Pattern: build<ConfigurationName> - Assembles the artifacts of a configuration.  
Pattern: upload<ConfigurationName> - Assembles and uploads the artifacts belonging to a configuration.  
  
To see all tasks and more detail, run gradle tasks --all  
To see more detail about a task, run gradle help --task <task>  
BUILD SUCCESSFUL  
Total time: 5.03 secs
```

2.3.4. Gradle dependency management

- Managing the classpath of Gradle projects: adding JAR files, directories or other projects to the build path of the application.
- Automatic download of Java library dependencies, specifying the dependency in Gradle build file.
- A Java library is identified by Gradle via its project's **groupId:artifactId:version** (also known as **GAV** in Maven).
- Adding a dependency: new entry in the dependency section in ***build.gradle*** file.

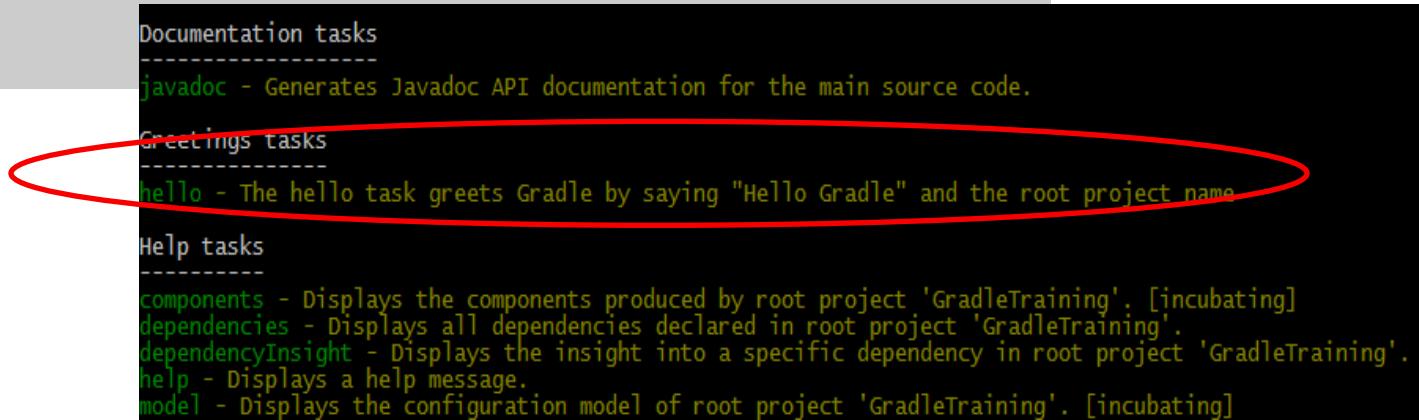
```
repositories {  
    jcenter()  
}  
  
dependencies {  
    compile fileTree(dir: 'libs',  
    include: ['*.jar'])  
    testCompile 'junit:junit:4.12'  
}
```

```
mepopa@MEPOPA-RO /d/Work/P006_IoT_GW/GradleTraining  
$ gradle clean build  
:clean UP-TO-DATE  
:compileJava UP-TO-DATE  
:processResources UP-TO-DATE  
:classes UP-TO-DATE  
:jar  
:assemble  
:compileTestJava UP-TO-DATE  
:processTestResources UP-TO-DATE  
:testClasses UP-TO-DATE  
:test UP-TO-DATE  
:check UP-TO-DATE  
:build  
  
BUILD SUCCESSFUL  
  
Total time: 4.357 secs
```

2.3.5. Gradle tasks

- Default Gradle tasks: tasks for introspection of Gradle itself (eg. task **tasks**).
- Custom Gradle tasks: tasks created by developers (eg. task **hello**). Running the **gradle tasks** task, the **hello** task will be listed under **Other tasks**. Tasks without a group are considered as private tasks. Groups can be applied with the **group** property and a description can be applied by using the **description** property .

```
task hello {  
  
    group 'greetings'  
    description 'The hello task greets Gradle by  
saying "Hello Gradle" and the root project name'  
  
    doLast {  
        println 'Hello Gradle'  
        println 'Project name: ' + rootProject.name  
    }  
}
```



2.3.6. Building Java projects

- Java plug-in: provides tasks to **compile** Java source code, run **unit tests**, create **Javadoc** and create a **JAR** file.
- Default project layout:
 - Java source code: **src/main/java**.
 - Java tests: **src/test/java**.
- Different project structure: **sourceSets**.
- Start the build: **gradle build** for **HelloWorld.java**.

```
sourceSets {  
    main {  
        java {  
            srcDir 'src'  
        }  
    }  
    test {  
        java {  
            srcDir 'test'  
        }  
    }  
}
```

```
mepopa@MEPOPA-R0 /d/work/P006_IoT_GW/GradleTraining  
$ gradle build  
T am not invoked, but I always get printed  
:compileJava  
:processResources UP-TO-DATE  
:classes  
:jar  
:assemble  
:compileTestJava UP-TO-DATE  
:processTestResources UP-TO-DATE  
:testClasses UP-TO-DATE  
:test UP-TO-DATE  
:check UP-TO-DATE  
:build  
  
BUILD SUCCESSFUL  
Total time: 5.434 secs
```

2.4 Jenkins/Hudson – DevOps Build Automation

```
export JAVA_HOME=/opt/software/java/jdks/jdk1.8.0_161
export CLASSPATH=.:$JAVA_HOME/jre/lib
export CATALINA_HOME=/opt/software/apache-tomcat-9.0.4
export
PATH=.:$JAVA_HOME/bin:$CATALINA_HOME/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/
bin:/sbin:/bin:/usr/games
```

```
cd $CATALINA_HOME
bin/startup.sh
```

```
http://127.0.0.1:8080
http://127.0.0.1:8080/jenkins/
```

```
bin/shutdown.sh
```

The screenshot shows the Jenkins dashboard interface. On the left, there is a sidebar with links: New Item, People, Build History, Manage Jenkins, My Views, Credentials, and New View. The main area displays a table of build items. One item is listed: 'buildAndRunC02'. The table has columns for Status (S), Workstation (W), Name, Last Success, Last Failure, and Last Duration. The 'Name' column is sorted by name. The 'Last Success' and 'Last Failure' columns show the date and time of the last successful and failed builds respectively. Below the table, there are links for 'Legend', 'RSS for all', 'RSS for failures', and 'RSS for just latest builds'. At the bottom, there are two sections: 'Build Queue' (which is empty) and 'Build Executor Status' (which shows 1 Idle and 2 Idle executors).

S	W	Name	Last Success	Last Failure	Last Duration
blue circle icon	cloud icon	buildAndRunC02	22 hr - #7	22 hr - #5	7.5 sec

Icon: S M L

Legend RSS for all RSS for failures RSS for just latest builds

Build Queue

No builds in the queue.

Build Executor Status

1 Idle
2 Idle

2.4 Jenkins/Hudson – DevOps Build Automation

buildAndRunC02 Config [Jenkins] - Mozilla Firefox

LibreOffice Writer

127.0.0.1:8080/jenkins/job/buildAndRunC02/configure

General Source Code Management Build Triggers Build Environment **Build** Post-build Actions

Execute shell

Command

```
cd /home/stud/jenkinsworkspace
if [ ! -d "./javase" ]; then
    # Control will enter here if $DIRECTORY doesn't exist.
    git clone https://github.com/critoma/javase.git
fi

cd ./javase
git pull

export JAVA_HOME=/opt/software/java/jdks/jdk1.8.0_161
export PATH=$JAVA_HOME/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/
export CLASSPATH=.:$JAVA_HOME/jre/lib
export JSE=/home/stud/jenkinsworkspace/javase/lectures

cd $JSE/c02/src
javac eu/ase/ooparrays/Student.java
javac eu/ase/ooparrays/ProgMainOopArrays.java

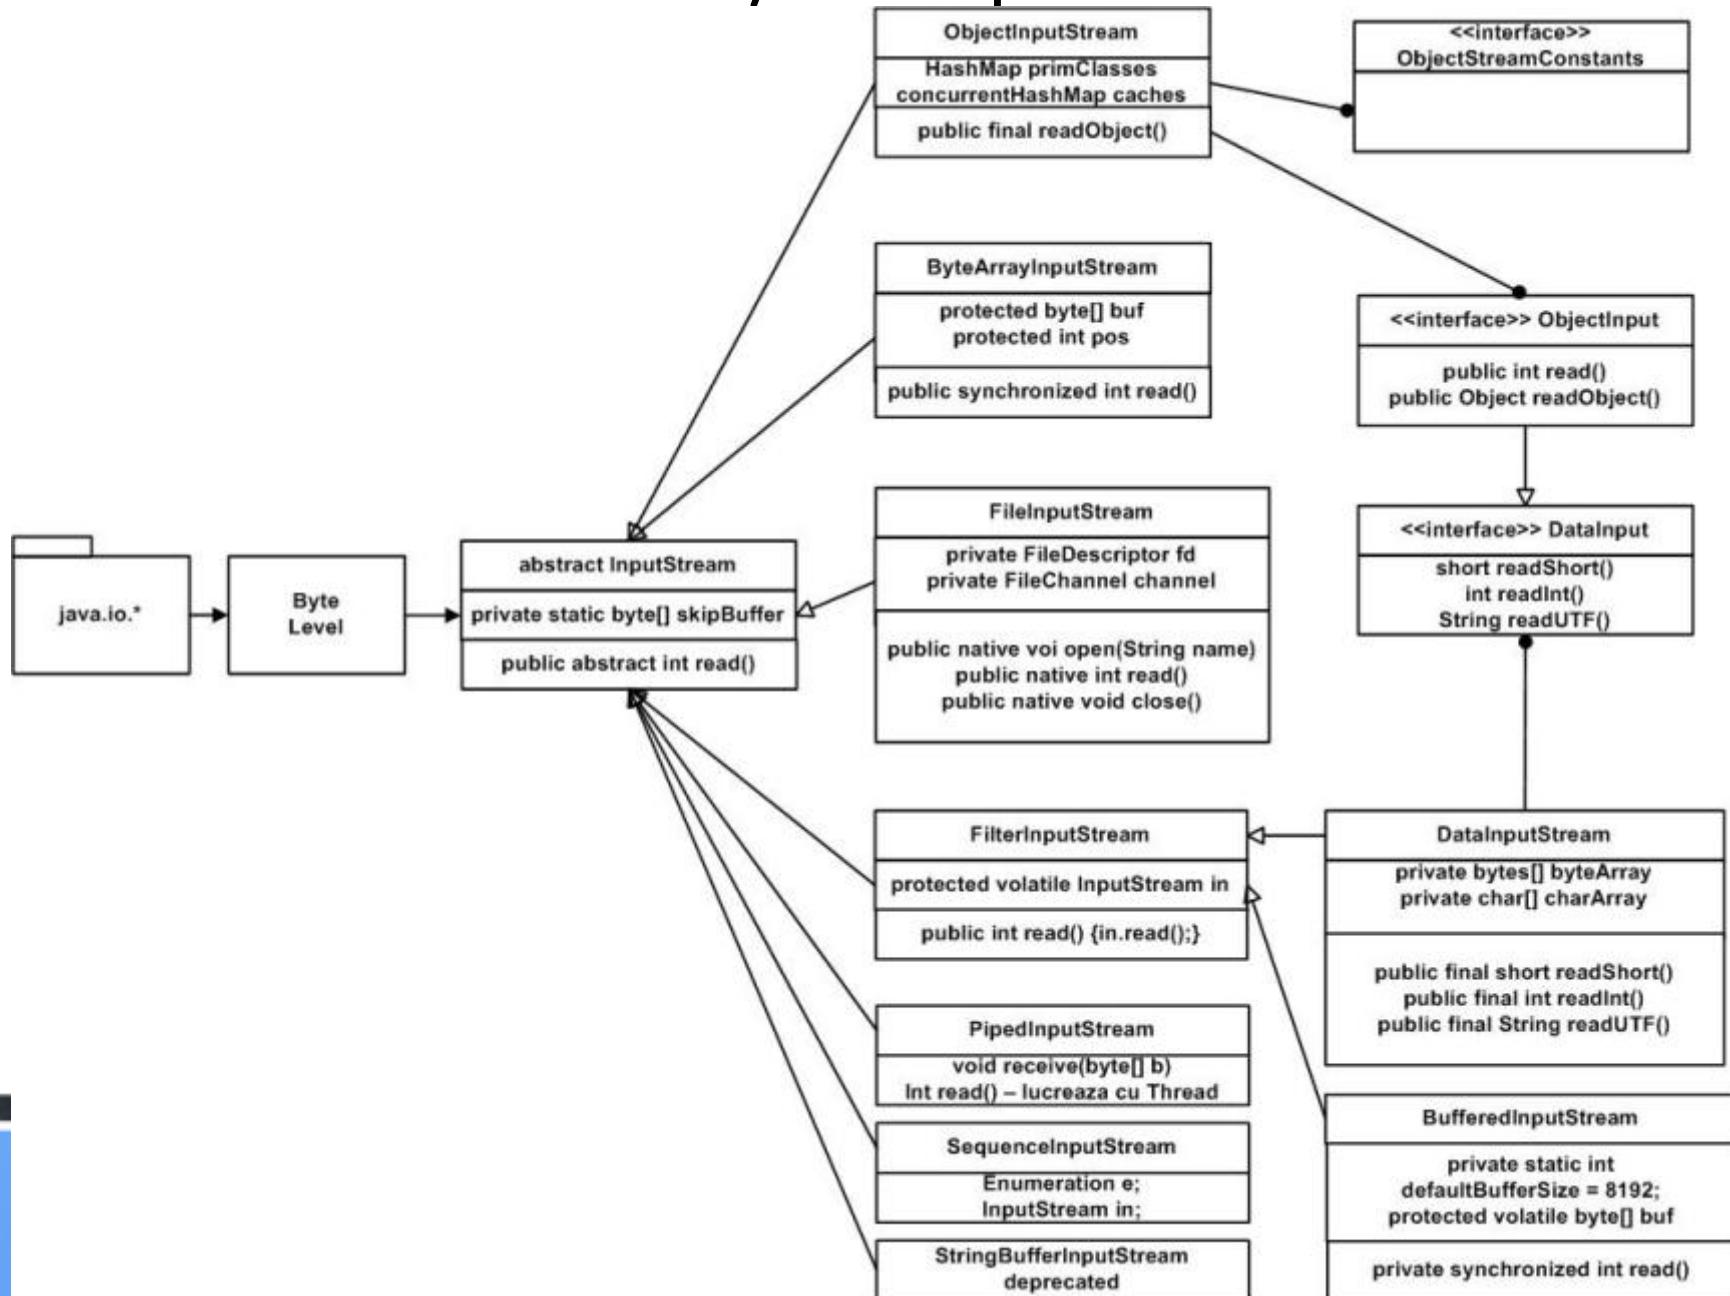
# full JAR deployment
jar -cmvf ./META-INF/MANIFEST.MF ../student.jar eu/ase/ooparrays/*.class
rm eu/ase/ooparrays/*.class

java -jar ../student.jar eu/ase/ooparrays/ProgMainOopArrays
rm ../student.jar
```

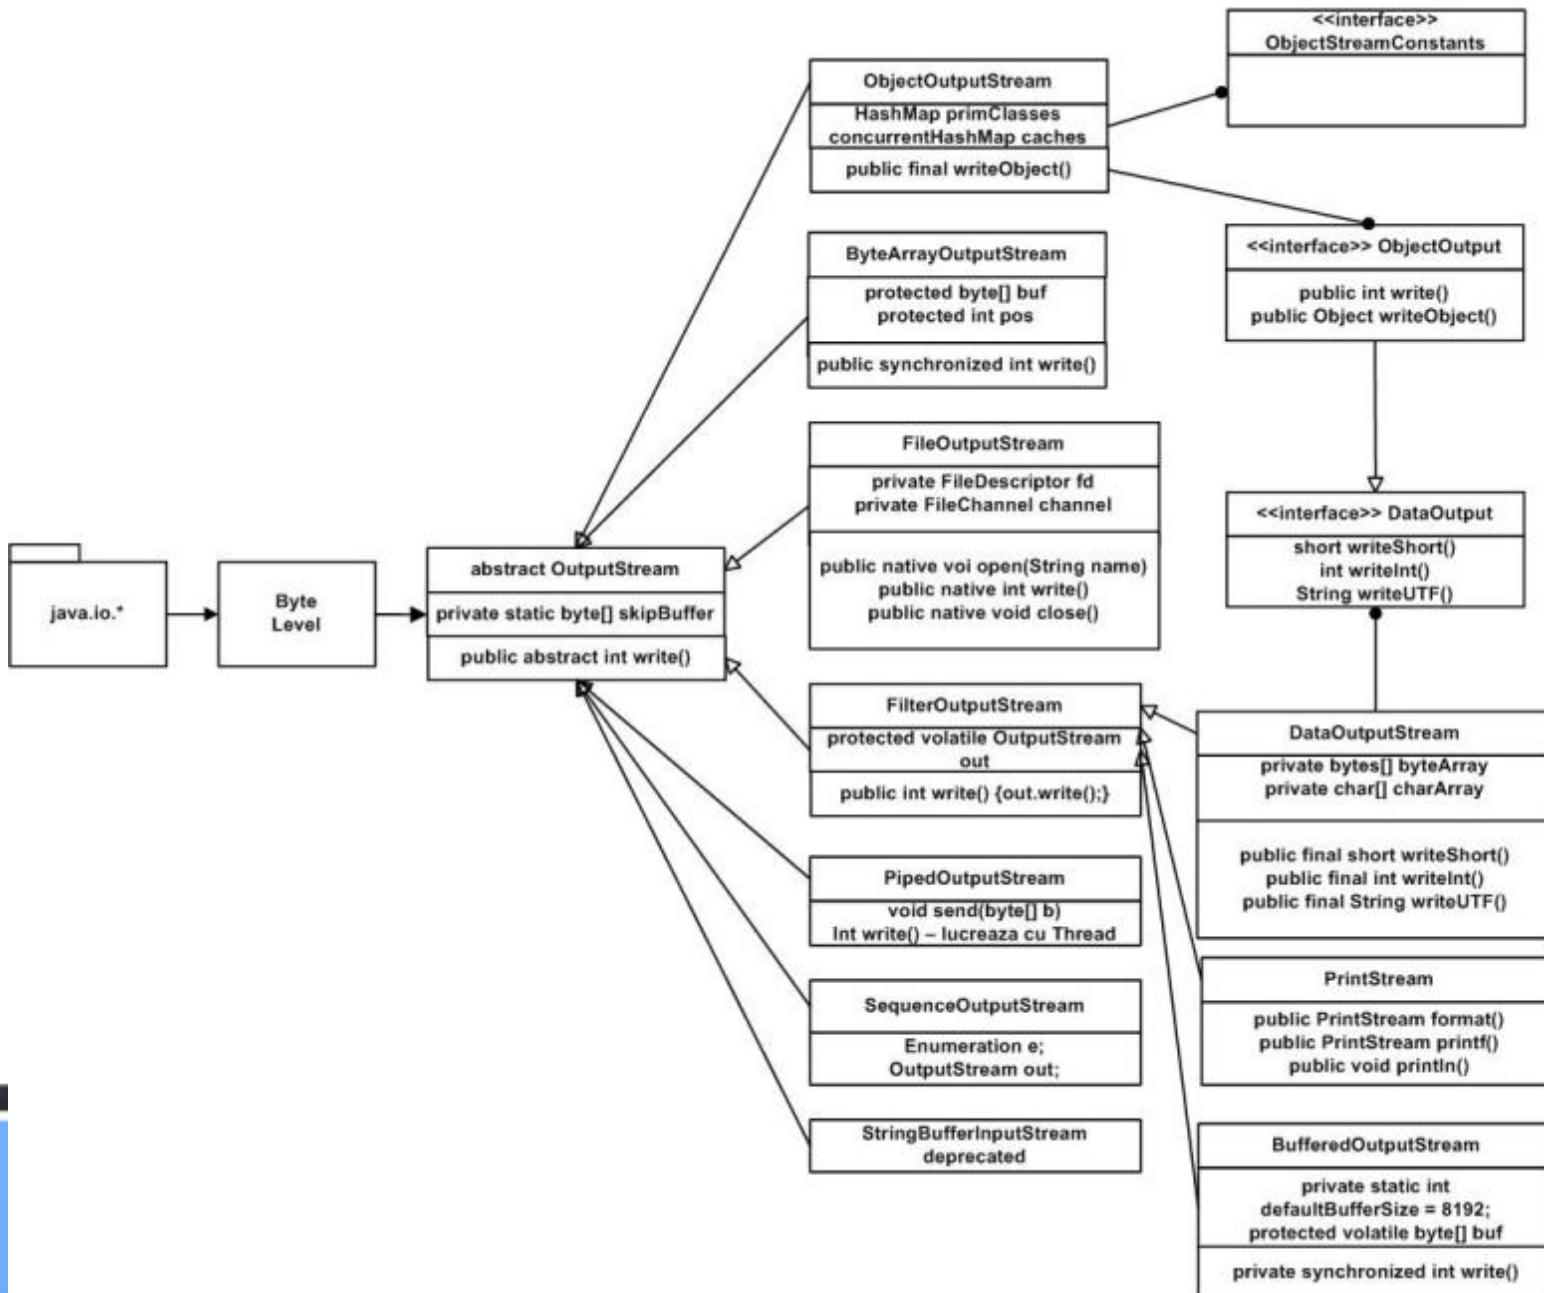
Save Apply

Plain Text ▾ Tab Width: 8 ▾ Ln 203, Col 1 ▾ INS

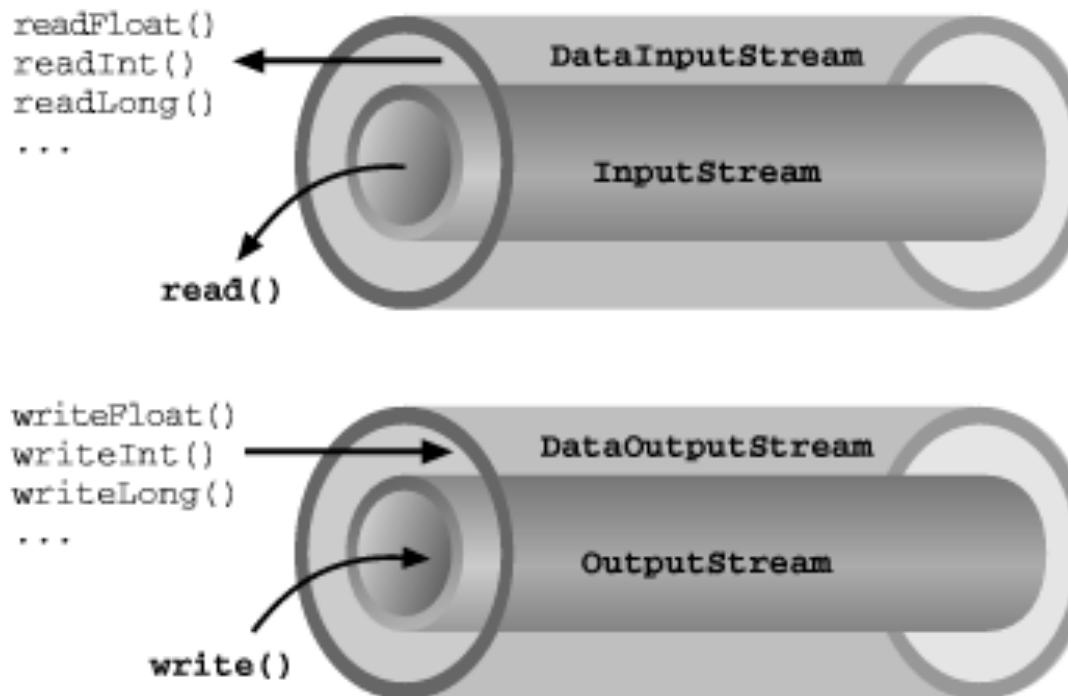
2.5 Java I/O – Input Stream



2.5 Java I/O – Output Stream

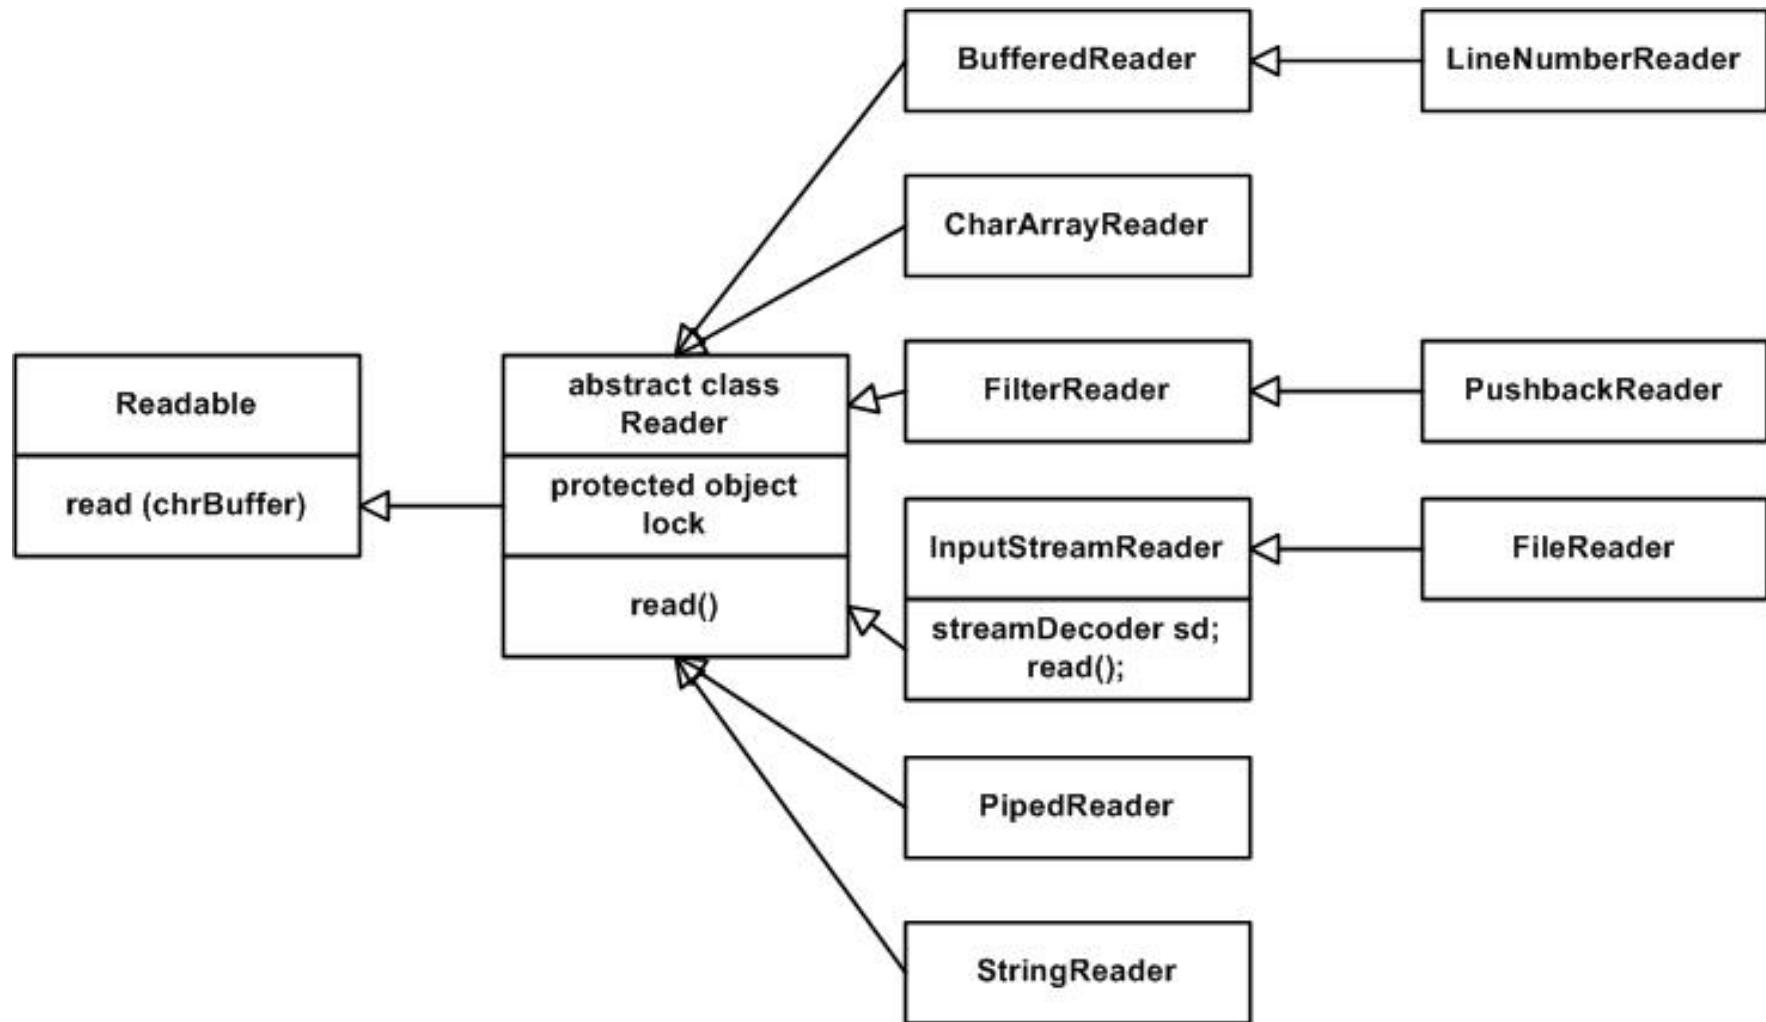


2.5 Java I/O – Streams Encapsulation

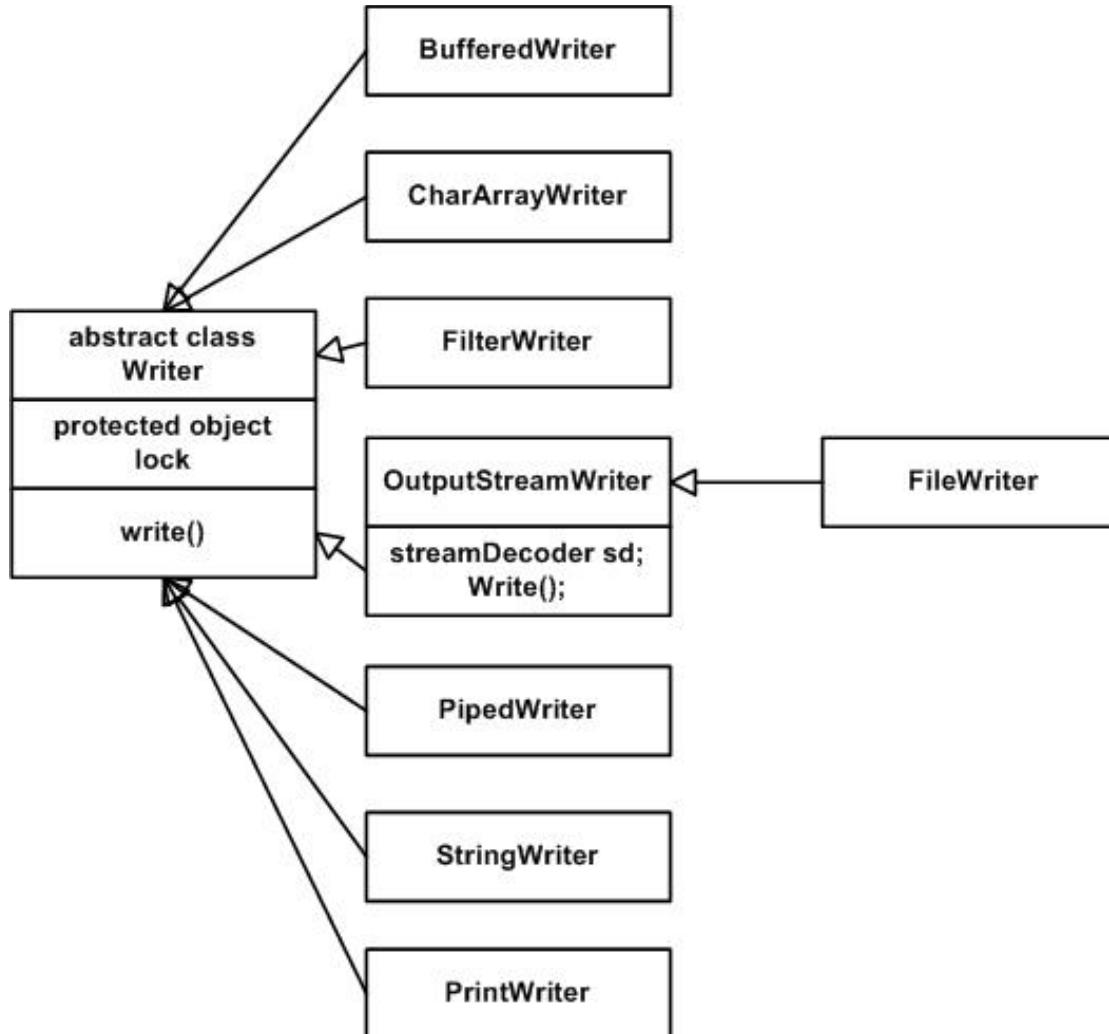


http://doc.sumy.ua/prog/java/exp/ch10_01.htm

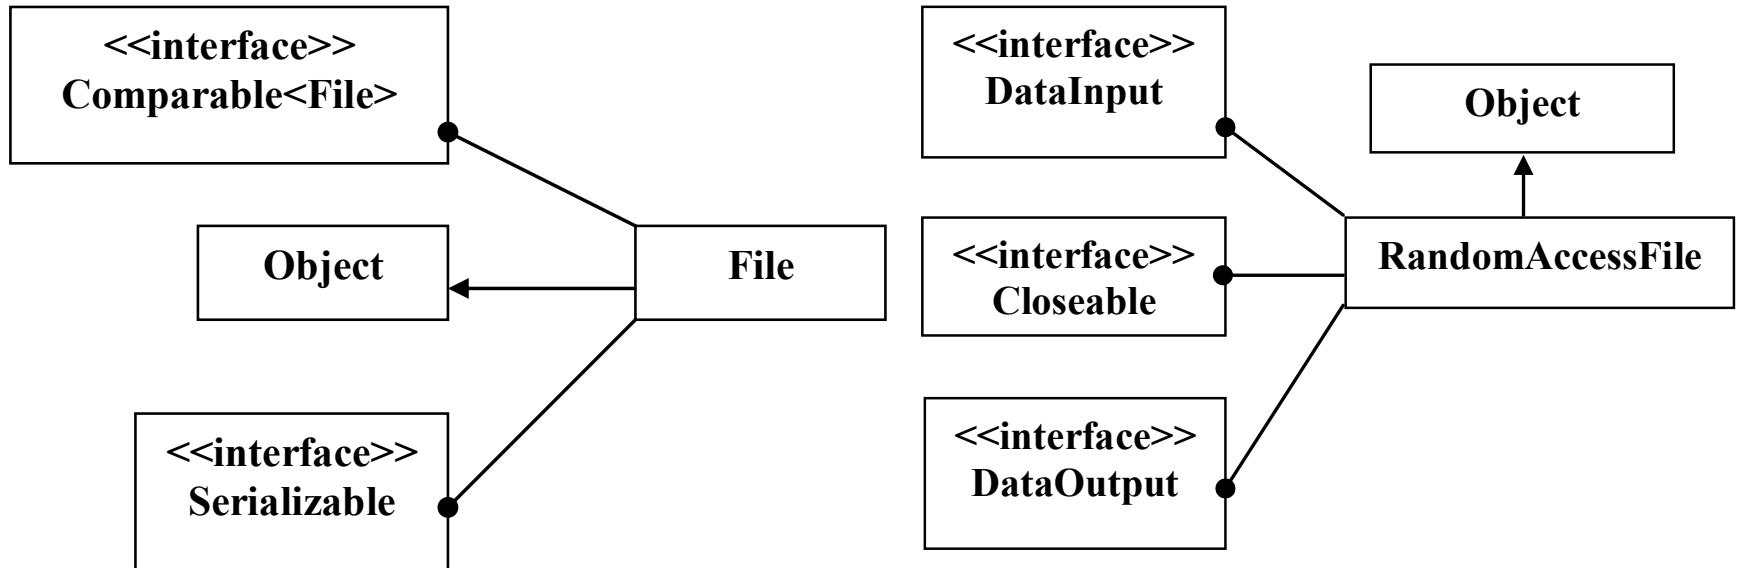
2.5 Java I/O – char level reading



2.5 Java I/O – char level writing

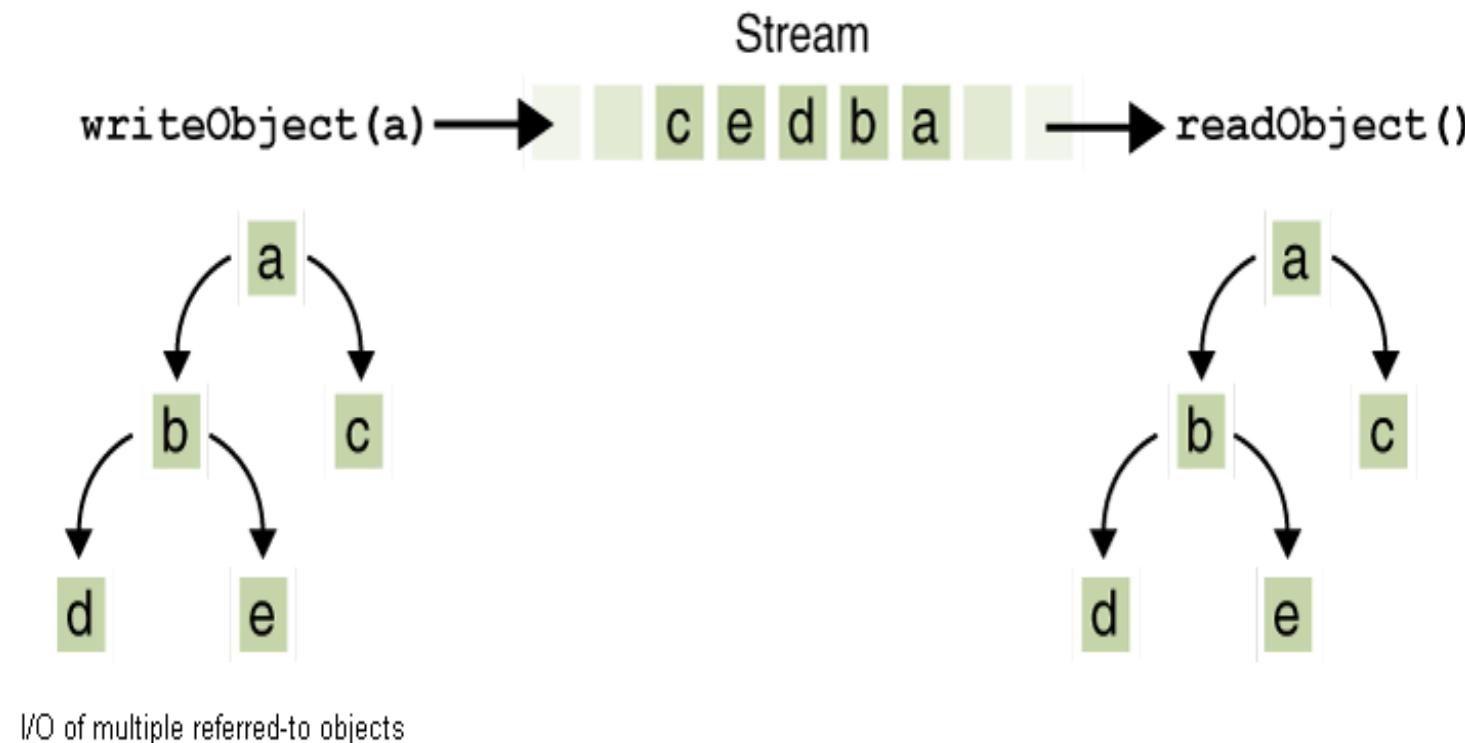


2.5 Java I/O – File Access



2.5 Java I/O – Serialization

This is demonstrated in the following figure, where `writeObject` is invoked to write a single object named `a`. This object contains references to objects `b` and `c`, while `b` contains references to `d` and `e`. Invoking `writeObject(a)` writes not just `a`, but all the objects necessary to reconstitute `a`, so the other four objects in this web are written also. When `a` is read back by `readObject`, the other four objects are read back as well, and all the original object references are preserved.

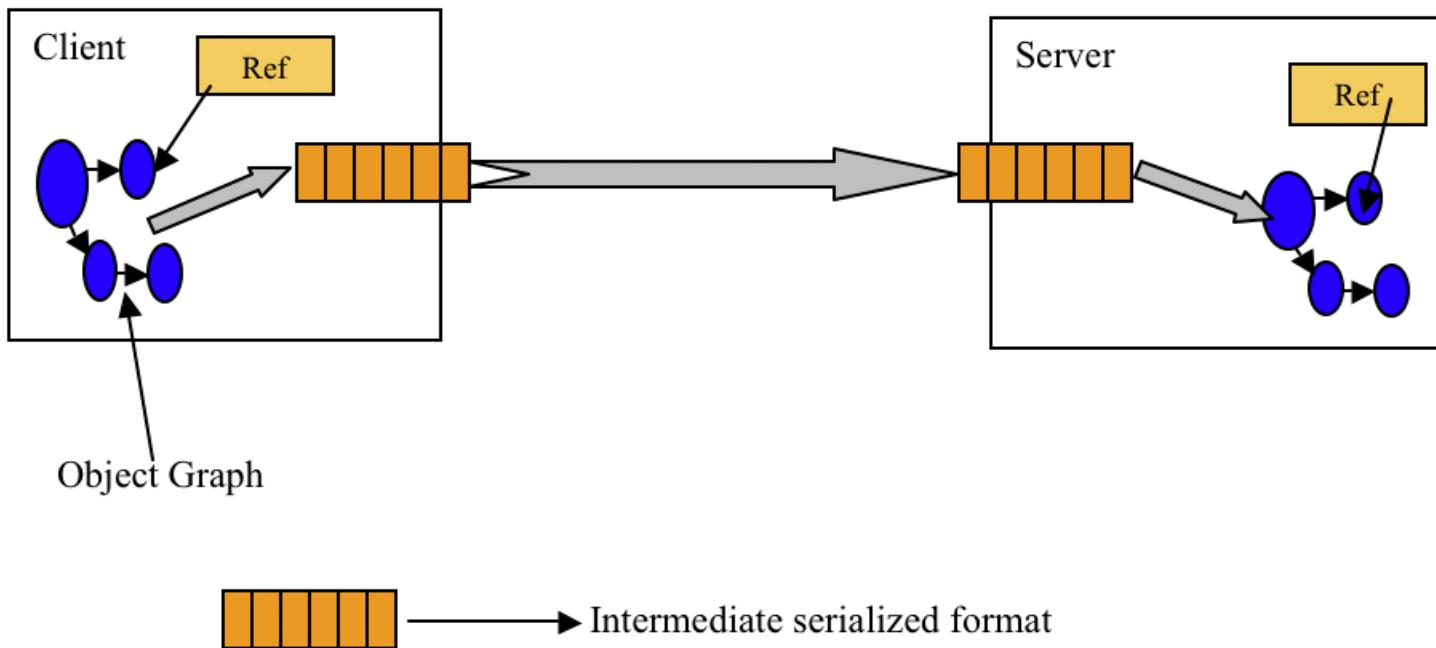


2.5 Java I/O – Serialization

What is going to be saved and restored by serialization in Java?

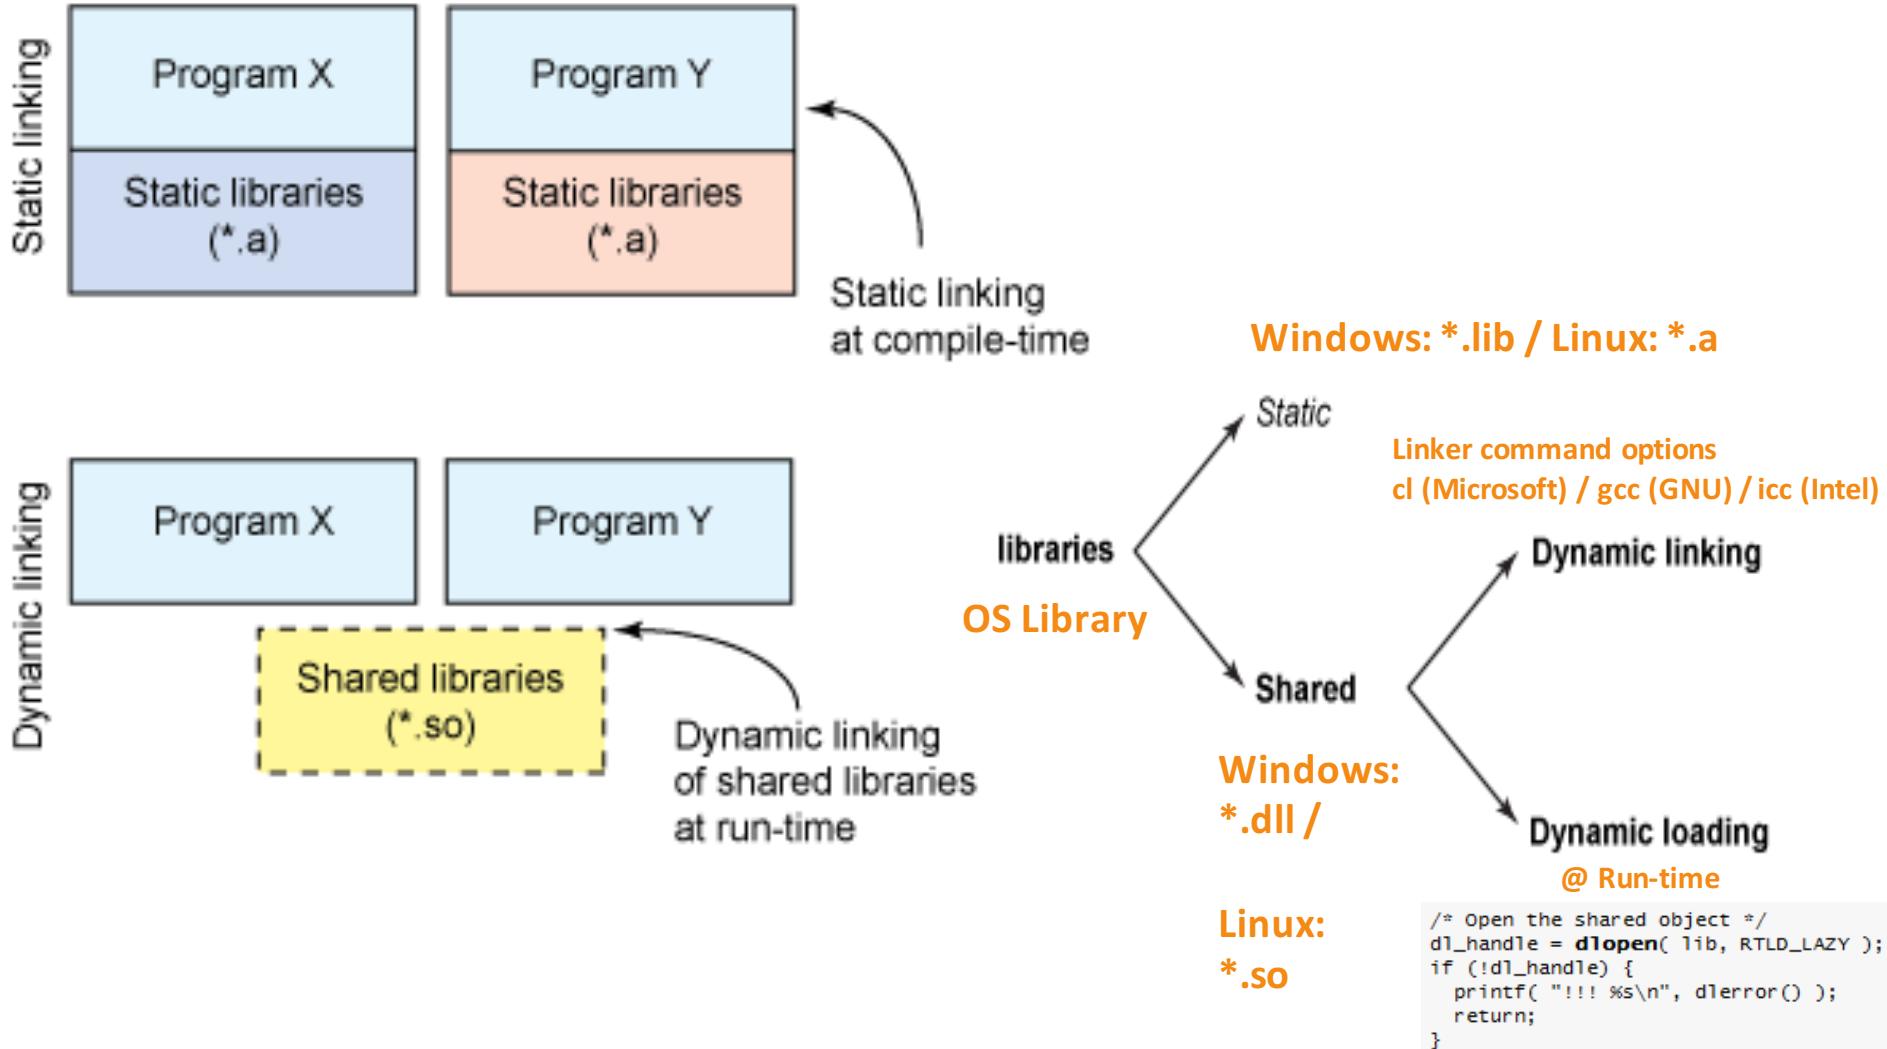
- Non-static fields? Static fields?
- Transient fields?
- Private and public fields and/or methods?
- Prototype / signature of the methods and / or the implementation of the methods?

<http://www.javaworld.com/community/node/2915>

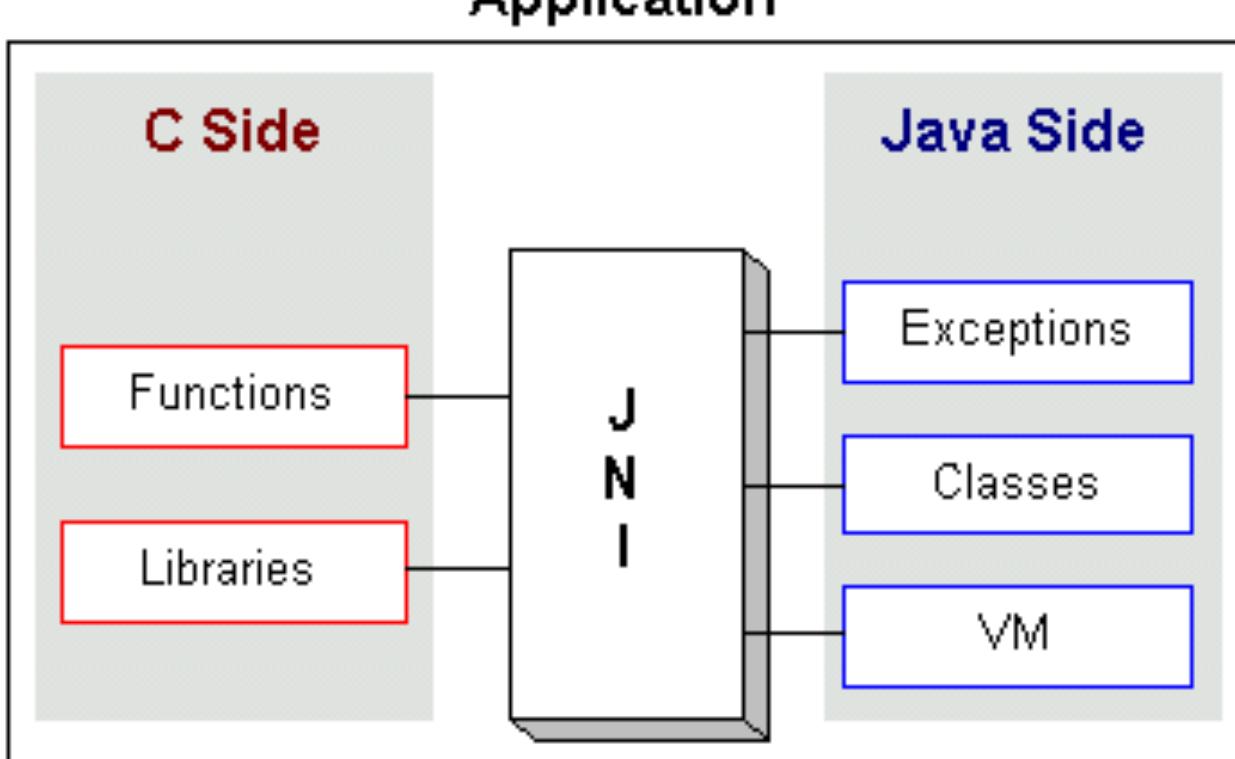


2.6 JNI – Linux vs. Win libraries

<http://www.ibm.com/developerworks/linux/library/l-dynamic-libraries/>

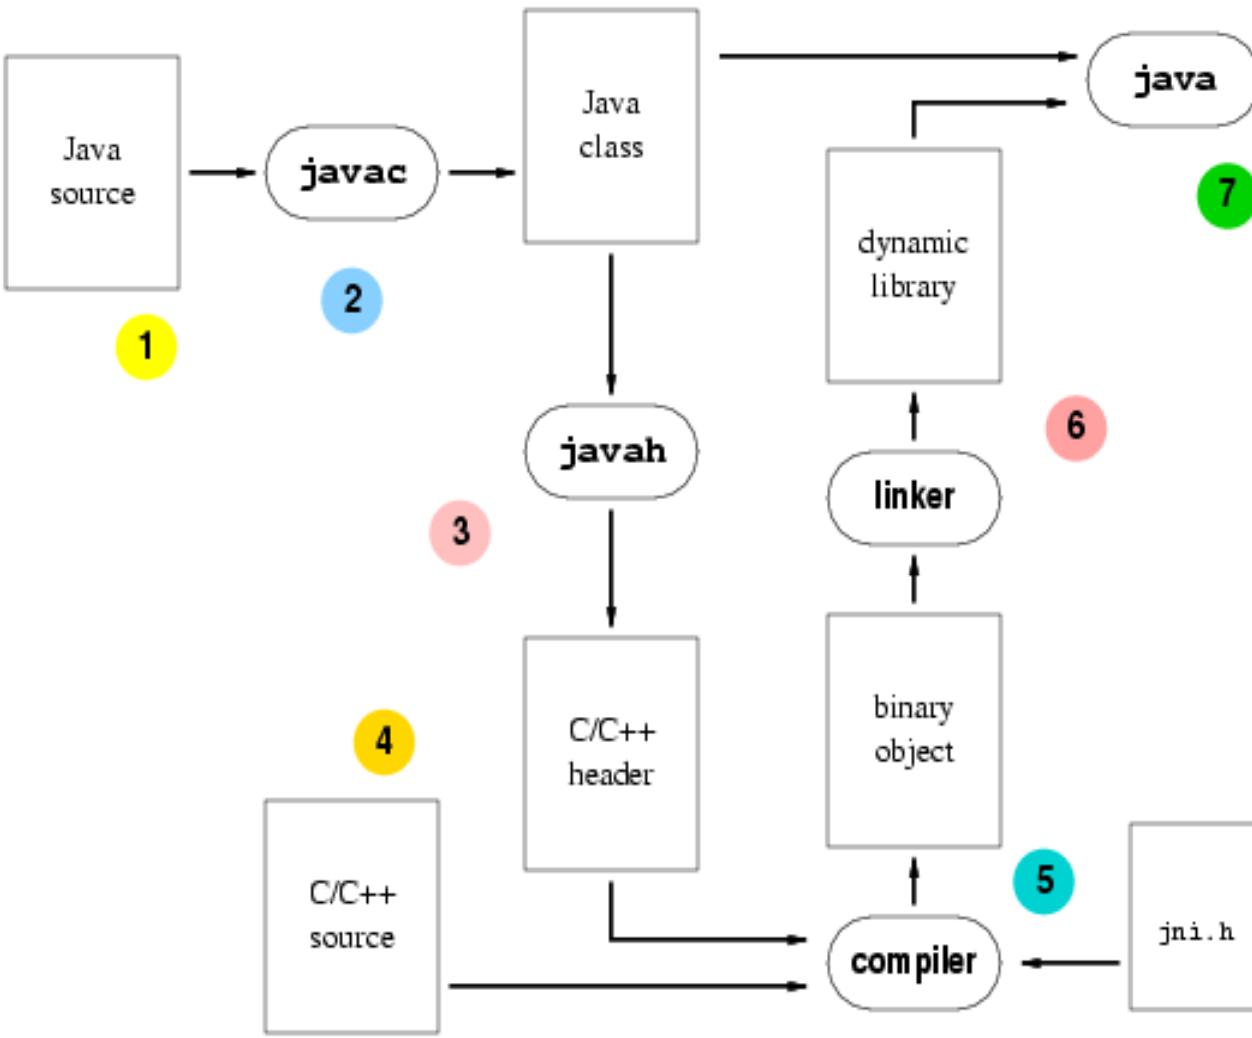


2.6 Java Native Interface



2.6 Java Native Interface

<http://cs.fit.edu/~ryan/java/language/jni.html>



1. Create Java source code with native methods

native *return type* method (*arguments*);

2. Compile Java source code and obtain the class files

3. Generate C/C++ headers for the native methods; `javah` gets the info it needs from the class files
4. Write the C/C++ source code for the native method using the function prototype from the generated include file and the typedefs from `include/jni.h`

5. Compile the C/C++ with the right header files

6. Use the linker to create a dynamic library file

7. Execute a Java program that loads the dynamic library

```
static {  
    System.loadLibrary("dynamic  
    library"); }
```

2.7 Summary of MS Windows Memory

Native EXE File on HDD
MS Windows:



EXE File Beginning – 'MZ'

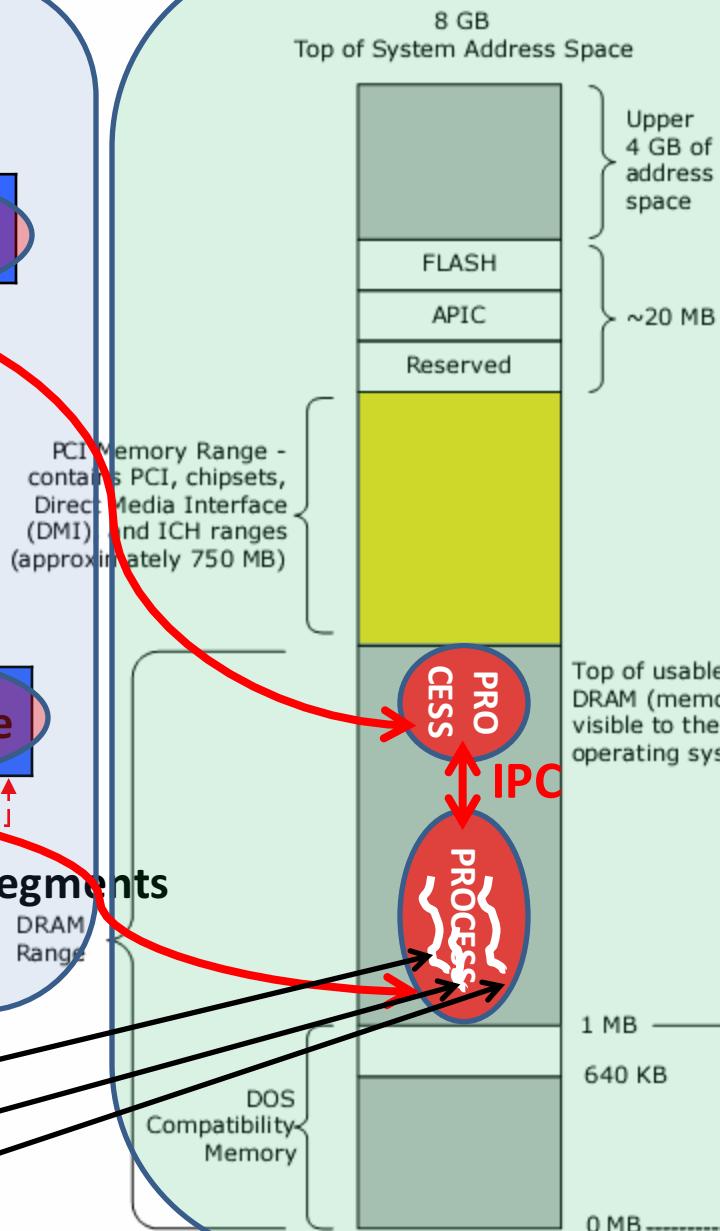
EXE 16, 32 bits Headers



References / pointers to the segments

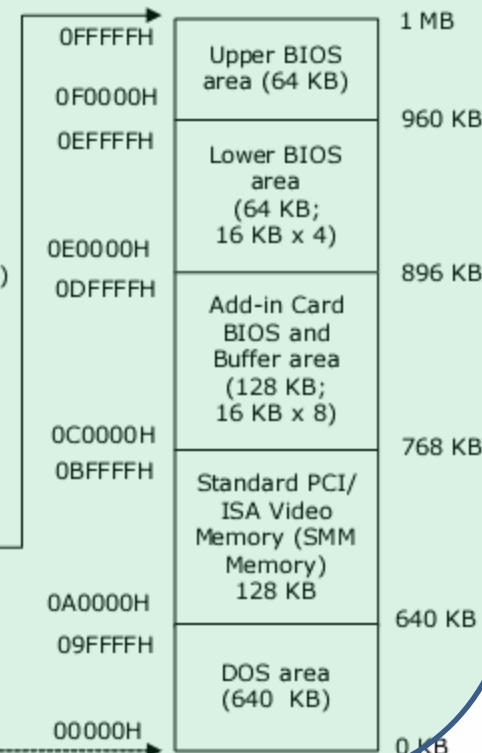
Relocation Pointer Table

Optional – Thread 1
Optional – Thread 2
... Optional – Thread n

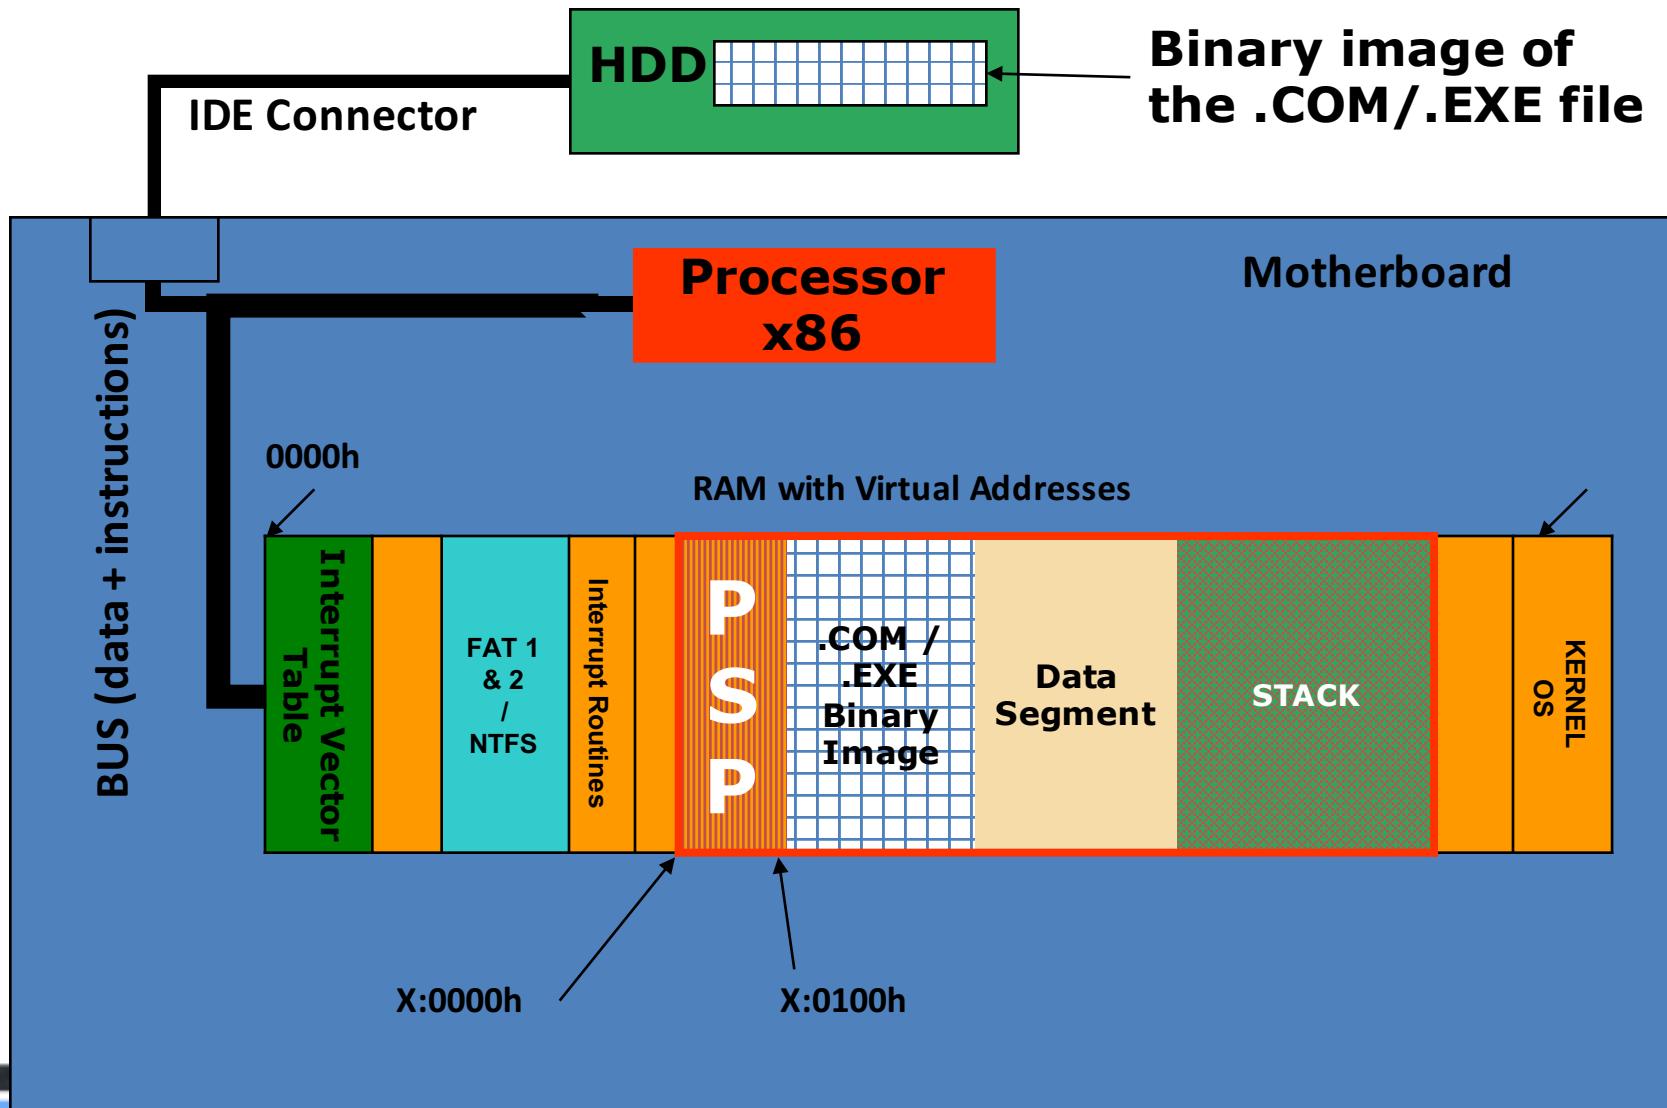


RAM Memory Layout
MS Windows:

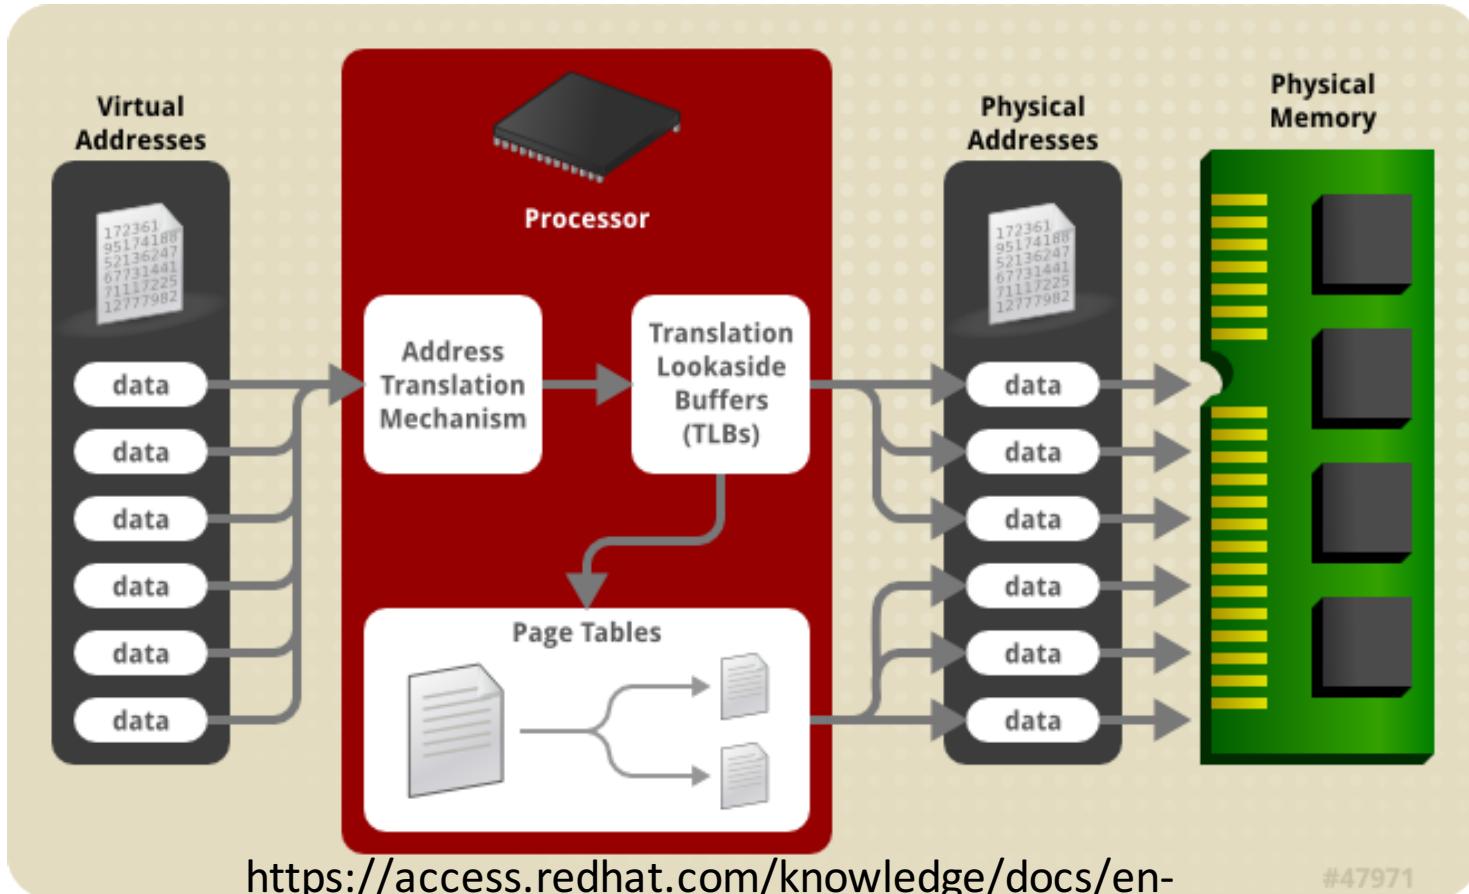
<http://www.codinghorror.com/blog/2007/03/where-where-my-4-gigabytes-of-ram.html>



2.7 Summary of MS Windows Process



2.7 Summary of Linux/Windows Virtual Memory

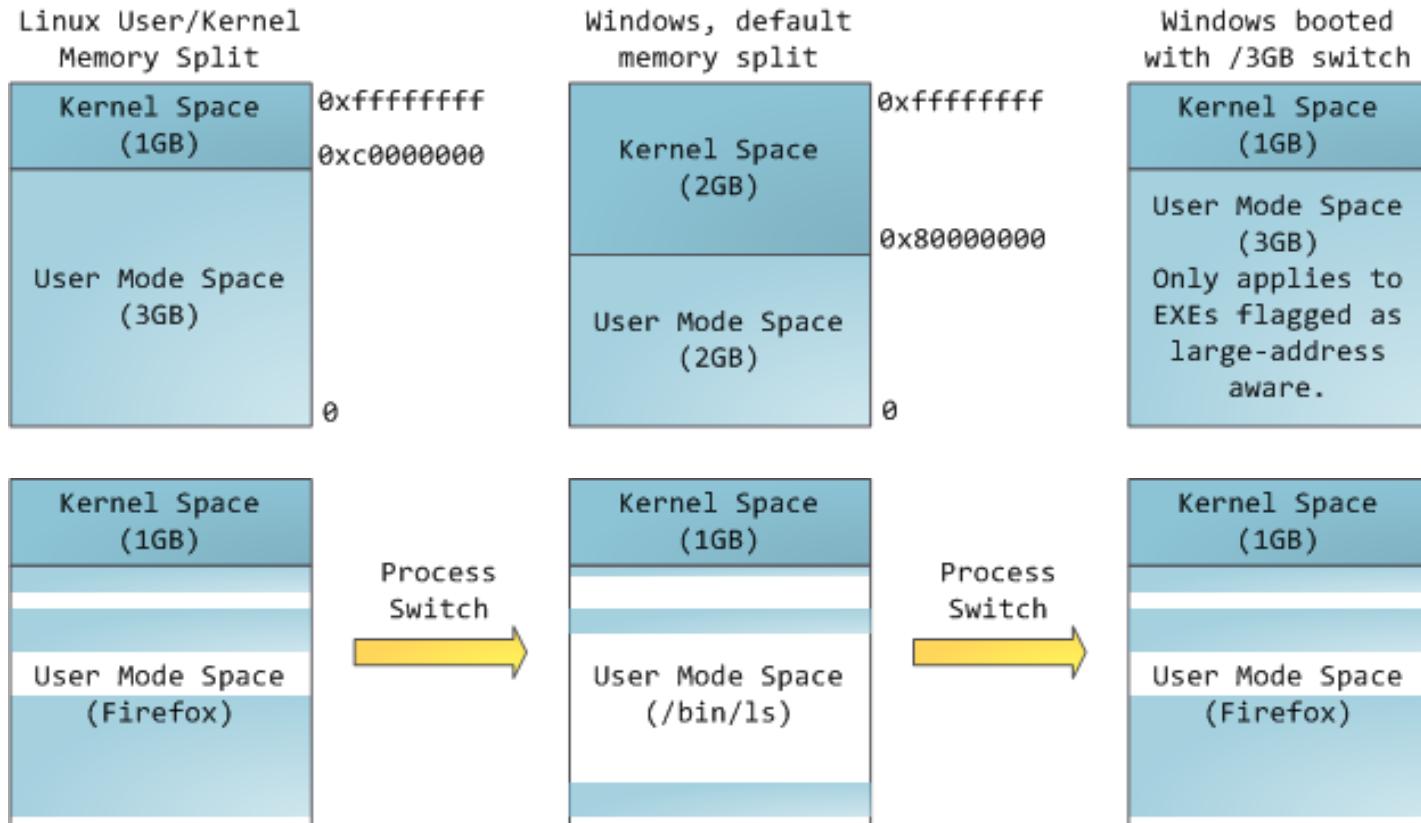


https://access.redhat.com/knowledge/docs/en-US/Red_Hat_Enterprise_MR...

#47971

US/Red_Hat_Enterprise_MR/2/html/Realtime_Reference_Guide/chap-
Realtime_Reference_Guide-Memory_allocation.html

2.7 Summary of Linux/Windows Virtual Memory

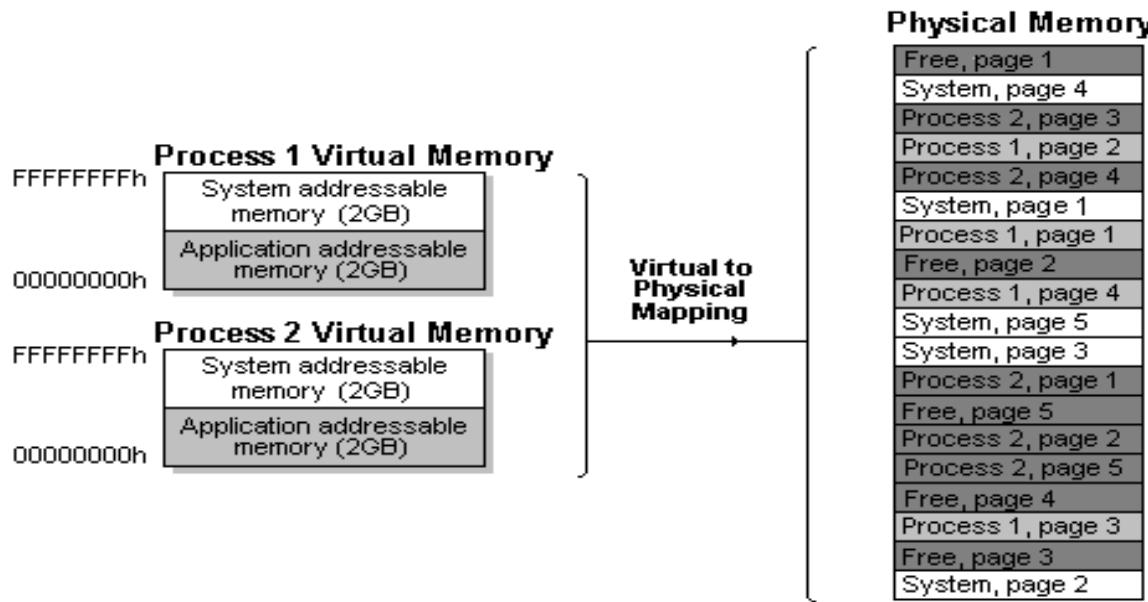


"Blue regions represent virtual addresses that are mapped to physical memory, whereas white regions are unmapped. In the example above, Firefox has used far more of its virtual address space due to its legendary memory hunger. The distinct bands in the address space correspond to **memory segments** like the heap, stack, and so on. Keep in mind these segments are simply a range of memory addresses and *have nothing to do with Intel-style segments*."

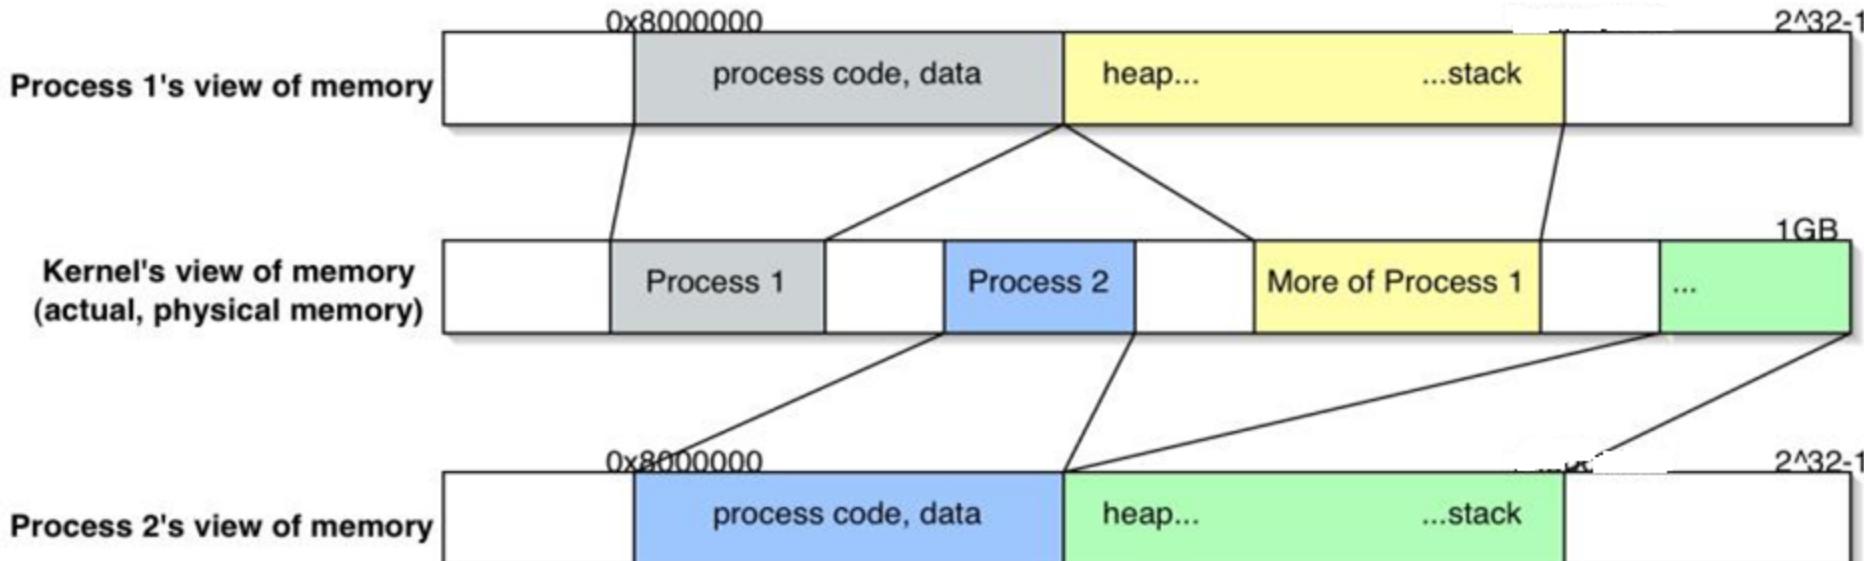
2.7 Summary of Linux/Windows Virtual Memory

MS Windows:

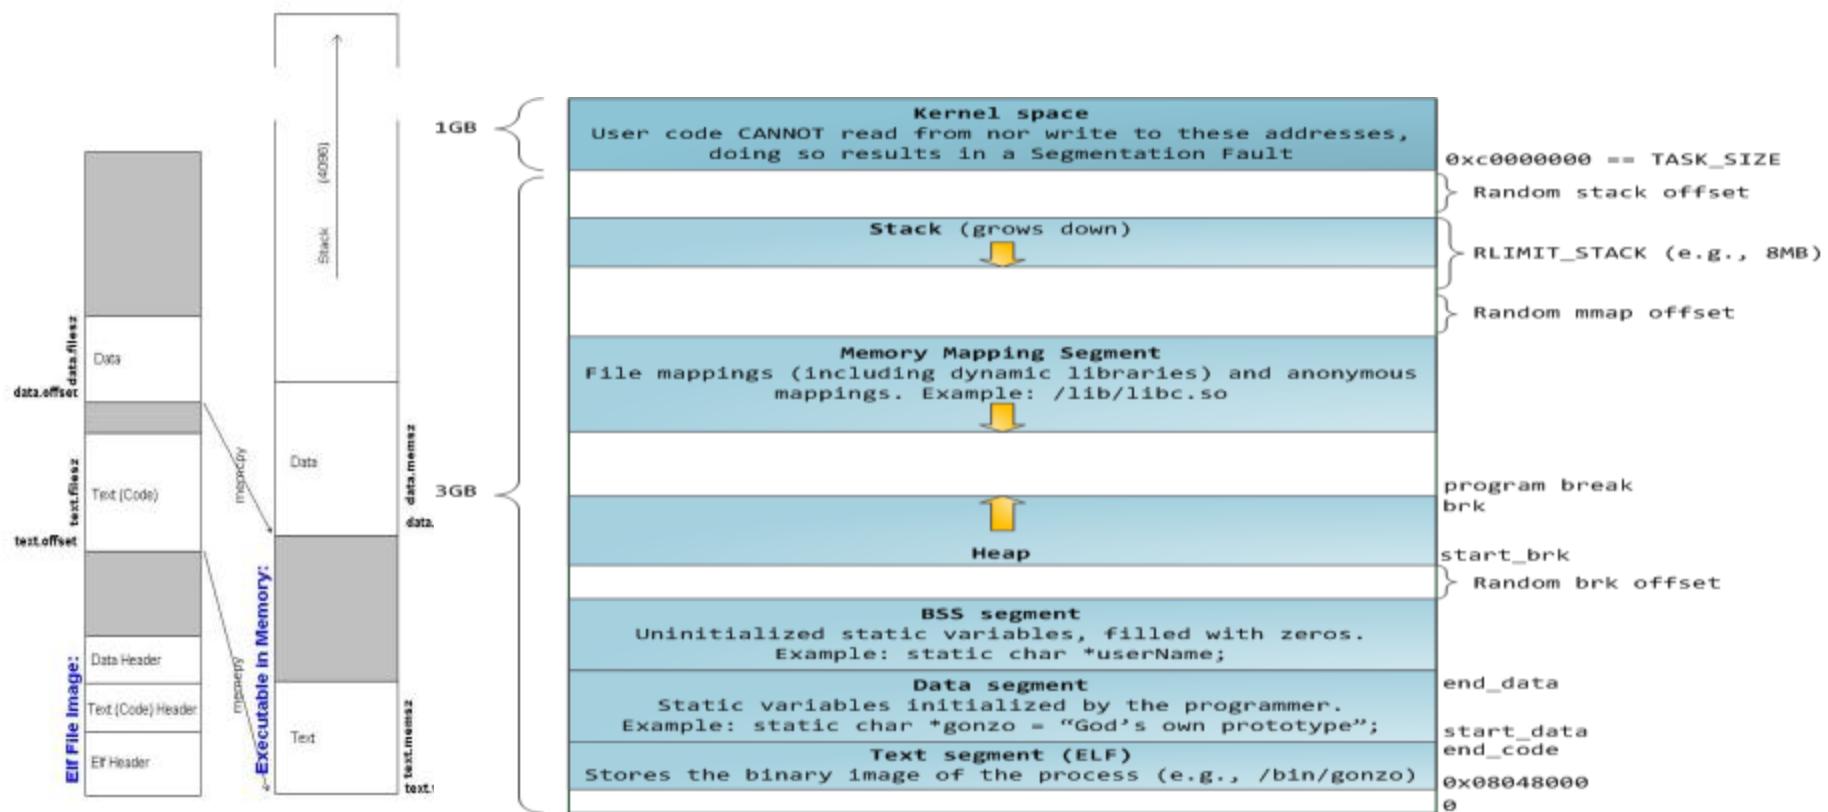
<http://technet.microsoft.com/en-us/library/cc751283.aspx>



LINUX: <http://www.read.cs.ucla.edu/111/2007fall/notes/lec4>



2.7 Summary of Linux executable ELF to memory - Process



<http://www.cs.umd.edu/~hollings/cs412/s04/proj1/index.html#cast>

2.7 Summary of Processes & IPC in Linux

<http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>

<https://computing.llnl.gov/tutorials/pthreads/>

<http://www.advancedlinuxprogramming.com/alp-folder/>

Before understanding a thread, one first needs to understand a UNIX process. A process is created by the operating system, and requires a fair amount of "overhead".

Processes contain information about program resources & program execution state, including:

- *Process ID, process group ID, user ID, and group ID;*
- *Environment;*
- *Working directory;*
- *Program instructions;*
- *Registers;*
- *Stack;*
- *Heap;*
- *File descriptors;*
- *Signal actions;*
- *Shared libraries;*
- *Inter-process communication tools (such as message queues, pipes, semaphores, or shared memory)*

2.7 Summary of Processes & IPC in Linux

Processes

- Fork
- Signals

Pipes

FIFO

File-locking

OS Message
Queues

Semaphores

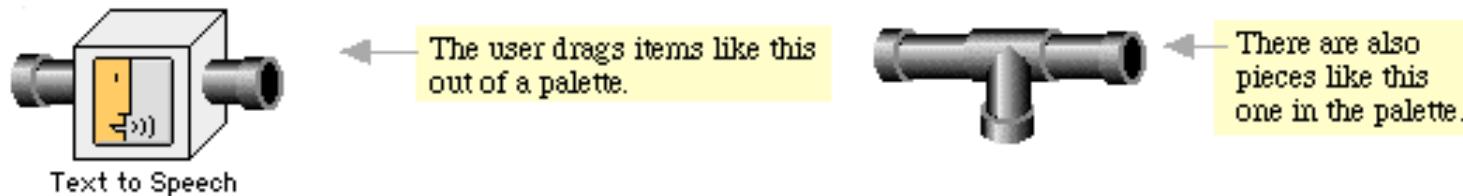
Shared
Memory

Memory
Mapped Files

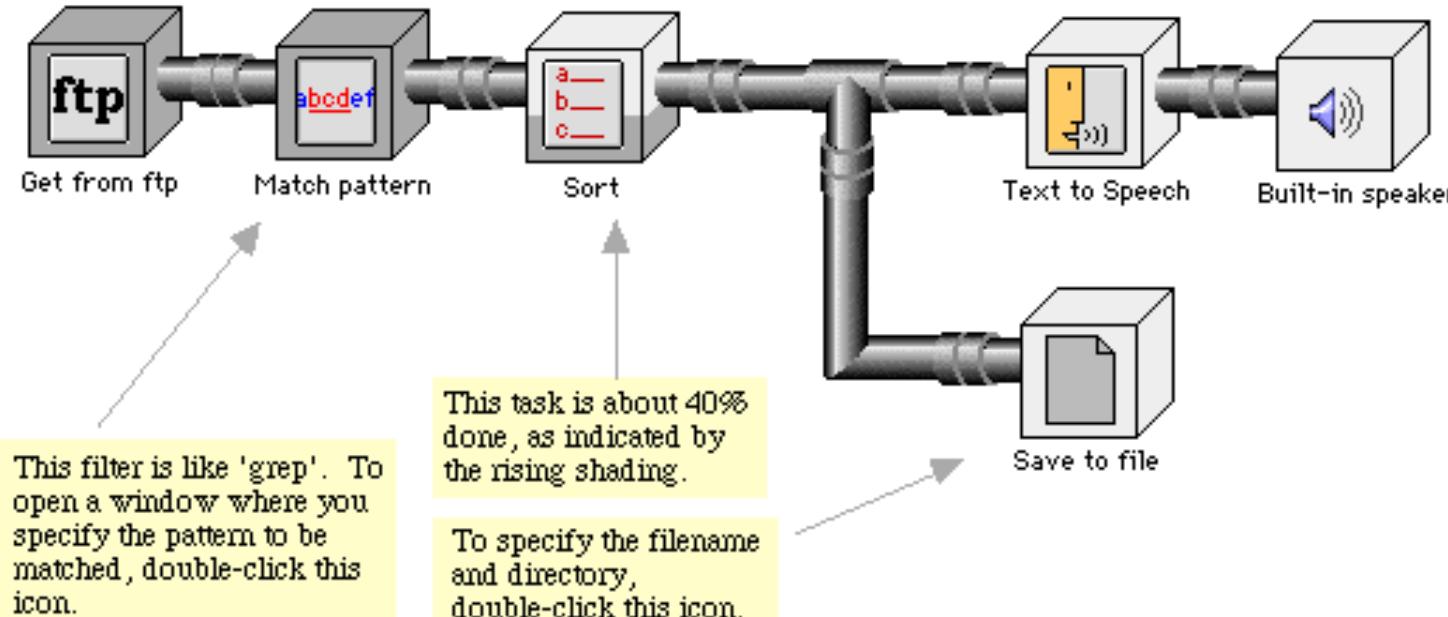
Sockets

2.7 Summary of IPC in Linux – Why Pipes?

http://www.sean-crist.com/personal/pages/visual_pipes/



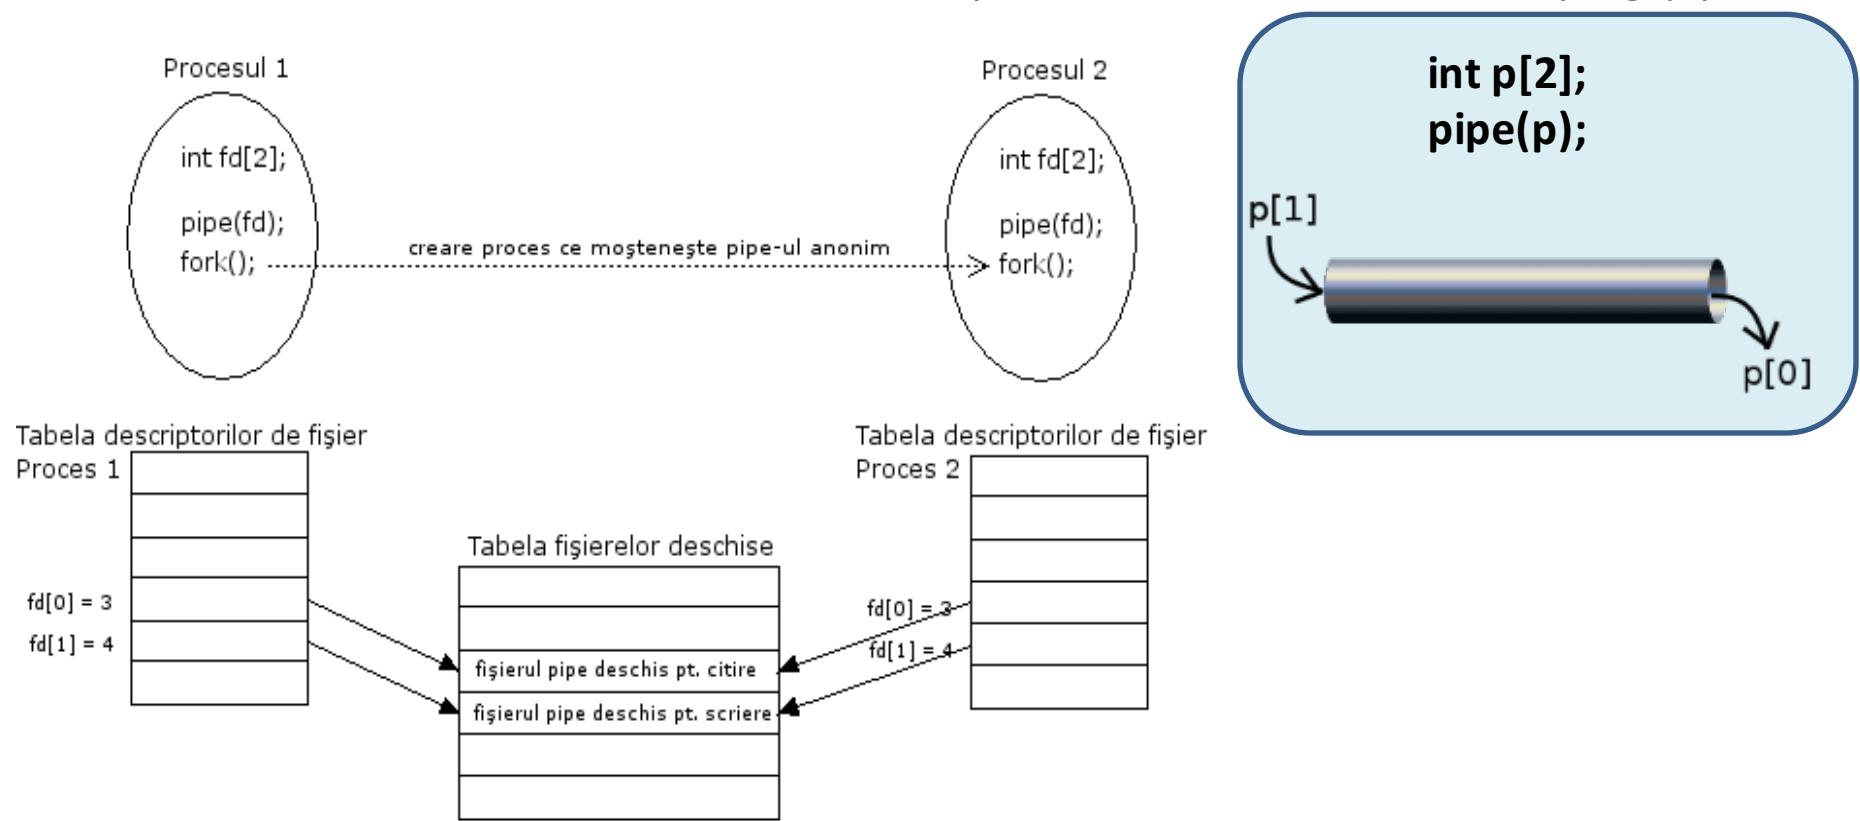
In the example below, the user has instructed the computer: 1) to get a file thru ftp; 2) to select just those lines matching a certain pattern; 3) to sort the results; and 4) to both save the results to file and also read them thru the loudspeaker.



Many people have observed that Linux is difficult for casual users to learn, and that Linux would have a better chance of general acceptance as a desktop platform if it were made easier to use. Pipes are at the root of the great flexibility of Unix, and representing them graphically makes this functionality better accessible to the casual user.

2.7 Summary of IPC in Linux – Fork & Pipes

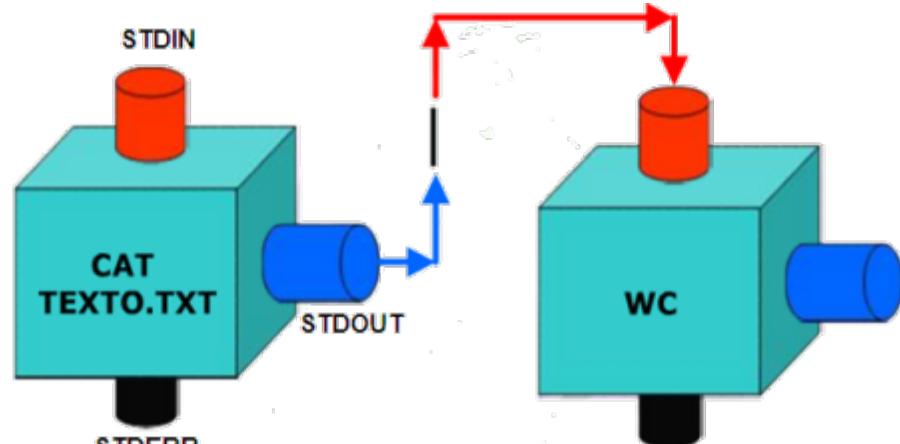
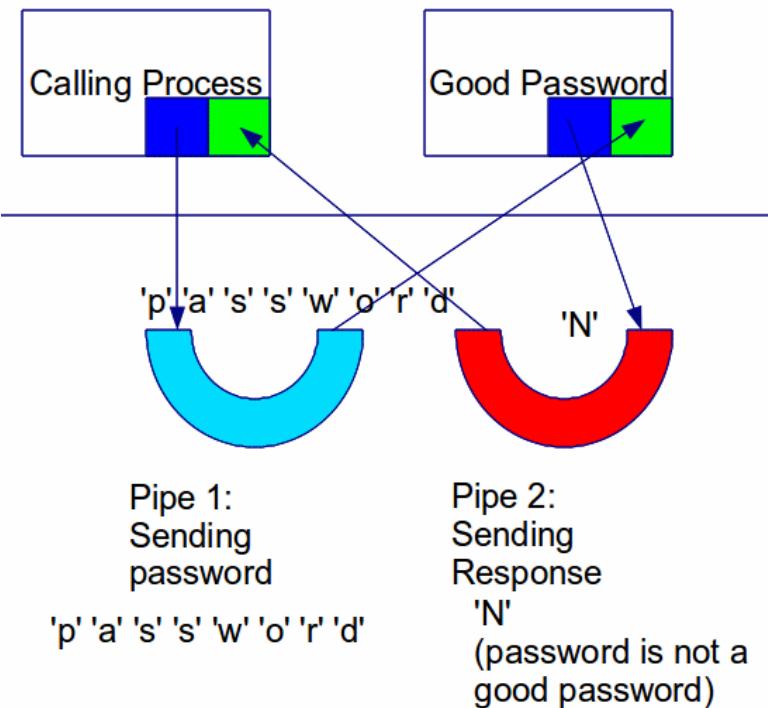
<http://www.reloco.com.ar/linux/prog/pipes.html>



<http://os.obs.utcluj.ro/OS/Lab/08.Linux%20Pipes.html>

2.7 Summary of IPC in Linux – Fork & Pipes

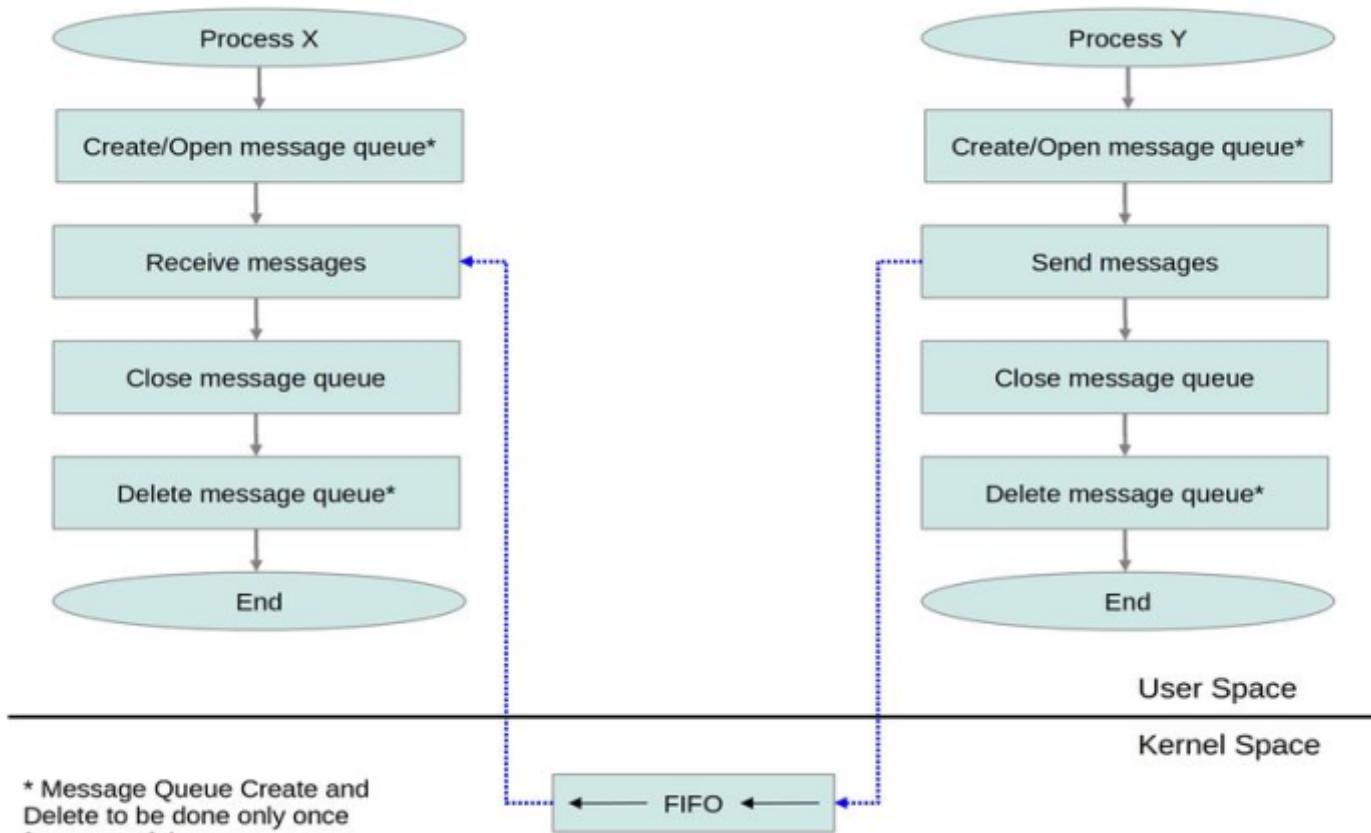
<http://www.vivaolinux.com.br/dica/Pipes-no-Linux>



http://www.read.cs.ucla.edu/111/_media/notes/ipc_pipes_1.gif

2.7 Summary of IPC in Linux – Message Queues

Linux C - System V API / POSIX API



http://www.linuxpedia.org/index.php?title=Linux_POSIX_Message_Queue

2.7 Summary of IPC in Linux – Message Queues

Linux C - System V API / POSIX API

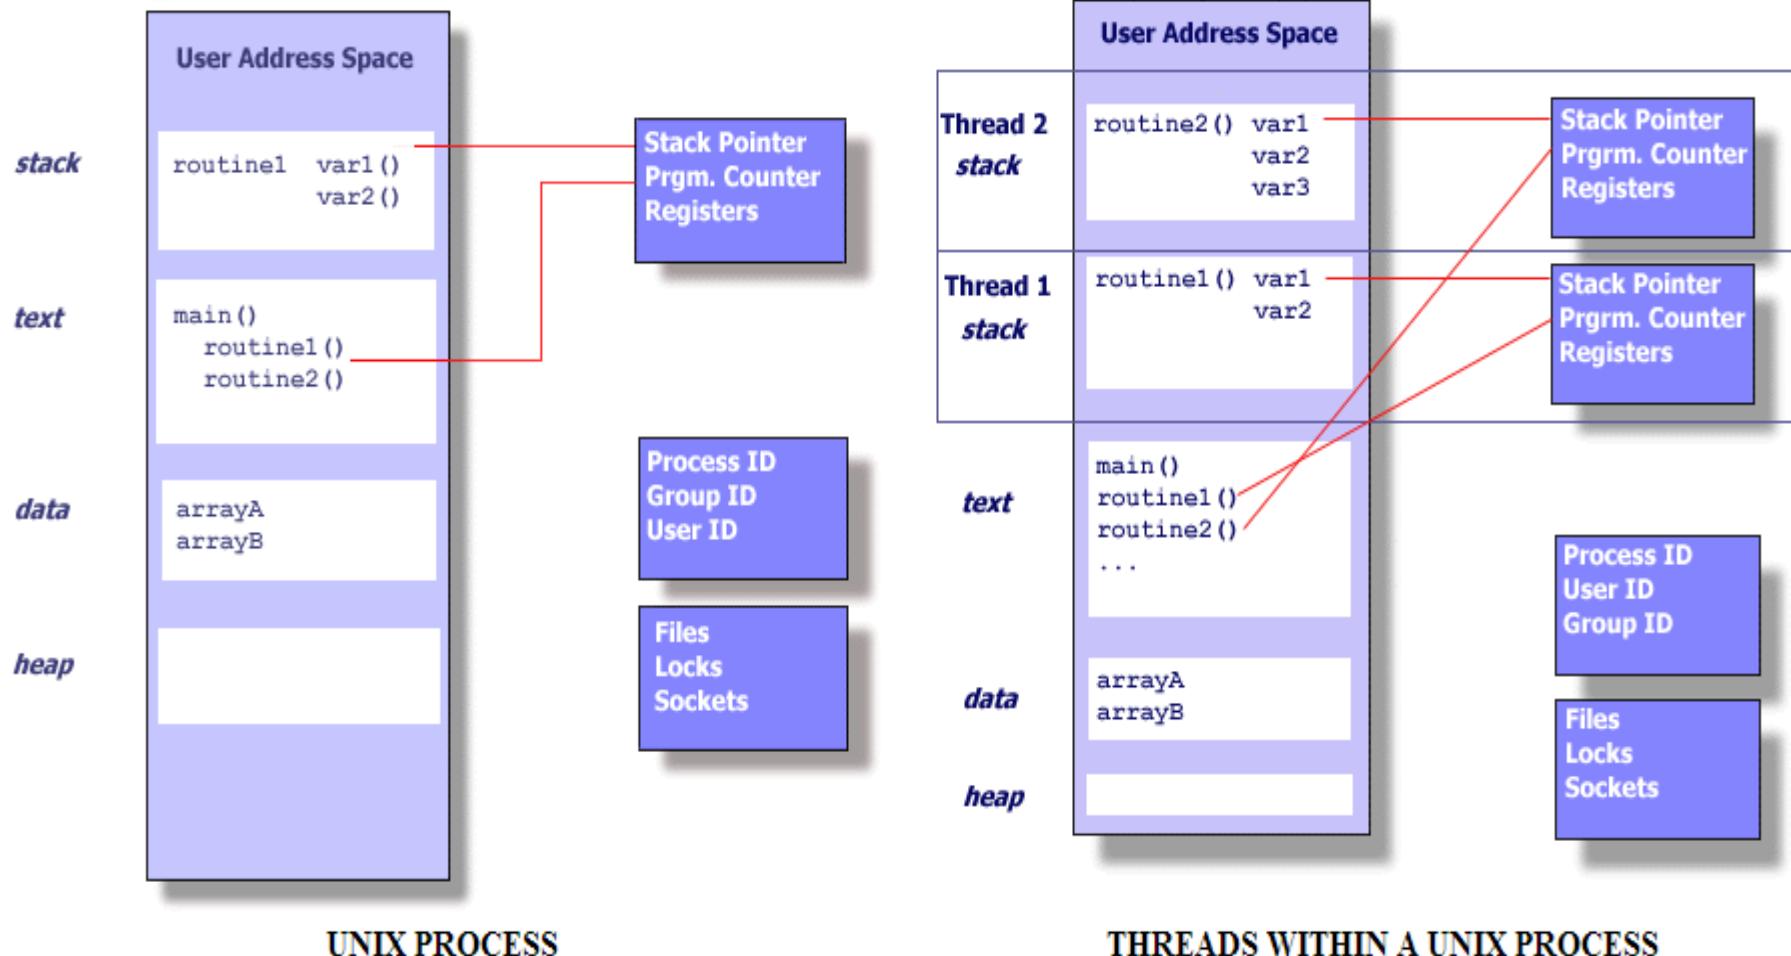
Operation	POSIX Function	SVR4 Function
Gain access to a queue, creating it if it does not exist.	mq_open(3)	msgget(2)
Query attributes of a queue and number of pending messages.	mq_getattr(3)	msgctl(2)
Change attributes of a queue.	mq_setattr(3)	msgctl(2)
Give up access to a queue.	mq_close(3)	n.a.
Remove a queue from the system.	mq_unlink(3), rm(1)	msgctl(2), ipcrm(1)
Send a message to a queue.	mq_send(3)	msgsnd(2)
Receive a message from a queue.	mq_receive(3)	msgrcv(2)
Request asynchronous notification of a message arriving at a queue.	mq_notify(3)	NA

http://menehune.opt.wfu.edu/Kokua/More_SGI/007-2478-008/sgi_html/ch06.html

http://www.users.pjwstk.edu.pl/~jms/qnx/help/watcom/clibref/mq_overview.html

2.8 Summary of Multi-threading in C vs. Java

Multi-threading vs. Multi-process development in UNIX/Linux:



2.8 Summary of Multi-threading in C vs. Java

Threads Features:

Thread operations include thread creation, termination, synchronization (joins, blocking), scheduling, data management and process interaction.

A thread does not maintain a list of created threads, nor does it know the thread that created it.

All threads within a process share the same address space.

Threads in the same process share:

- Process instructions
- Most data
- open files (descriptors)
- signals and signal handlers
- current working directory
- User and group id

Each thread has a unique:

- Thread ID
- set of registers, stack pointer
- stack for local variables, return addresses
- signal mask
- priority
- Return value: errno

pthread functions return "0" if OK.

2.8 Summary of Multi-threading in C vs. Java

Multi-threading in C/C++ with pthread (Is “counter++” an atomic operation?):

Without Mutex	With Mutex
<pre>1 int counter=0; 2 3 /* Function C */ 4 void functionC() { 5 6 counter++ 7 8 }</pre>	<pre>01 /* Note scope of variable and mutex are the same */ 02 03 pthread_mutex_t mutex1 = 04 PTHREAD_MUTEX_INITIALIZER; 05 int counter=0; 06 07 /* Function C */ 08 void functionC() 09 { 10 pthread_mutex_lock(&mutex1); 11 counter++; 12 pthread_mutex_unlock(&mutex1); 13 }</pre>

Possible execution sequence

Thread 1	Thread 2	Thread 1	Thread 2
counter = 0	counter = 0	counter = 0	counter = 0
counter = 1	counter = 1	counter = 1	Thread 2 locked out. Thread 1 has exclusive use of variable counter
			counter = 2

2.8 Summary of Multi-threading in C vs. Java

Multi-threading vs. Multi-process mini-terms:

Mutexes are used to prevent data inconsistencies due to race conditions.

A race condition often occurs when two or more threads need to perform operations on the same memory area, but the results of computations depends on the order in which these operations are performed.

Mutexes are used for serializing shared resources. Anytime a global resource is accessed by more than one thread the resource should have a Mutex associated with it.

One can apply a **mutex** to protect a segment of memory ("critical region") from other threads.

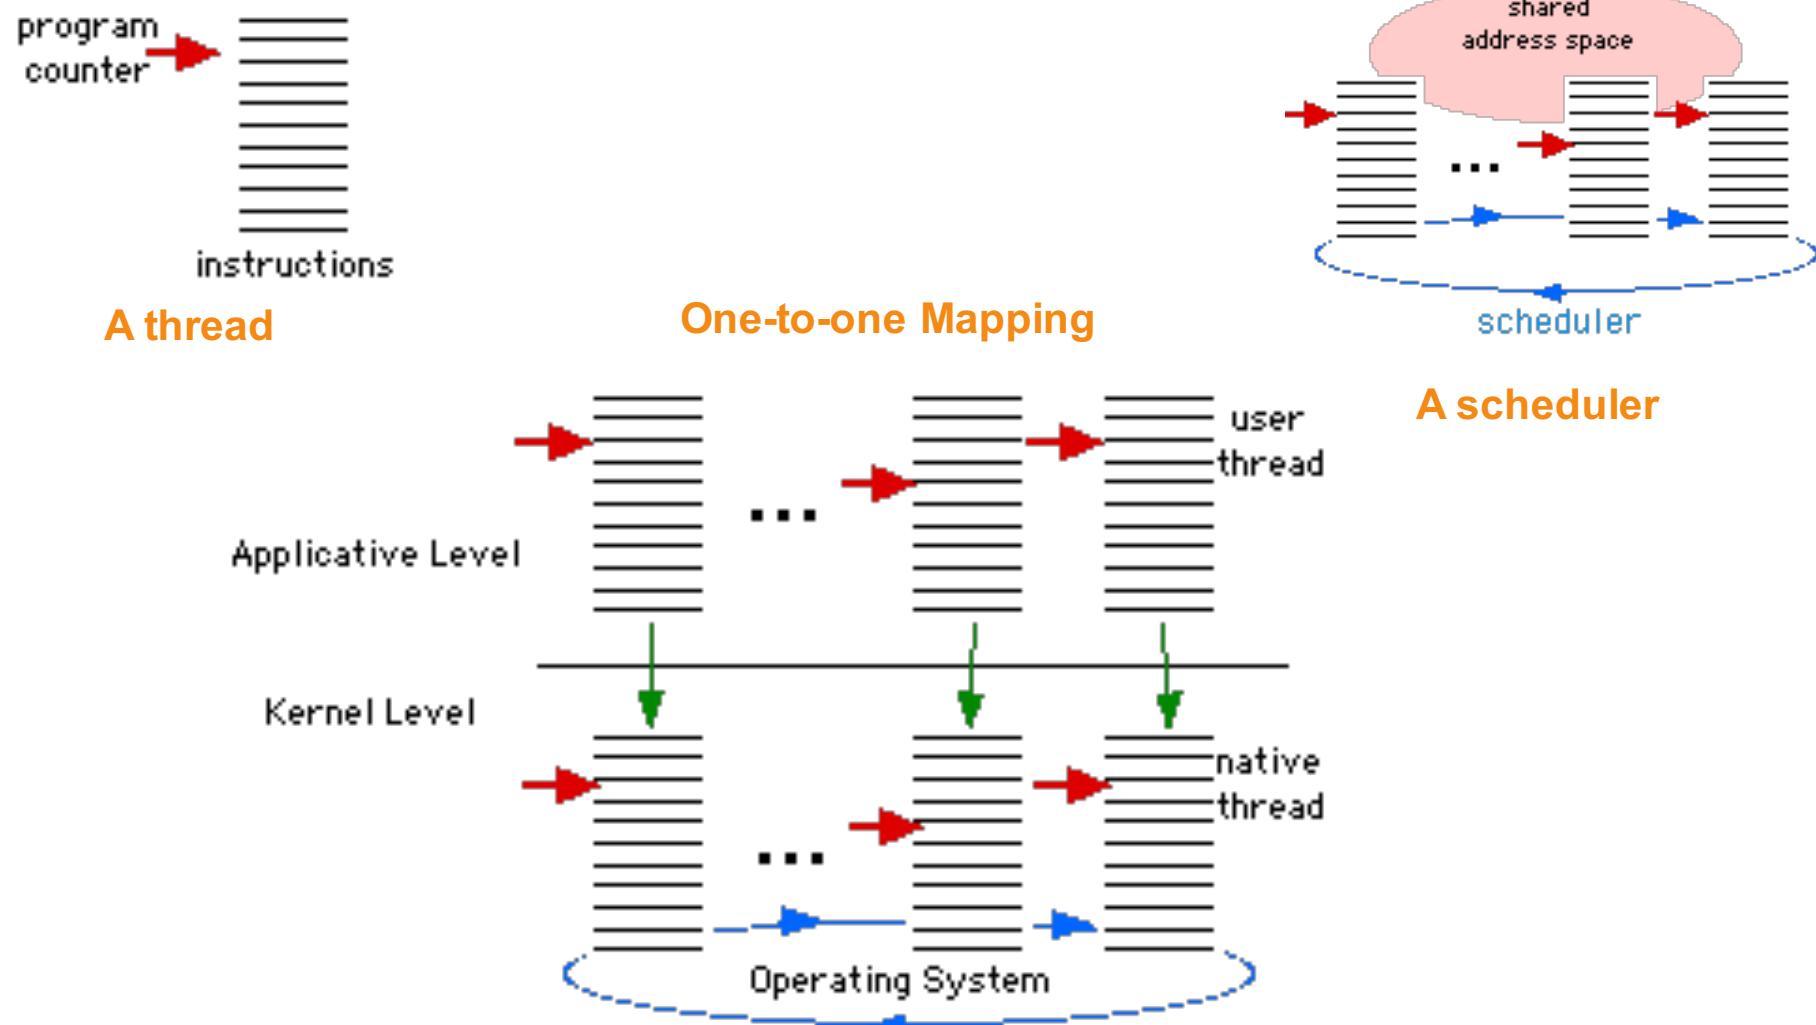
Mutexes can be applied only to threads in a single process and do not work between processes as do **semaphores**.

In Java **Mutex** is quite \Leftrightarrow **synchronized**

2.8 Summary of Multi-threading in C vs. Java

Multi-threading Models:

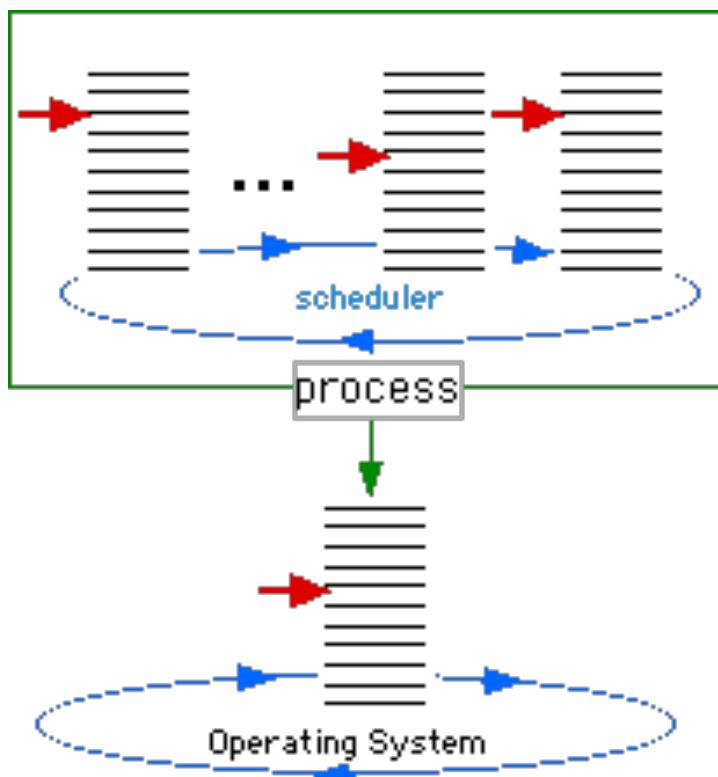
<http://www-sop.inria.fr/inde/rp/FairThreads/FTJava/documentation/FairThreads.html#One-one-mapping>



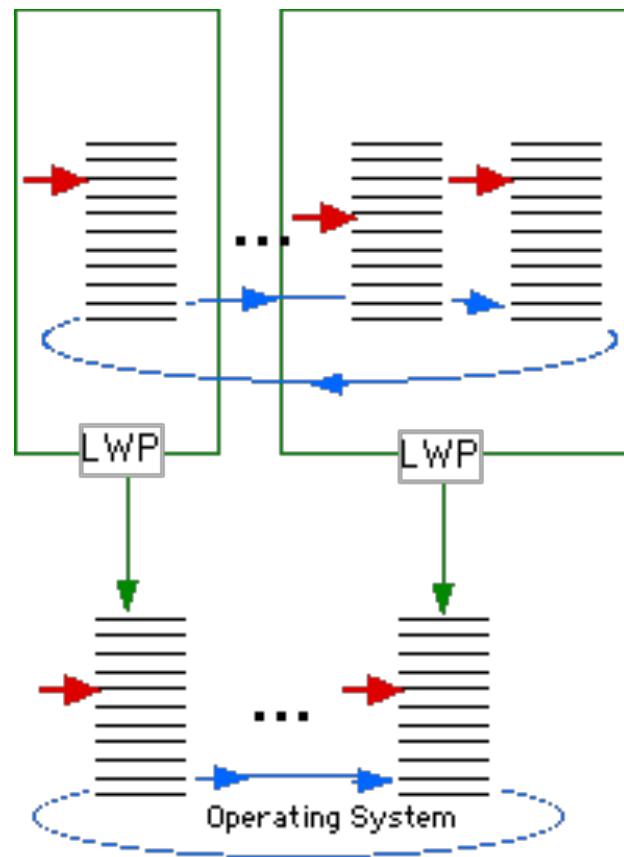
2.8 Summary of Multi-threading in C vs. Java

Multi-threading Models:

<http://www-sop.inria.fr/inde/rp/FairThreads/FTJava/documentation/FairThreads.html#One-one-mapping>



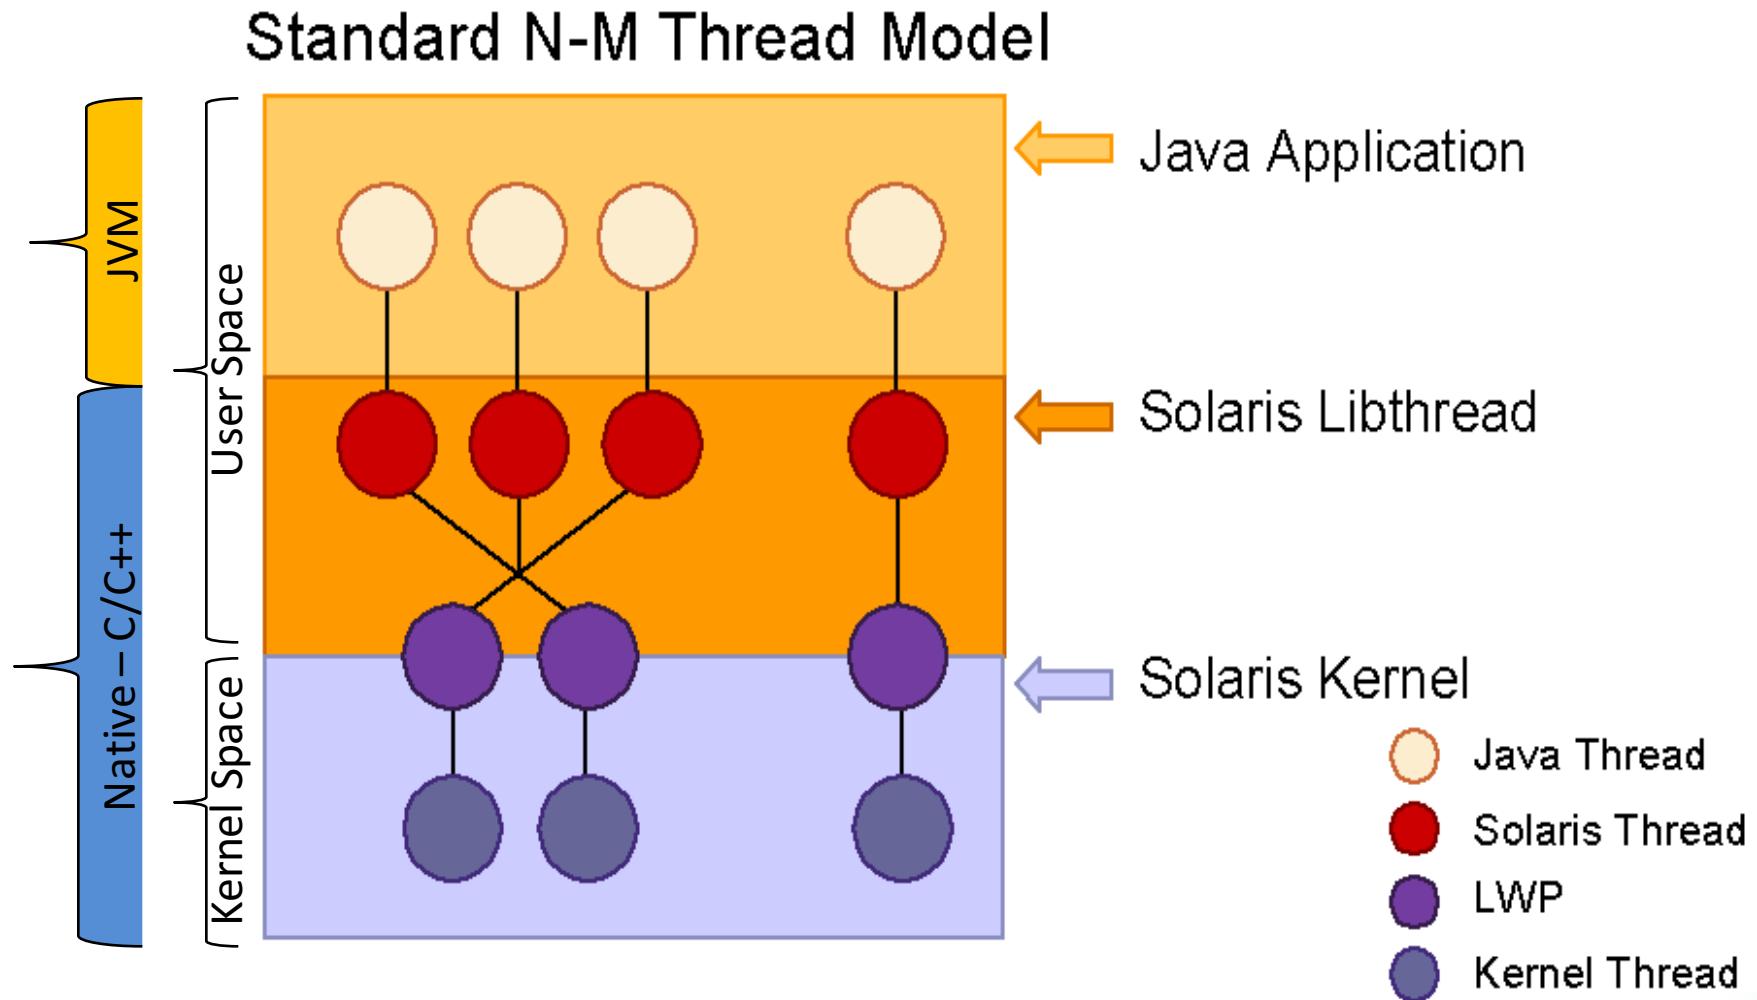
Many-to-one Mapping



Many-to-many Mapping

2.8 Summary of Multi-threading in C vs. Java

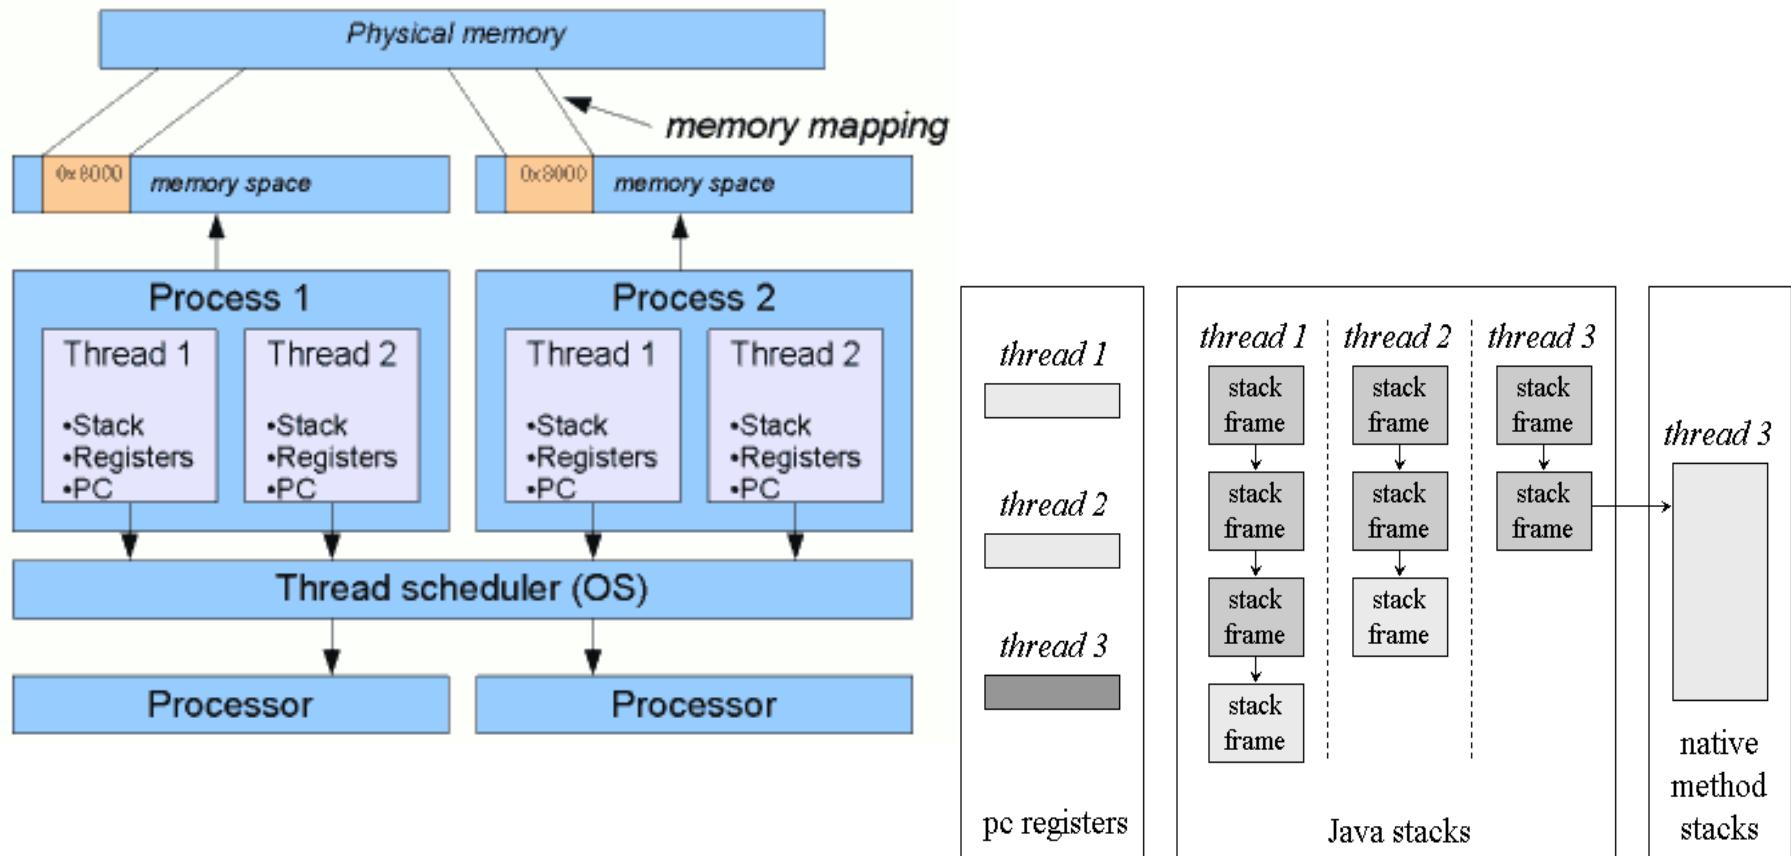
JVM Multi-threading Model mapping to OS native threads:



2.8 Summary of Multi-threading in C vs. Java

JVM Multi-threading Model mapping to OS native threads:

http://www.javamex.com/tutorials/threads/how_threads_work.shtml



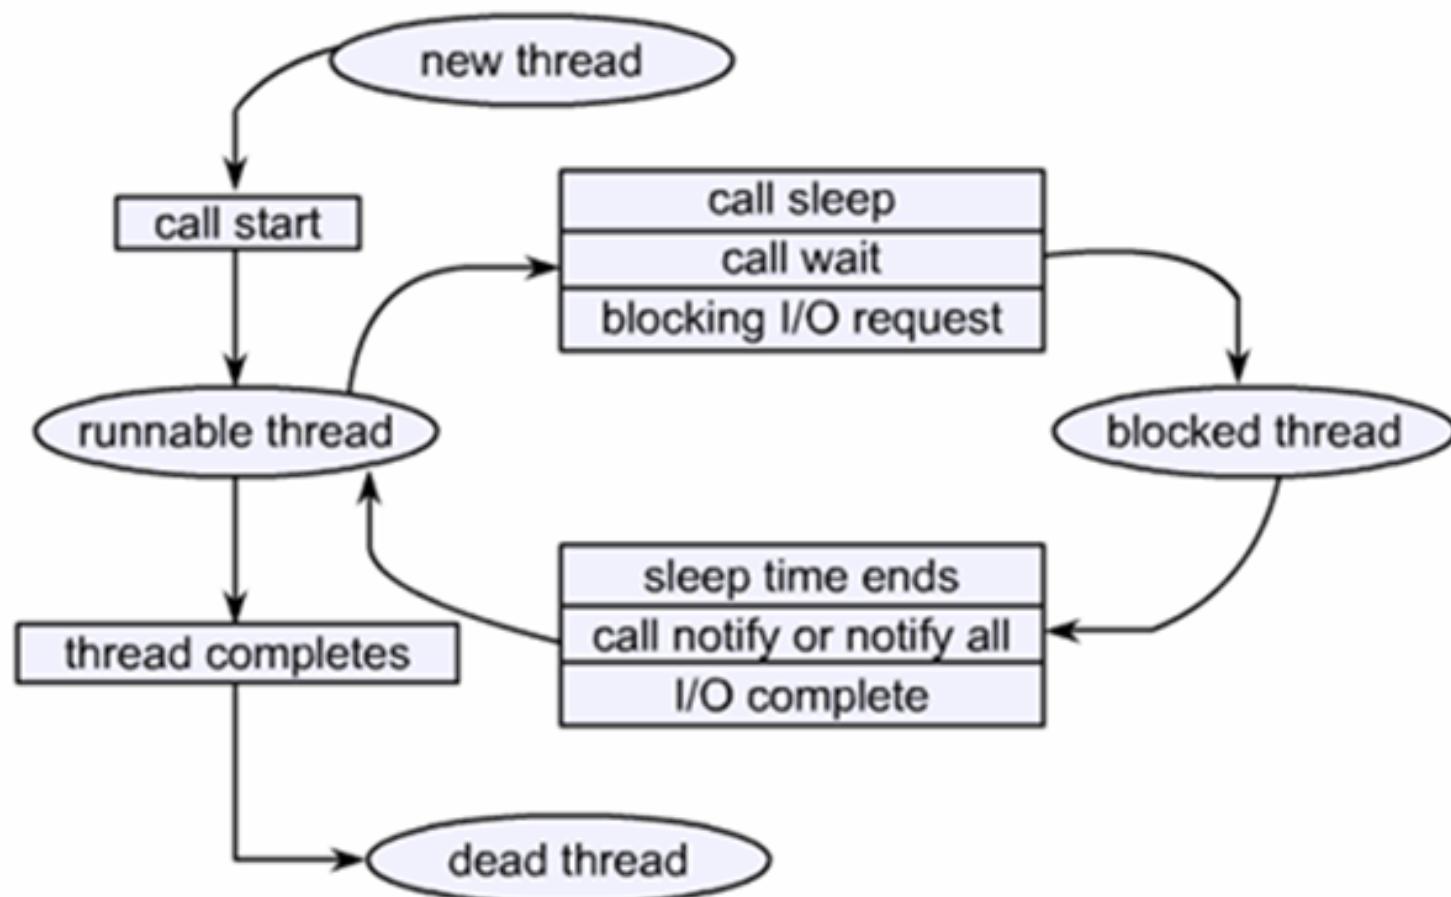
2.9 Summary of Multi-threading in Java

Java Multi-threading API

Option 1 – Java API for Multi-Threading Programming	Option 2 – Java API for Multi-Threading Programming
Defining the classes:	
class Fir extends Thread { public void run() {...} }	class Fir extends Ceva implements Runnable { public void run() {...} }
Instantiates the objects	
Fir f = new Fir();	Fir obf = new Fir(); Thread f = new Thread(obf);
Set the thread in ‘Runnable’ state	
f.start();	f.start();
Specific Thread methods calls	
public void run() { Thread.sleep(...); String fName = this.getName(); ... }	public void run() { Thread.sleep(...); Thread t = Thread.currentThread(); String fName = t.getName(); ... }

2.9 Summary of Multi-threading in Java

Java Thread States



Section Conclusions

All threads in a program must run the same executable. A child process, on the other hand, may run a different executable by calling an exec function.

An errant thread can harm other threads in the same process because threads share the same virtual memory space and other resources. For instance, a wild memory write through an uninitialized pointer in one thread can corrupt memory visible to another thread.

An errant process, on the other hand, cannot do so because each process has a copy of the program's memory space.

Copying memory for a new process adds an additional performance overhead relative to creating a new thread. However, the copy is performed only when the memory is changed, so the penalty is minimal if the child process only reads memory.

Threads should be used for programs that need fine-grained parallelism. For example, if a problem can be broken into multiple, nearly identical tasks, threads may be a good choice. Processes should be used for programs that need coarser parallelism.

Sharing data among threads is trivial because threads share the same memory. However, great care must be taken to avoid race conditions. Sharing data among processes requires the use of IPC mechanisms. This can be more cumbersome but makes multiple processes less likely to suffer from concurrency bugs.

Multi-threading & IPC Summary
for easy sharing



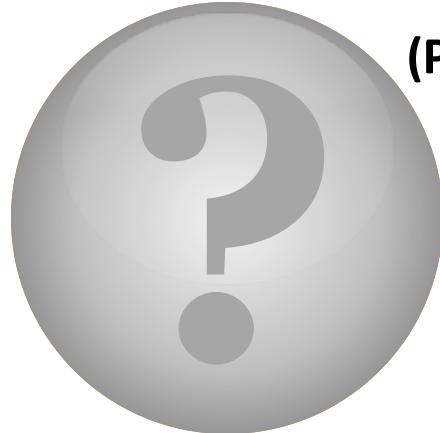
Share knowledge, Empowering Minds

Communicate & Exchange Ideas



Some “myths”:

(Distributed Systems).Equals(Distributed Computing) == true?



(Parallel System).Equals(Parallel Computing) == true?

(Parallel System == Distributed System) != true?

**(Sequential vs. Parallel vs. Concurrent vs.
Distributed Programming) ?(Different) : (Same)**

Questions & Answers!

But wait...

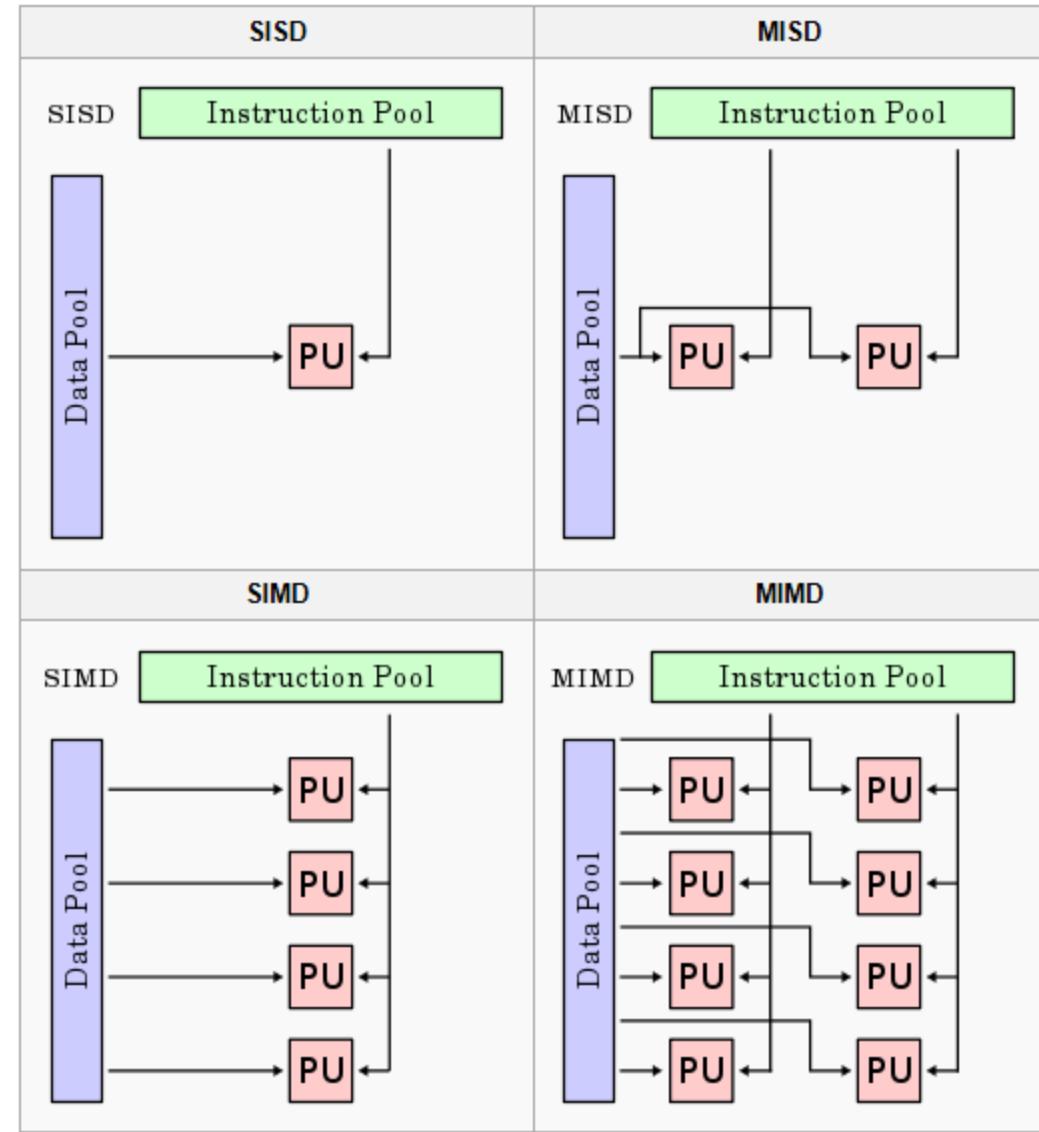
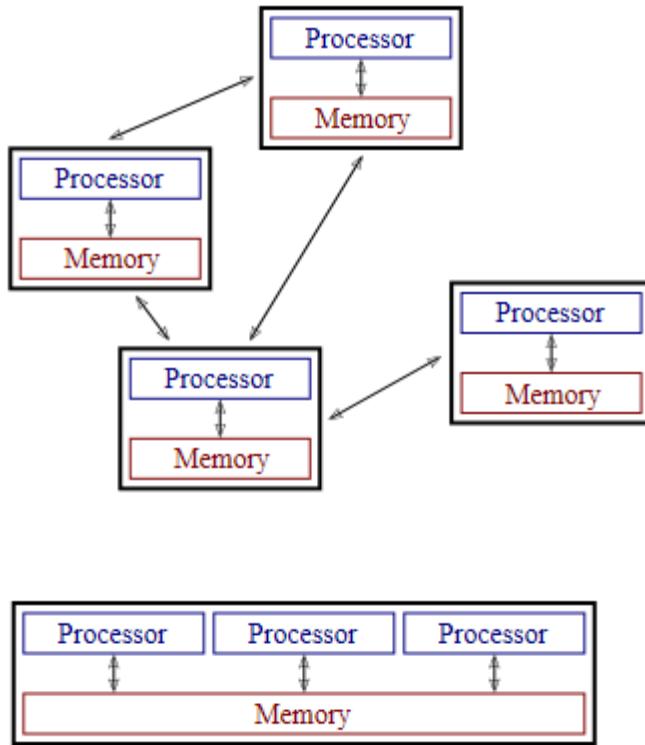
There's More!

if (HTC != HPC)
 HTC (High Throughput Computing) >
 MTC (Many Task Computing) >
 HPC (High Performance Computing);

... Will be continued! - In the next lectures ...

Flynn Taxonomy Parallel vs. Distributed Systems

Parallel vs. Distributed Computing / Algorithms



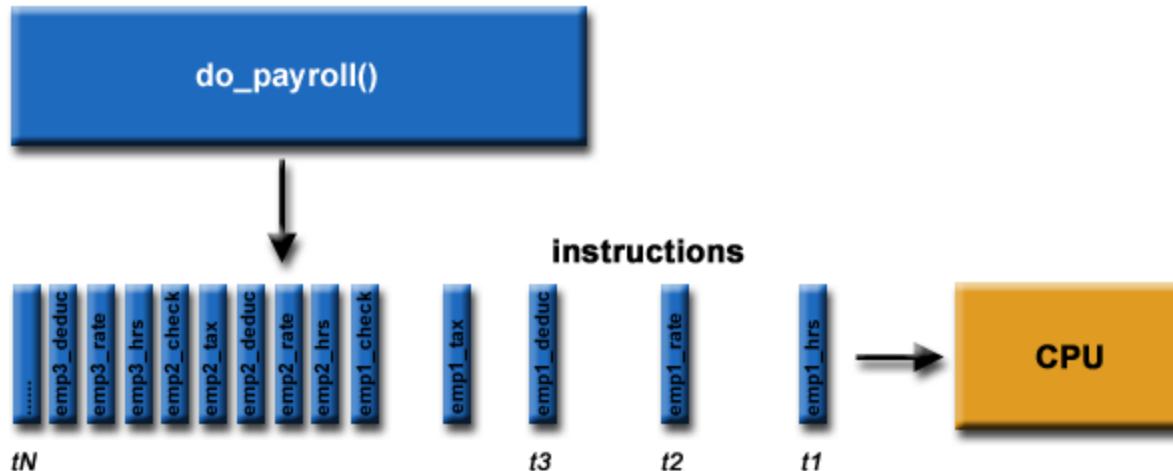
Where is the picture for:
Distributed System and **Parallel System**?

http://en.wikipedia.org/wiki/Distributed_computing
http://en.wikipedia.org/wiki/Flynn's_taxonomy

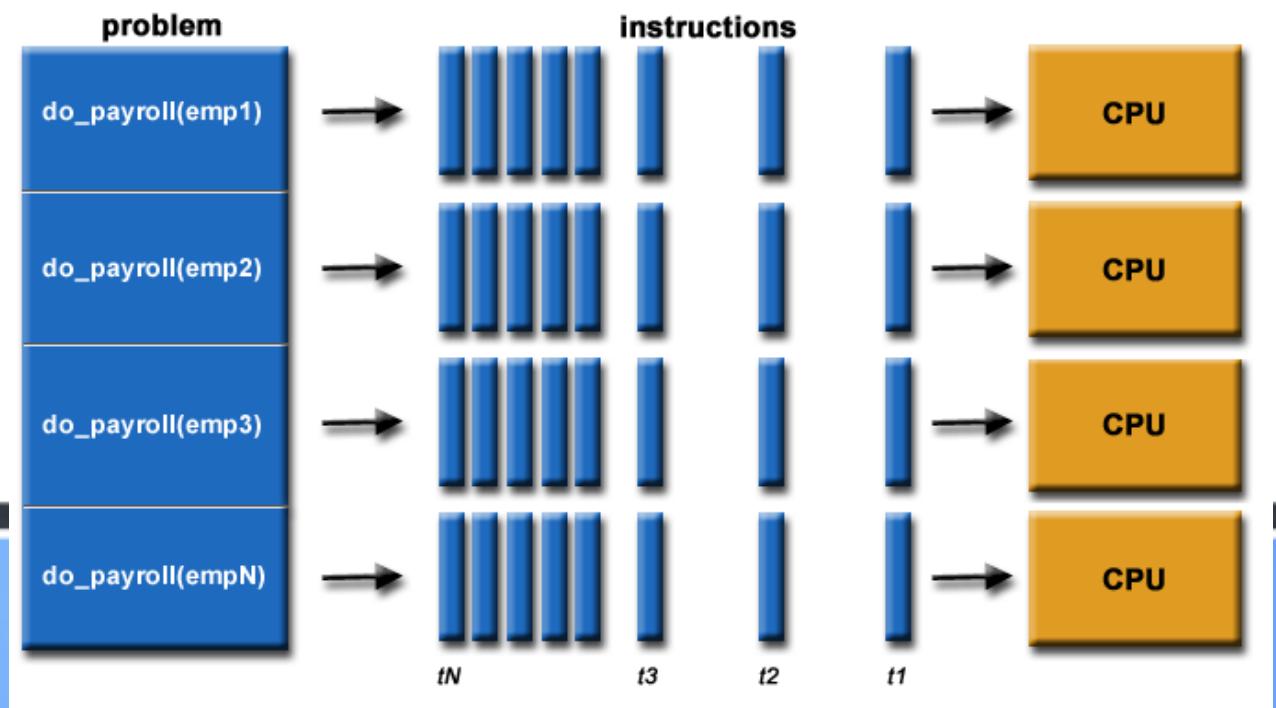
Parallel Computing & Systems - Intro

https://computing.llnl.gov/tutorials/parallel_comp/

Serial Computing



Parallel Computing



Parallel Computing & Systems - Intro

https://computing.llnl.gov/tutorials/parallel_comp/

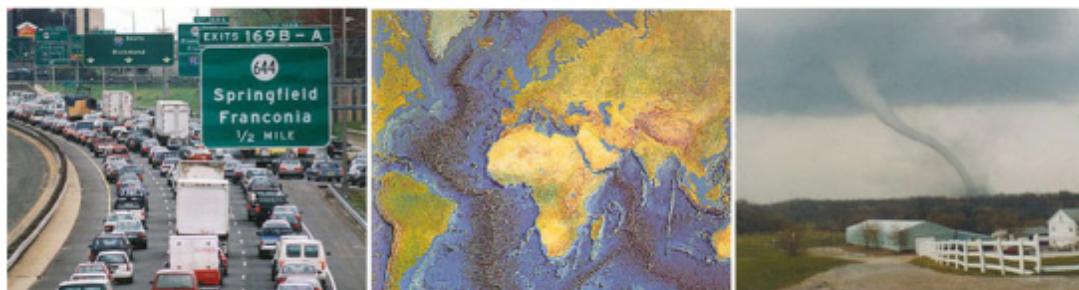
The Real World is Massively Parallel



Galaxy Formation

Planetary Movements

Climate Change



Rush Hour Traffic

Plate Tectonics

Weather



Auto Assembly

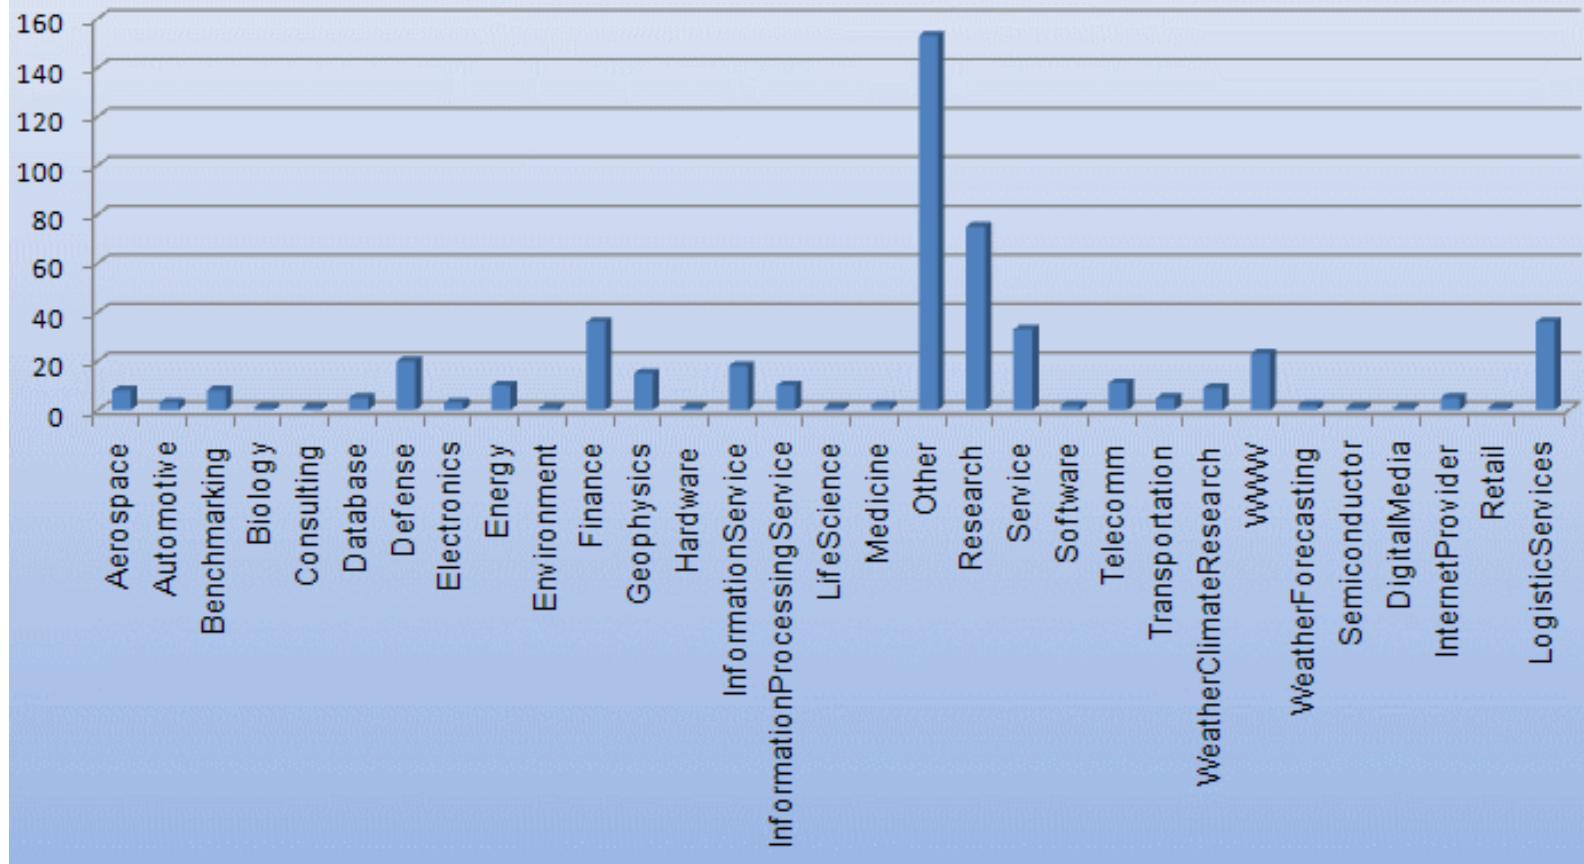
Jet Construction

Drive-thru Lunch

Parallel Computing & Systems - Intro

https://computing.llnl.gov/tutorials/parallel_comp/

Top500 HPC Application Areas



Intel Accepted Technologies Overview

Parallel Programming

C/C++ in Linux with:

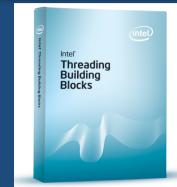
MP – Multi-processing
Programming
(OpenMP)

MPI – Message
Passing Interface
(OpenMPI)

TBB – Thread
Building Blocks
(Intel TBB)

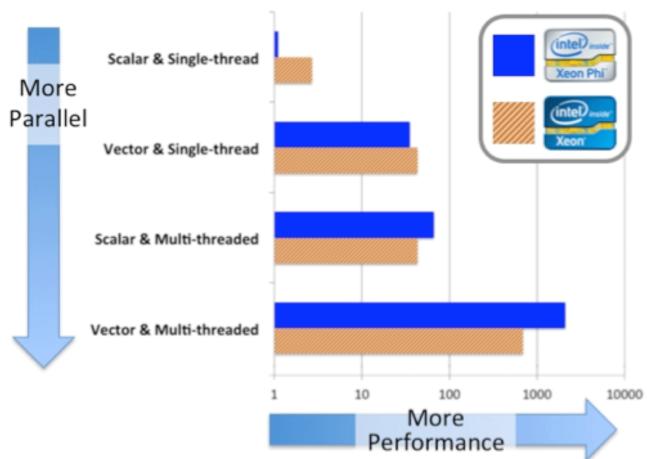
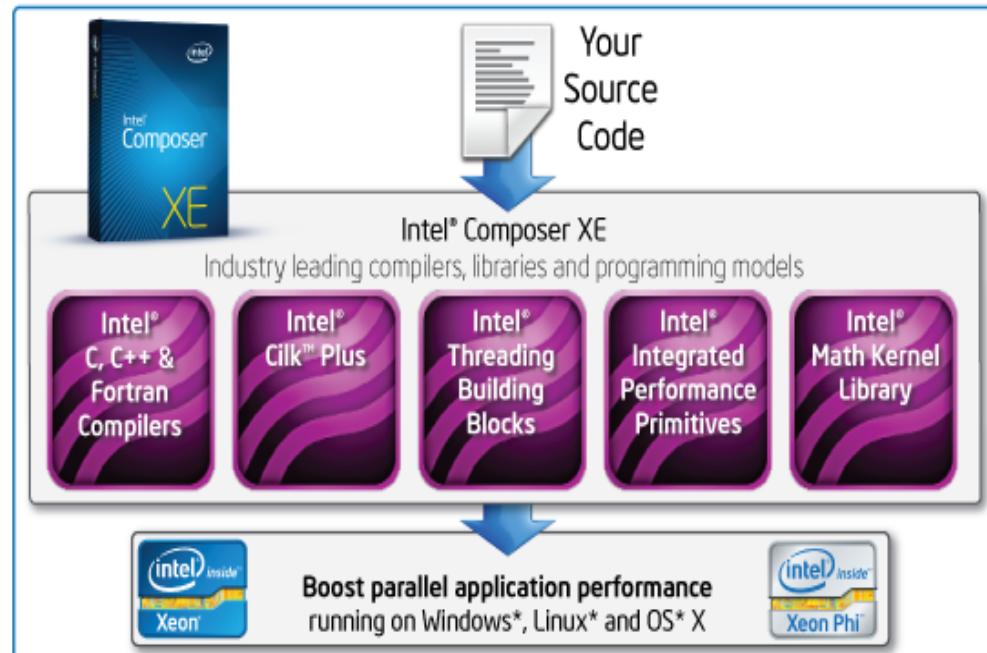
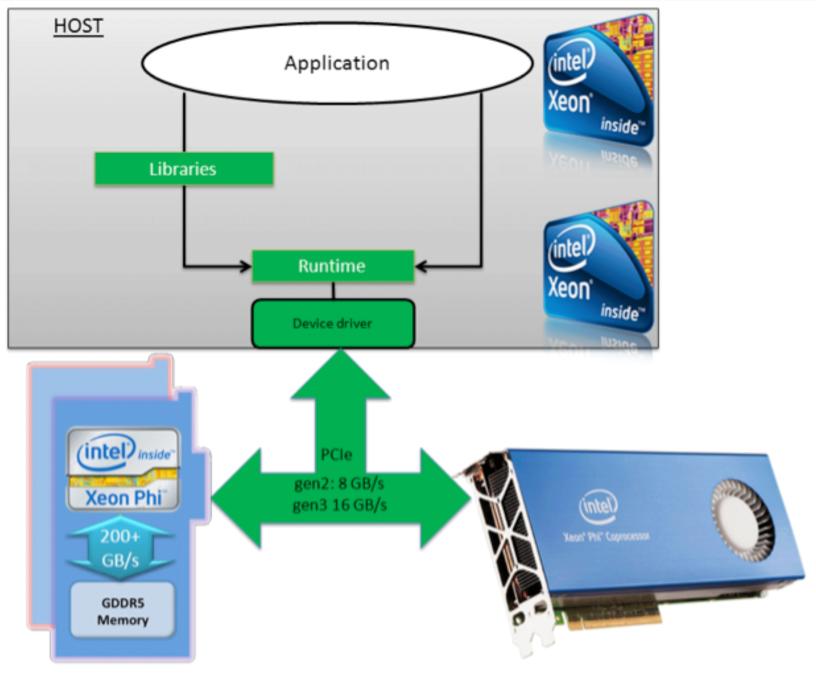
OpenCL – Open
Computing
Language (Intel
OpenCL SDK)

Multi-threaded
Parallel
Computing (Intel
Cilk Plus, POSIX
Threads, C++'11
Multithread)



Intel® Cilk™ Plus
C/C++ compiler extension for simplified parallelism

HW & SW Platform

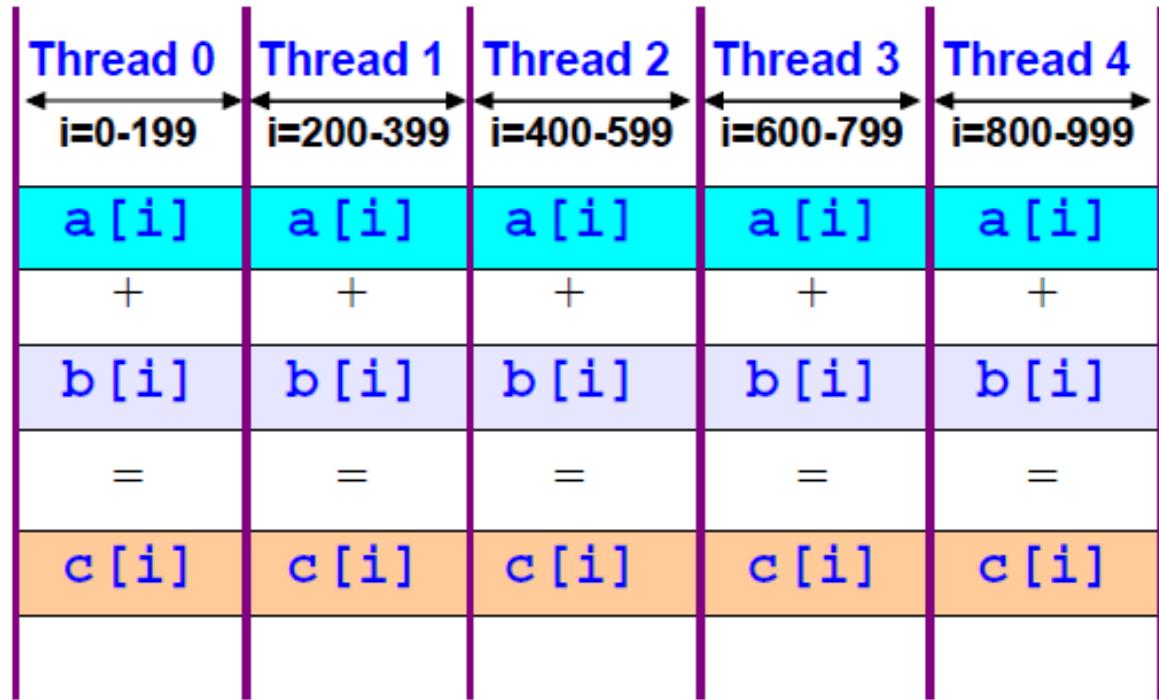


Alternative to:
1. C/C++ Nvidia CUDA
2. C/C++ OpenCL
programming on GPU – video boards

Vector Adding with Parallel Computing

<http://ism.ase.ro/> | <http://acs.ase.ro/java>

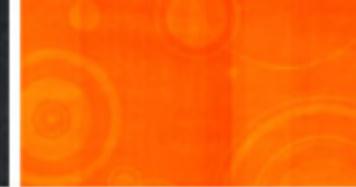
Download the VM-Ware virtual machine with Linux Ubuntu x64 - Intel 2 cores, RAM 6048MB, HDD 25 GB



Adding two vectors sample:

- POSIX Threads
- C++'11 Threads
- Java Threads
- C++ (in OpenMP) – OpenMP mini Tutorial

Parallel Programming Restrictions



Hardware : We use a variety of servers and Xeon PHI accelerators (60 cores, 240 threads). Your code has to be portable between the execution environments, but the large workloads will be run on the Xeon PHI only.

Xeon PHI has its own small operating system and libraries, all you can use is : intel compiler, C/C++, OpenMP, TBB, MPI, Cilk (also available on servers). We only accept native Xeon PHI code for this contest, no offloading. *Late edit : we discourage you from using OpenCL, it's not optimal for this problem and without offloading.*

<http://www.drdobbs.com/parallel/programming-intels-xeon-phi-a-jumpstart/240144160>

Parallel Programming Restrictions

<http://www.drdobbs.com/parallel/programming-intels-xeon-phi-a-jumpstart/240144160>

The source code for C/C++ can be compiled without modification by the Intel compiler (icc), to run in the following modes:

- **Native:** The entire application runs on the Intel Xeon Phi.
- **Offload:** The host processor runs the application and offloads compute intensive code and associated data to the device as specified by the programmer via pragmas in the source code.
- **Host:** Run the code as a traditional OpenMP application on the host.

```
# compile for host-based OpenMP
icc -mkl -O3 -no-offload -openmp -Wno-unknown-pragmas -std=c99 -vec-report3 \
matrix.c -o matrix.omp

# compile for offload mode
icc -mkl -O3 -offload-build -Wno-unknown-pragmas -std=c99 -vec-report3 \
matrix.c -o matrix.off

# compile to run natively on the Xeon Phi
icc -mkl -O3 -mmic -openmp -L /opt/intel/lib/mic -Wno-unknown-pragmas \
-std=c99 -vec-report3 matrix.c -o matrix.mic -liomp5
```

```
1  /* matrix.c (Rob Farber) */
2  #ifndef MIC_DEV
3  #define MIC_DEV 0
4  #endif
5
6  #include <stdio.h>
7  #include <stdlib.h>
8  #include <omp.h>
9  #include <mkl.h>
10 #include <math.h>
11
12 // An OpenMP simple matrix multiply
13 void doMult(int size, float (* restrict A)[size],
14             float (* restrict B)[size], float (* restrict C)[size])
15 {
16 #pragma offload target(mic:MIC_DEV) \
17     in(A:length(size*size)) in( B:length(size*size)) \
18     out(C:length(size*size))
19 {
20     // Zero the C matrix
21 #pragma omp parallel for default(None) shared(C,size)
22     for (int i = 0; i < size; ++i)
23         for (int j = 0; j < size; ++j)
24             C[i][j] = 0.f;
25
26     // Compute matrix multiplication.
27 #pragma omp parallel for default(None) shared(A,B,C,size)
28     for (int i = 0; i < size; ++i)
29         for (int k = 0; k < size; ++k)
30             for (int j = 0; j < size; ++j)
31                 C[i][j] += A[i][k] * B[k][j];
32 }
33 }
```

Open MP

Open specifications for Multi Processing

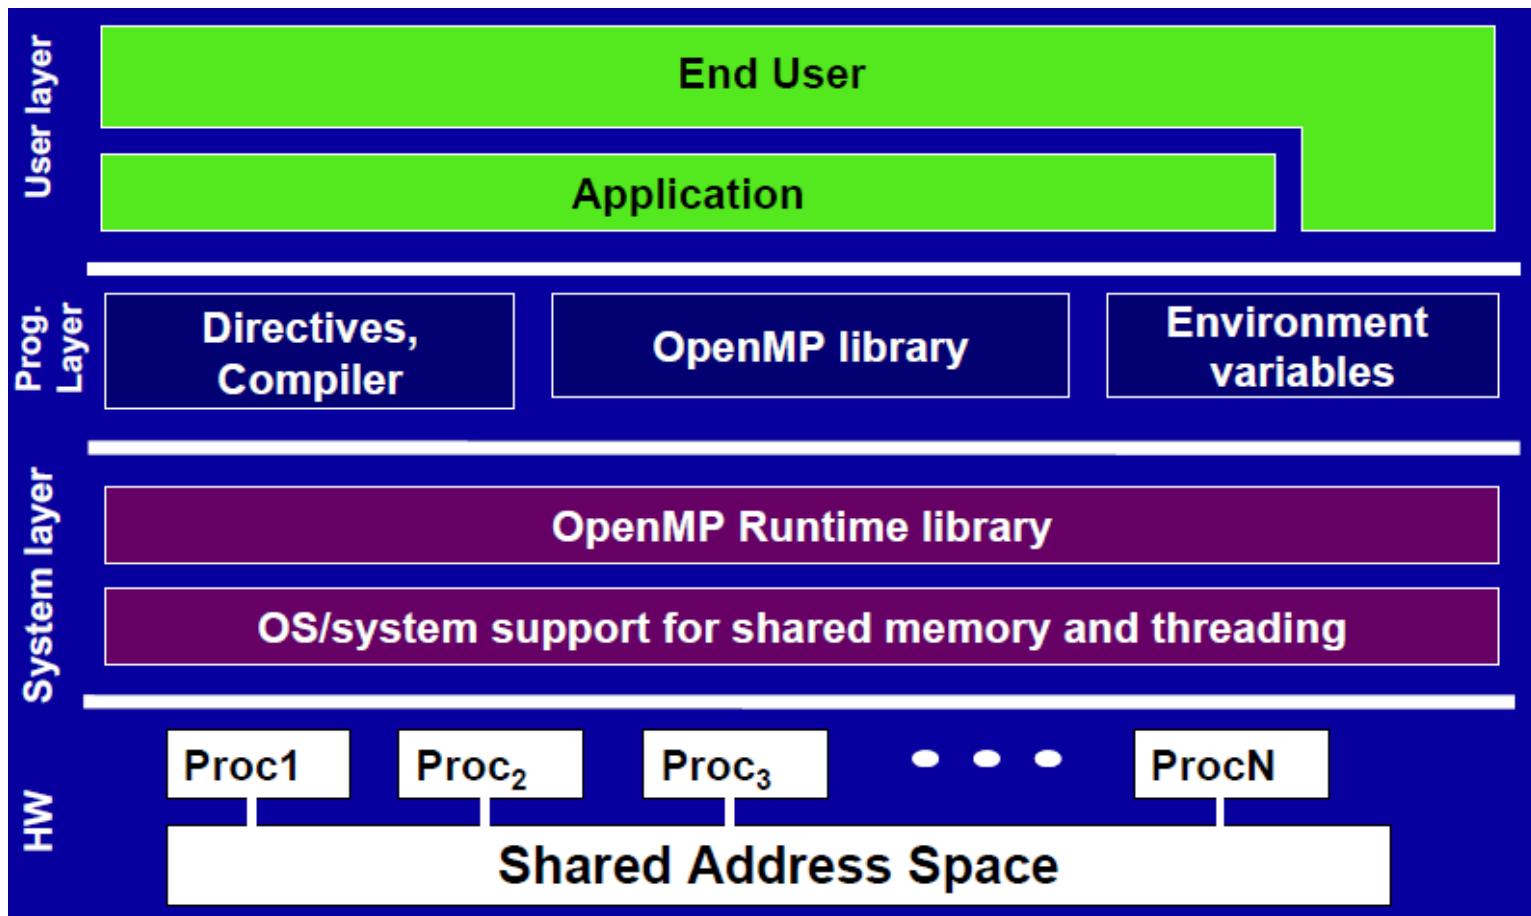
Long version: Open specifications for Multi-Processing via collaborative work between interested parties from the hardware and software industry, government and academia.

- An Application Program Interface (API) that is used to explicitly direct multi-threaded, shared memory parallelism.
- API components:
 - Compiler directives (Compilers that supports: GNU & Intel C/C++ compilers - **gcc/g++ & icc**)
 - Runtime library routines
 - Environment variables
- Portability
 - API is specified for C/C++ and Fortran
 - Implementations on almost all platforms including Unix/Linux and Windows
- Standardization
 - Jointly defined and endorsed by major computer hardware and software vendors
 - Possibility to become ANSI standard

Partial Copyright:

<http://www3.nd.edu/~zxu2/acms60212-40212-S12/Lec-11-01.pdf> | <https://computing.llnl.gov/tutorials/openMP/>

OpenMP Architecture – Version 3



OpenMP Mini-Tutorial – Version 3

Thread

- A **process** is an instance of a computer program that is being executed. It contains the program code and its current activity.
- A **thread of execution** is the smallest unit of processing that can be scheduled by an operating system.
- Differences between threads and processes:
 - A thread is contained inside a process. Multiple threads can exist within the same process and share resources such as memory. The threads of a process share the latter's instructions (code) and its context (values that its variables reference at any given moment).
 - Different processes do not share these resources.

[http://en.wikipedia.org/wiki/Process_\(computing\)](http://en.wikipedia.org/wiki/Process_(computing)) |

Process

- A process contains all the information needed to execute the program:
 - Process ID
 - Program code
 - Data on run time stack
 - Global data
 - Data on heap

Each process has its own address space.

- In multitasking, processes are given time slices in a round robin fashion.
 - If computer resources are assigned to another process, the status of the present process has to be saved, in order that the execution of the suspended process can be resumed at a later time.

OpenMP - Mini-Tutorial – Version 3

Threads Features:

- Thread model is an extension of the process model.
- Each process consists of multiple independent instruction streams (or threads) that are assigned computer resources by some scheduling procedure.
- Threads of a process share the address space of this process.
 - Global variables and all dynamically allocated data objects are accessible by all threads of a process.
- Each thread has its own run time stack, register, program counter.
- Threads can communicate by reading/writing variables in the common address space.

Thread operations include thread creation, termination, synchronization (joins, blocking), scheduling, data management and process interaction.

A thread does not maintain a list of created threads, nor does it know the thread that created it.

All threads within a process share the same address space.

Threads in the same process share:

- Process instructions
- Most data
- open files (descriptors)
- signals and signal handlers
- current working directory
- User and group id

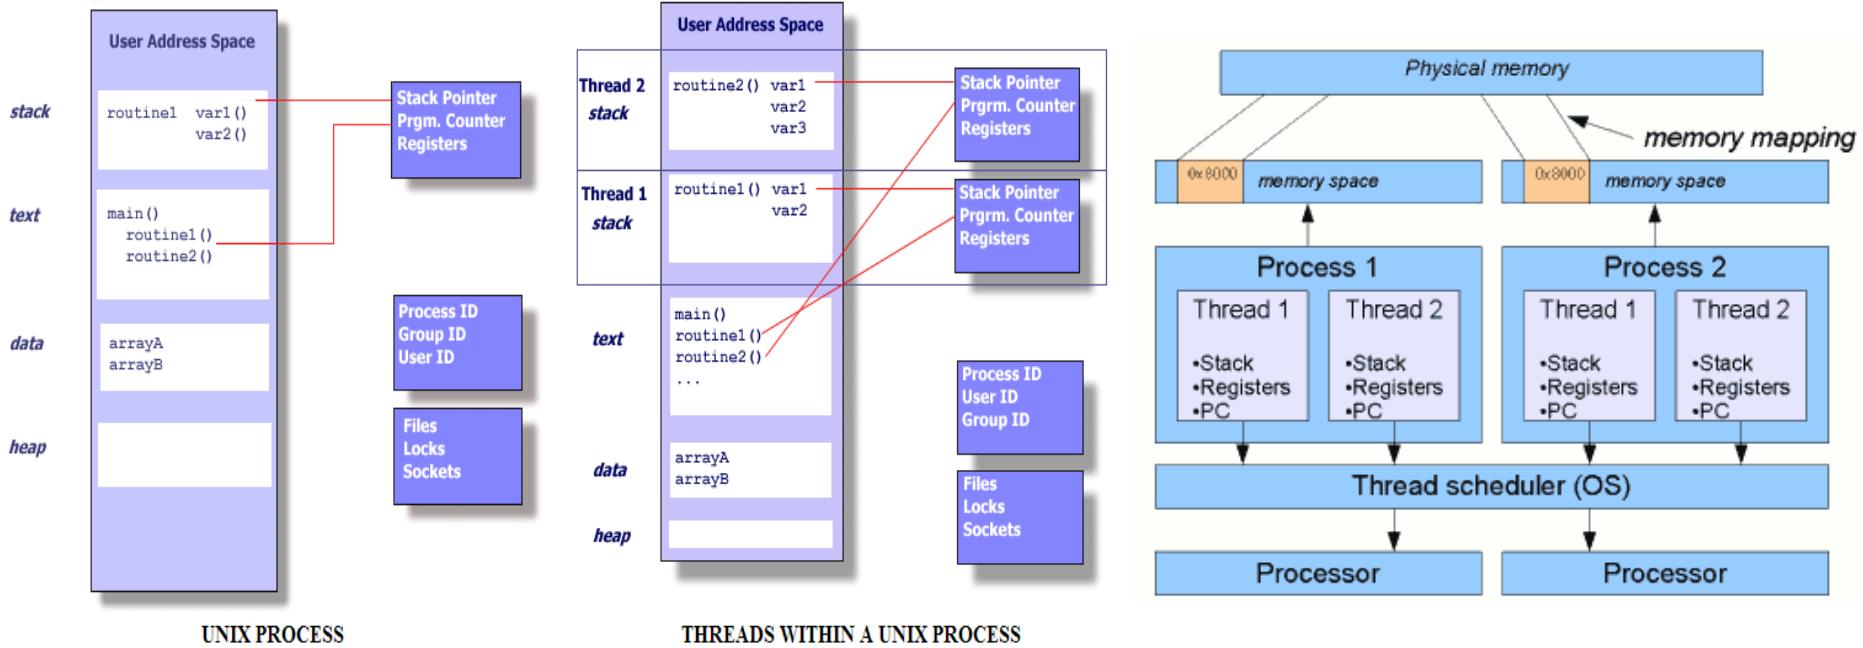
Each thread has a unique:

- Thread ID
- set of registers, stack pointer
- stack for local variables, return addresses
- signal mask
- priority
- Return value: errno

pthread functions return "0" if OK.

OpenMP - Mini-Tutorial – Version 3

Multi-threading vs. Multi-process development in UNIX/Linux:



<https://computing.llnl.gov/tutorials/pthreads/>

http://www.javamex.com/tutorials/threads/how_threads_work.shtml

OpenMP - Mini-Tutorial – Version 3

OpenMP Programming Model:

- Shared memory, thread-based parallelism
 - OpenMP is based on the existence of multiple threads in the shared memory programming paradigm.
 - A shared memory process consists of multiple threads.
- Explicit Parallelism
 - Programmer has full control over parallelization. OpenMP is not an automatic parallel programming model.
- Compiler directive based
 - Most OpenMP parallelism is specified through the use of compiler directives which are embedded in the source code.

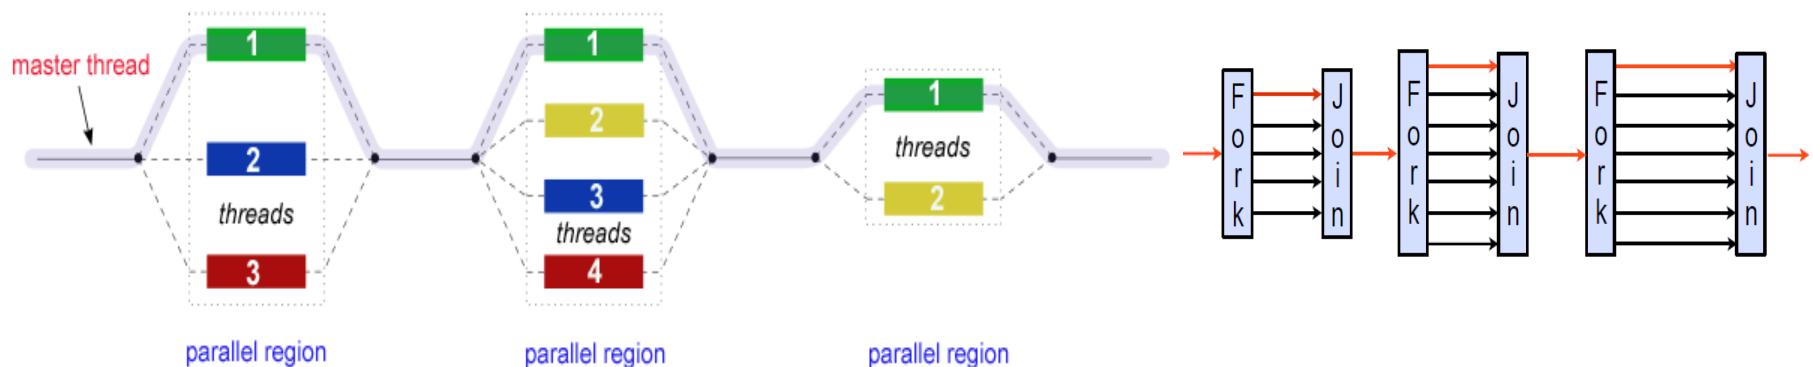
OpenMP is NOT:

- Necessarily implemented identically by all vendors
- Meant for distributed-memory parallel systems (it is designed for shared address spaced machines)
- Guaranteed to make the most efficient use of shared memory
- Required to check for data dependencies, data conflicts, race conditions, or deadlocks
- Required to check for code sequences
- Meant to cover compiler-generated automatic parallelization and directives to the compiler to assist such parallelization
- Designed to guarantee that input or output to the same file is synchronous when executed in parallel.

OpenMP - Mini-Tutorial – Version 3

OpenMP - Fork-Join Parallelism Model:

- OpenMP program begin as a single process: the *master thread (in pictures in red/grey)*. The master thread executes sequentially until the first *parallel region* construct is encountered.
- When a parallel region is encountered, master thread
 - Create a group of threads by **FORK**.
 - Becomes the master of this group of threads, and is assigned the thread id 0 within the group.
- The statement in the program that are enclosed by the *parallel region* construct are then executed in parallel among these threads.
- **JOIN**: When the threads complete executing the statement in the *parallel region* construct, they synchronize and terminate, leaving only the master thread.



OpenMP - Mini-Tutorial – Version 3

I/O

- OpenMP does not specify parallel I/O.
- It is up to the programmer to ensure that I/O is conducted correctly within the context of a multi-threaded program.

Memory Model

- Threads can “cache” their data and are not required to maintain exact consistency with real memory all of the time.
- When it is critical that all threads view a shared variable identically, the programmer is responsible for insuring that the variable is updated by all threads as needed.

//OpenMP Code Structure

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"

int main()
{
    #pragma omp parallel
    {
        int ID = omp_get_thread_num();
        printf("Hello (%d)\n", ID);
        printf(" world (%d)\n", ID);
    }
}
```

Set # of threads for OpenMP:

- In csh:

```
setenv OMP_NUM_THREADS 8
```

- In bash:

```
set OMP_NUM_THREADS=8
export $OMP_NUM_THREADS
```

Compile: g++ -fopenmp hello.c

Run: ./a.out

Compiler / Platform	Compiler	Flag
Intel Linux Opteron/Xeon	icc icpc ifort	-fopenmp
PGI Linux Opteron/Xeon	pgcc pgCC pgf77 pgf90	-mp
GNU Linux Opteron/Xeon IBM Blue Gene	gcc g++ g77 gfortran	-fopenmp

OpenMP - Mini-Tutorial – Version 3

OpenMP Core Syntax

```
#include "omp.h"
void main ()
{
    int var1, var2, var3;
    // 1. Serial code
    ...
    // 2. Beginning of parallel section.
    // Fork a team of threads. Specify variable scoping
    #pragma omp parallel private(var1, var2) shared(var3)
    {
        // 3. Parallel section executed by all threads
        ...
        // 4. All threads join master thread and disband
    }
    // 5. Resume serial code ...
}
```

OpenMP C/C++ Directive Format

- OpenMP directive forms
 - C/C++ use compiler directives
- Prefix: #pragma omp ...
 - A directive consists of a directive name followed by *clauses*

Example:

```
#pragma omp parallel default (shared)
private (var1, var2)
```

OpenMP Directive Format - General Rules:

- Case sensitive
- Only one directive-name may be specified per directive
- Each directive applies to at most one succeeding statement, which must be a structured block.
- Long directive lines can be “continued” on succeeding lines by escaping the newline character with a backslash “\” at the end of a directive line.

OpenMP - Mini-Tutorial – Version 3

OpenMP *parallel* Region Directive

```
#pragma omp parallel [clause list]
```

Typical clauses in [clause list]

- Conditional parallelization
 - **if (scalar expression)**
 - Determine whether the parallel construct creates threads
- Degree of concurrency
 - **num_threads (integer expression)**
 - number of threads to create
- Date Scoping
 - **private (variable list)**
 - Specifies variables local to each thread
 - **firstprivate (variable list)**
 - Similar to the private
 - Private variables are initialized to variable value before the parallel directive
 - **shared (variable list)**
 - Specifies variables that are shared among all the threads
 - **default (data scoping specifier)**
 - Default data scoping specifier may be shared or none

Example:

```
#pragma omp parallel if (is_parallel == 1)
num_threads(8) shared (var_b) private (var_a)
firstprivate (var_c) default (none)
{
/* structured block */
}

▪ if (is_parallel == 1) num_threads(8)
    – If the value of the variable is_parallel is one, create 8 threads
▪ shared (var_b)
    – Each thread shares a single copy of variable var_b
▪ private (var_a) firstprivate (var_c)
    – Each thread gets private copies of variable var_a and var_c
    – Each private copy of var_c is initialized with the value of var_c in main thread when the parallel directive is encountered
▪ default (none)
    – Default state of a variable is specified as none (rather than shared)
    – Signals error if not all variables are specified as shared or private
```

OpenMP - Mini-Tutorial – Version 3

Number of Threads:

- The number of threads in a parallel region is determined by the following factors, in order of precedence:
 - 1.Evaluation of the `if` clause
 - 2.Setting of the `num_threads()` clause
 - 3.Use of the `omp_set_num_threads()` library function
 - 4.Setting of the `OMP_NUM_THREADS` environment variable
 - 5.Implementation default – usually the number of cores on a node
- Threads are numbered from 0 (master thread) to N-1

OpenMP - Mini-Tutorial – Version 3

Thread Creation: Parallel Region Example - Create threads with the parallel construct

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"

int main()
{
    int nthreads, tid;
    #pragma omp parallel num_threads(4) private(tid)
    {
        tid = omp_get_thread_num();
        printf("Hello world from (%d)\n", tid);
        if(tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("number of threads = %d\n", nthreads);
        }
    } // all threads join master thread and terminates
}
```

Clause to request threads

Each thread executes a copy of the code within the structured block

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"
```

```
int main(){
    int nthreads, A[100] , tid;
    // fork a group of threads with each thread having a private tid variable
    omp_set_num_threads(4);
    #pragma omp parallel private (tid)
    {
        tid = omp_get_thread_num();
        foo(tid, A);
    } // all threads join master thread and terminates
}
```

A single copy of A[] is shared between all threads

OpenMP - Mini-Tutorial – Version 3

Work-Sharing Construct:

- A parallel construct by itself creates a “Single Program/Instruction Multiple Data” (SIMD) program, i.e., each thread executes the same code.
- Work-sharing is to split up pathways through the code between threads within a team.
 - Loop construct (for/do)
 - Sections/section constructs
 - Single construct
- Within the scope of a parallel directive, work-sharing directives allow concurrency between iterations or tasks
- ***Work-sharing constructs do not create new threads.***
- A work-sharing construct must be enclosed dynamically within a parallel region in order for the directive to execute in parallel.
- ***Work-sharing constructs must be encountered by all members of a team or none at all.***
- Two directives to be presented
 - – **do/for**: concurrent loop iterations
 - – **sections**: concurrent tasks

OpenMP - Mini-Tutorial – Version 3

Work-Sharing do/for Directive

do/for:

- Shares iterations of a loop across the group
- Represents a “data parallelism”.*

for directive partitions parallel iterations across threads

do is the analogous directive in Fortran

- Usage:

```
#pragma omp for [clause list]
/* for loop */
```

- Implicit barrier at end of for loop

```
#include <stdlib.h>
#include <stdio.h>
#include "omp.h"
void main()
{
    int nthreads, tid;

    omp_set_num_threads(3);

    #pragma omp parallel private(tid)
    {
        int i;
        tid = omp_get_thread_num();
        printf("Hello world from (%d)\n", tid);
        #pragma omp for
        for(i = 0; i <=4; i++)
        {
            printf("Iteration %d by %d\n", i, tid);
        }
    } // all threads join master thread and terminates
}
```

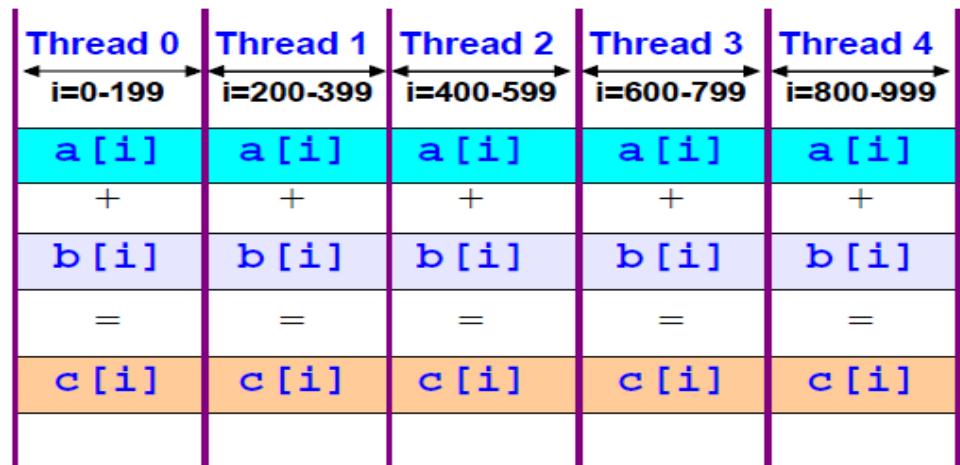
OpenMP - Mini-Tutorial – Version 3

```
//Sequential code to add two vectors:  
for(i=0;i<N;i++) {  
    c[i] = b[i] + a[i];  
}
```

```
//OpenMP implementation 1 (not desired):  
#pragma omp parallel  
{  
    int id, i, Nthrds, istart, iend;  
    id = omp_get_thread_num();  
    Nthrds = omp_get_num_threads();  
    istart = id*N/Nthrds;  
    iend = (id+1)*N/Nthrds;  
  
    if(id == Nthrds-1) iend = N;  
    for(i = istart; i<iend; i++) {  
        c[i] = b[i]+a[i];  
    }  
}
```

```
//A worksharing for construct to add vectors:  
#pragma omp parallel  
{  
    #pragma omp for  
    {  
        for(i=0; i<N; i++) { c[i]=b[i]+a[i]; }  
    }  
}
```

```
//A worksharing for construct to add vectors:  
#pragma omp parallel for  
for(i=0; i<N; i++) { c[i]=b[i]+a[i]; }
```



OpenMP - Mini-Tutorial – Version 3

C/C++ **for** Directive Syntax:

```
#pragma omp for [clause list]
    schedule (type [,chunk])
    ordered
    private (variable list)
    firstprivate (variable list)
    shared (variable list)
    reduction (operator: variable list)
    collapse (n)
    nowait
/* for_loop */
```

For Directive Restrictions

For the “*for loop*” that follows the *for* directive:

- It must not have a break statement
- The loop control variable must be an integer
- The initialization expression of the “*for loop*” must be an integer assignment.
- The logical expression must be one of <,≤,>,≥
- The increment expression must have integer increments or decrements only.

How to combine values into a single accumulation variable (avg)?

//Sequential code to do average value from an array-vector:

```
{
    double avg = 0.0, A[MAX];
    int i;
    ...
    for(i=0; i<MAX; i++) {
        avg += a[i];
    }
    avg /= MAX;
}
```

OpenMP - Mini-Tutorial – Version 3

Reduction Clause

- *Reduction (operator:variable list):*
specifies how to combine local copies of a variable in different threads into a single copy at the master when threads exit.
Variables in *variable list* are implicitly private to threads.
- Operators used in Reduction Clause: +, *, -, &, |, ^, &&, and ||
- Usage Sample:

```
#pragma omp parallel reduction(+: sums) num_threads(4)
{
    /* compute local sums in each thread */
}
/* sums here contains sum of all local instances of sum */
```

Reduction Operators/Initial-Values in C/C++ OpenMP

Operator	Initial Value
+	0
*	1
-	0
&	~0

Operator	Initial Value
	0
^	0
&&	1
	0

Reduction in OpenMP for:

- Inside a parallel or a work-sharing construct:
- A local copy of each list variable is made and initialized depending on *operator* (e.g. 0 for "+")
 - Compiler finds standard reduction expressions containing *operator* and uses it to update the local copy.
 - Local copies are reduced into a single value and combined with the original global value when returns to the master thread.

//A work-sharing for average value from a vector:

```
{
    double avg = 0.0, A[MAX];
    int i;
    ...
    #pragma omp parallel for reduction (+:avg)
    for(i =0; i<MAX; i++) {avg += a[i];}

    avg /= MAX;
}
```

OpenMP - Mini-Tutorial – Version 3

Matrix-Vector Multiplication

```
#pragma omp parallel default (none) \
shared (a, b, c, m,n) private(i,j,sum) num_threads(4)
for(i=0; i < m; i++)
{
    sum=0.0;
    for(j=0; j < n; j++)
        sum+=b[i][j]*c[j];
    a[i]=sum;
}
```

```
for (i=0,1,2,3,4)
    i = 0
    sum = b[i=0] [j]*c[j]
    a[0] = sum
    i = 1
    sum = b[i=1] [j]*c[j]
    a[1] = sum
```

Thread 0,

```
for (i=5,6,7,8,9)
    i = 5
    sum = b[i=5] [j]*c[j]
    a[5] = sum
    i = 6
    sum = b[i=6] [j]*c[j]
    a[6] = sum
```

Thread 1,

...etc...

OpenMP - Mini-Tutorial – Version 3

Matrix-Vector | Matrix-Matrix Multiplication

schedule clause

- Describe how iterations of the loop are divided among the threads in the group. The default schedule is implementation dependent.
- Usage: `schedule (scheduling_class[, parameter])`.
 - **static** - Loop iterations are divided into pieces of size chunk and then statically assigned to threads. If chunk is not specified, the iteration are evenly (if possible) divided contiguously among the threads.
 - **dynamic** - Loop iterations are divided into pieces of size chunk and then dynamically assigned to threads. When a thread finishes one chunk, it is dynamically assigned another. The default chunk size is 1.
 - **guided** - For a chunk size of 1, the size of each chunk is proportional to the number of unassigned iterations divided by the number of threads, decreasing to 1. For a chunk size with value $k(k>1)$, the size of each chunk is determined in the same way with the restriction that the chunks do not contain fewer than k iterations (except for the last chunk to be assigned, which may have fewer than k iterations). The default chunk size is 1.
 - **runtime** - The scheduling decision is deferred until runtime by the environment variable OMP_SCHEDULE. It is illegal to specify a chunk size for this clause
 - **auto** - The scheduling decision is made by the compiler and/or runtime system.

Static scheduling - 16 iterations, 4 threads:

Thread	0	1	2	3
no chunk*	1-4	5-8	9-12	13-16
chunk = 2	1-2 9-10	3-4 11-12	5-6 13-14	7-8 15-16

```
// Static schedule maps iterations to threads at compile time
// static scheduling of matrix multiplication loops
#pragma omp parallel default(private) \
shared (a, b, c, dim) num_threads(4)
#pragma omp for schedule(static)
for(i=0;i < dim;i++)
{
    for(j=0;j < dim;j++)
    {
        c[i][j] = 0.0;
        for(k=0;j < dim;k++)
            c[i][j] += a[i][k]*b[k][j];
    }
}
```



Questions & Answers!

But wait...
There's More!

- 1. DAD - Is what you expected?**
- 2. How many hours per week are you going to invest in order to achieve DAD goals?**
- 3. How many of you are working in IT field – SW Dev., Admin., Designers?**
- 4. What bachelor programs are you graduated from?**
- 5. How many students get the payment scholarship from the companies vs. how many are/aren't paying the studies?**
- 6. In what disciplines did we collaborate together?**



Thanks!



DAD – Distributed Application Development
End of Lecture 1

